

BZAN 6354

Live Lecture 5

February 19, 2024

Dr. Mark Grimes, Ph.D.
gmgrimes@bauer.uh.edu

UNIVERSITY of
HOUSTON

C. T. BAUER COLLEGE of BUSINESS
Department of Decision & Information Sciences

Agenda

- Administration
- Module 3.2 – Decomposed Design-Specific ERD models
- Module 3.3 – Data Modeling Errors
- Building on the MGHH Model
- Break
- Module 6.1 – The Relational Data Model
- Module 6.2 – Characteristics of a Relation
- Module 6.3 – Data Integrity Constraints ← We may not get quite this far, we'll see...

Assignment 1

- Grades are posted
 - Average 23 / 25
- Overall good work
- I left brief comments on most submissions - if you want detailed feedback send me an email
 - Attributes like “Phone Number” should be text, not numeric
 - Attributes like “Age” should be derived, not stored (Date of Birth should be stored)
 - Neatness counts!
- Any questions/thoughts on the assignment?

Assignment 2

- Was posted last week, but you need some knowledge from today
 - Weak entities and partial keys
- Creating an ERD for “Shampooch Doggy Spa”
 - Different than “Dave’s Dog Wash”
- Just like with assignment 1, there is a walkthrough video to help you along the way / let you check your work
- Due one week from today – February 26

Exam 1

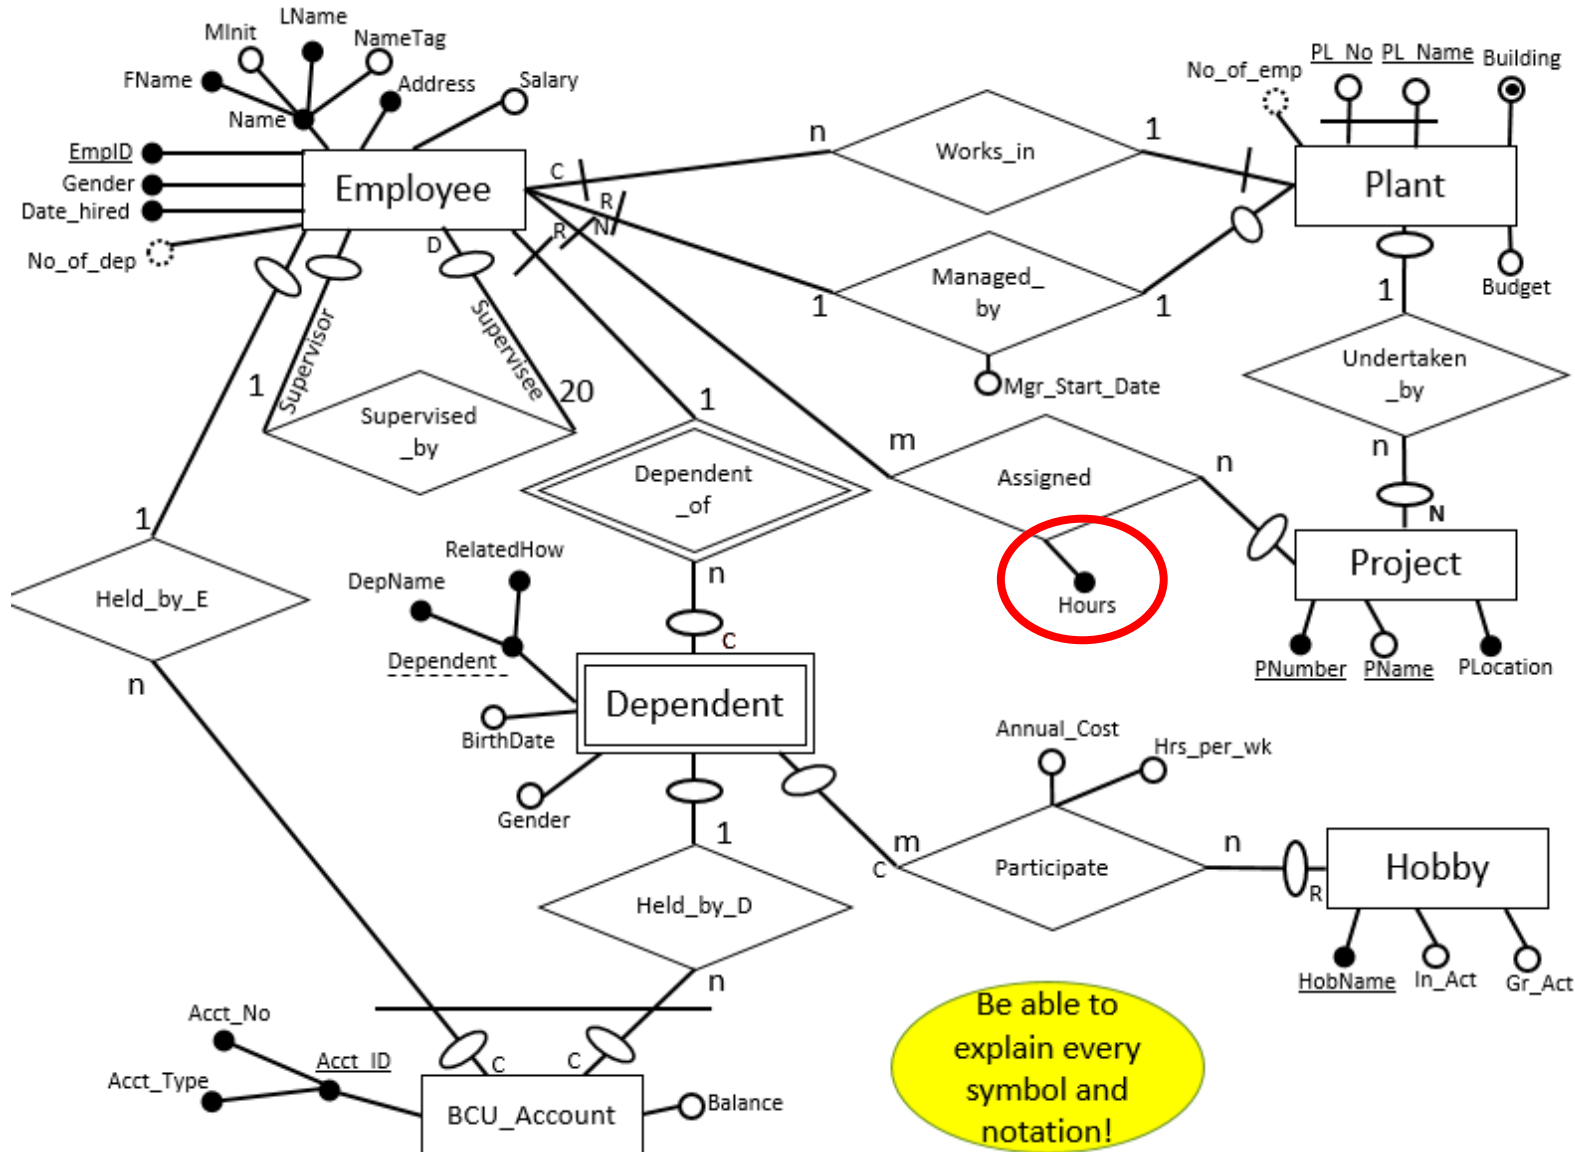
- Coming up in two weeks!
 - Learning objectives from each module are useful as a “study guide”
 - Weekly Progress Quizzes are good practice
- Will be about
 - 1/3 Multiple Choice
 - 1/3 Fill in the blank/matching/etc.
 - 1/3 Drawing an ERD

Module 3.2.4

Decomposed Design-Specific ERD

- Attribute Placement
- Weak entities
- Capturing Multi-value attributes
- Modeling m:n relationships




Attribute Placement



Be able to explain every symbol and notation!


Attribute Placement

- Attributes may not only be part of an entity, but may also be part of the relationship
- Rent on a dorm room is the same no matter who is in it
- A student pays the same rent no matter what dorm room they have
- Rent varies based on student and room type




Attribute Placement

- Room A costs \$1,200 per semester
 - Room A is really nice
- Room B costs \$1,000 per semester
 - Room B is a normal dorm room
- Room C costs \$800 per semester
 - Room C is a smaller room in an older dormitory



Attributes in a relationship


- Doug pays \$600 per semester, no matter what room
 - Maybe scholarship? Financial need?
- Eugene pays \$1,000 per semester, no matter what room
 - Perhaps a discount for having good grades
- Fiona pays \$1,400 per semester, no matter what room
 - Normal rate



Attributes in a relationship

- Rate depends both on how nice the room is and the student's details
- Which you pick depend on your.... Business Rules

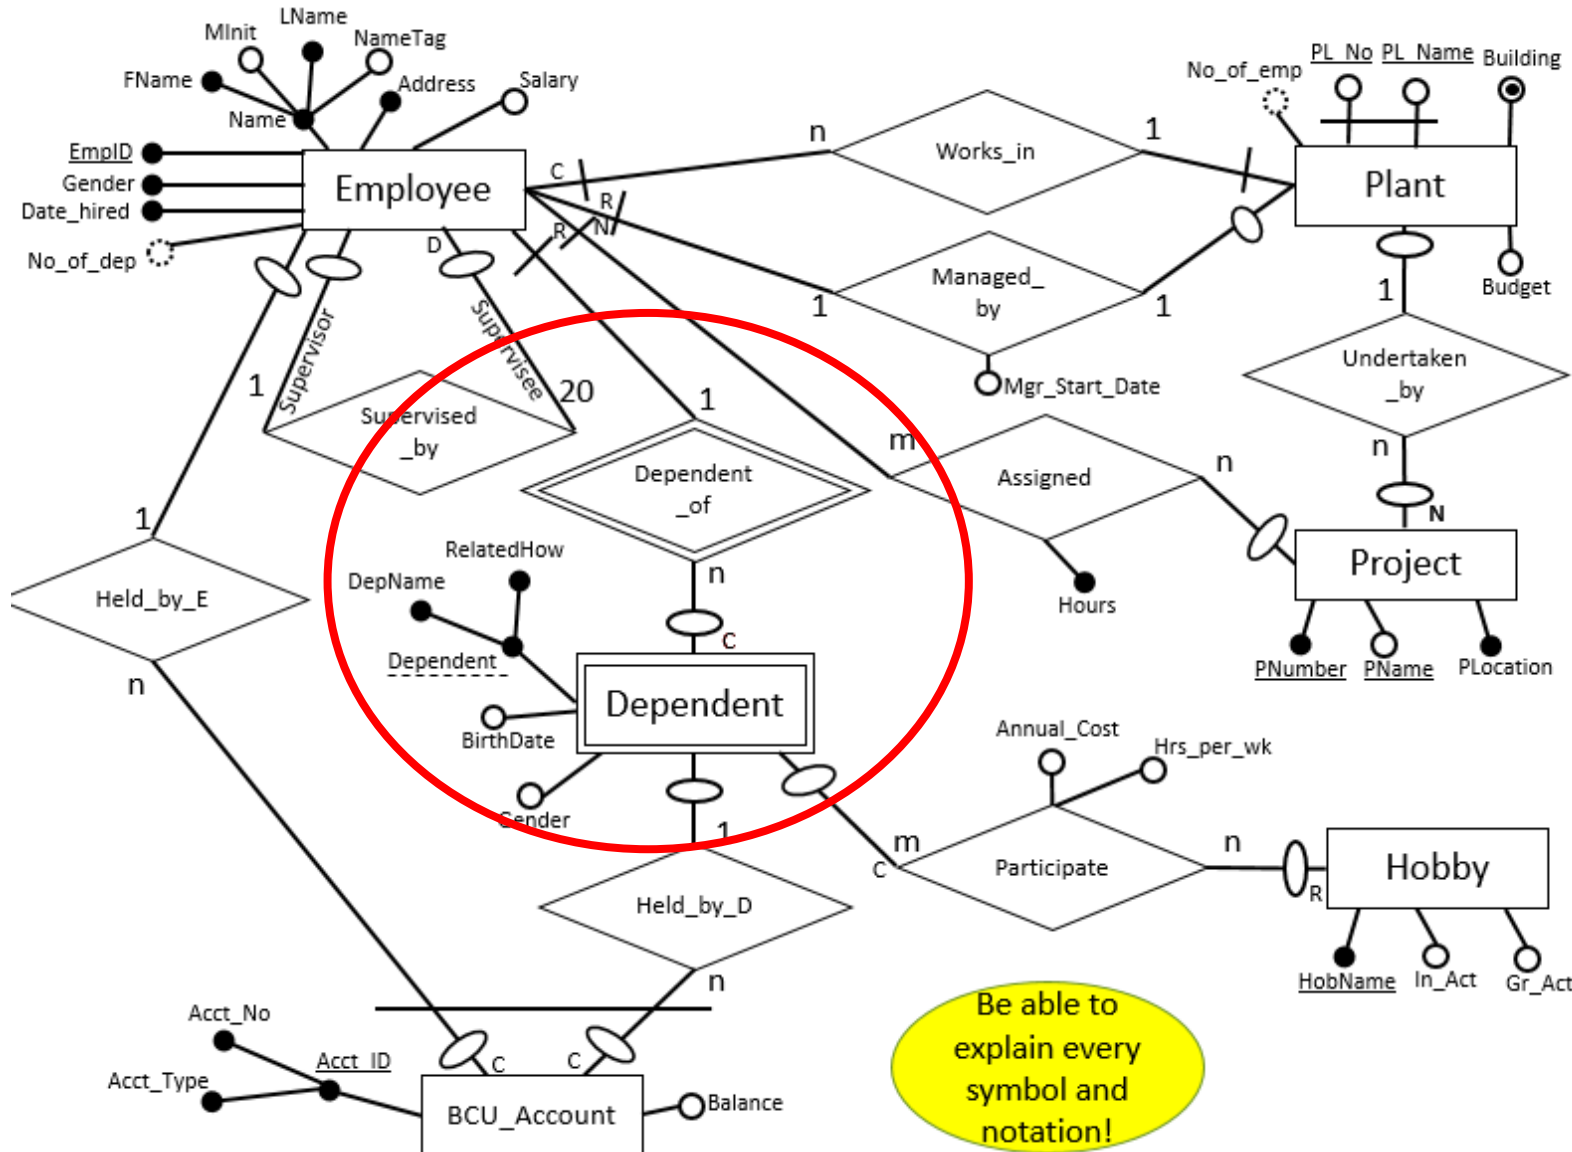
	Room A	Room B	Room C
Doug	\$800	\$700	\$500
Eugene	\$1,400	\$1,200	\$800
Fiona	\$1,600	\$1,400	\$1,000



Base/Strong vs. Weak Entity Types


- Base (or strong) entity types are those where the entities have independent existence (i.e., each entity of this entity type can be uniquely identified)
 - A base entity type has a unique identifier.
- Weak entity types are those where entities do not have an independent existence
 - Duplicate entity instances may be present
 - A weak entity type does not have a unique identifier
 - Must be related to a base (strong) entity to be identified – known as an “identifying relationship”
 - A weak entity type has a “**partial key**” – also known as a “discriminator”

Base/Strong vs. Weak Entity Types



Be able to explain every symbol and notation!

Weak entity type: Classroom as an example



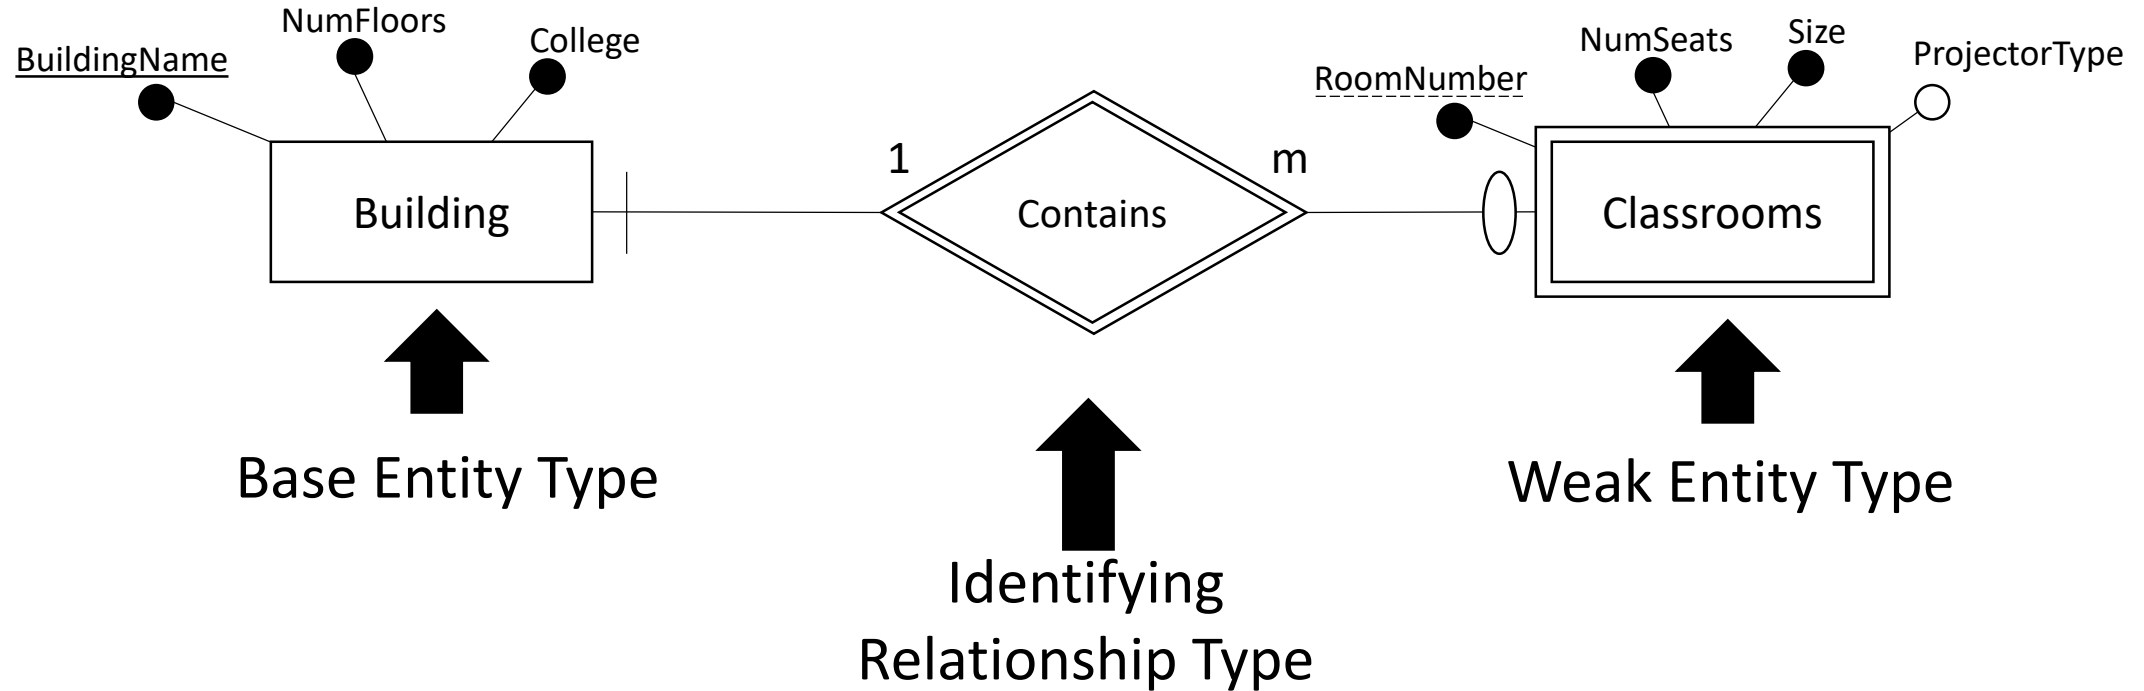
Building		
BuildingName	NumFloors	College
Melcher	3	Business
CBB	5	Business
Cemo	2	Business



Classrooms			
RoomNumber	NumSeats	Size	ProjectorType
110	66	900	Epson MX324
111	66	900	Epson MX324
112	66	900	ViewSonix P32
108	25	400	Epson Slidemaster
110	20	450	
214	120	2200	
226	80	1100	ViewSonix P32
109	85	1100	Epson MX324
111	120	1350	
112	80	900	


Note: A similar example with “apartments” is on page 53 of the book

Base/Strong vs. Weak Entity Types




- Double square = weak entity
- Double diamond = identifying relationship

Partial key (Discriminator) Defined



- An attribute, atomic or composite, in a weak entity type, which in conjunction with a unique identifier(s) of the parent entity type(s) in the identifying relationship type(s), uniquely identifies weak entities is called the partial key (discriminator) of a weak entity type.
- Primary key (solid underline) for Building table: **BuildingName**
- Partial key (dotted underline) for Classrooms table: **RoomNumber**
- To uniquely identify a classroom , you must specify **BOTH** the **BuildingName** AND **RoomNumber**
- This is the Identifying relationship


A reminder



- Database design is both a SCIENCE and an ART
- There are other ways to model this depending on:
 - Business rules
 - Performance considerations
 - Personal preference

Two issues that change the way you look at ERDs

- Multi-value attributes don't really exist*
 - Exist in conceptual models, but most DBMS will only allow for a single value
 - Can be solved by creating a 1:m or m:n relationship
- M:N relationships don't really exist*
 - Remember, they are really two 1:m and m:1 relationships



* They exist CONCEPTUALLY, but cannot be technically implemented in a database

Modeling a multi-value attribute

- We understand values as being at the intersection of a column and a row
- ...but what about when we have “multi-value” attributes?


Horses				
Name	Color	Spots	Sex	Weight
Sam	Brown	No	F	1500
Erica	Yellow	Yes	F	920
John	Grey	No	M	1800
Trotty	Brown	Yes	M	1300
Rio	Grey	No	F	1700
Robin	Yellow	No	M	1100
Katy	Brown	No	F	1200
Pegasus	Brown	No	M	1750

Values



The diagram illustrates the concept of values as being at the intersection of a column and a row. Red arrows point from the word "Values" (underlined) to specific cells in the table. The arrows point to the following cells: (Name, Katy), (Color, Brown), (Spots, No), (Sex, F), (Weight, 1200), (Name, Pegasus), (Color, Brown), (Spots, No), (Sex, M), and (Weight, 1750). This demonstrates how a single value (like "Brown" or "No") can appear in multiple rows, and how a single row contains multiple values.


Modeling a multi-value attribute



Plant			
<u>PL_No</u>	PL_Name	Budget	Building
10	Underwood	3000000	1
			2
			3
11	Garnett	3000000	1
			2
12	Belmont	3500000	1
13	Vanderbilt	3500000	1
			2

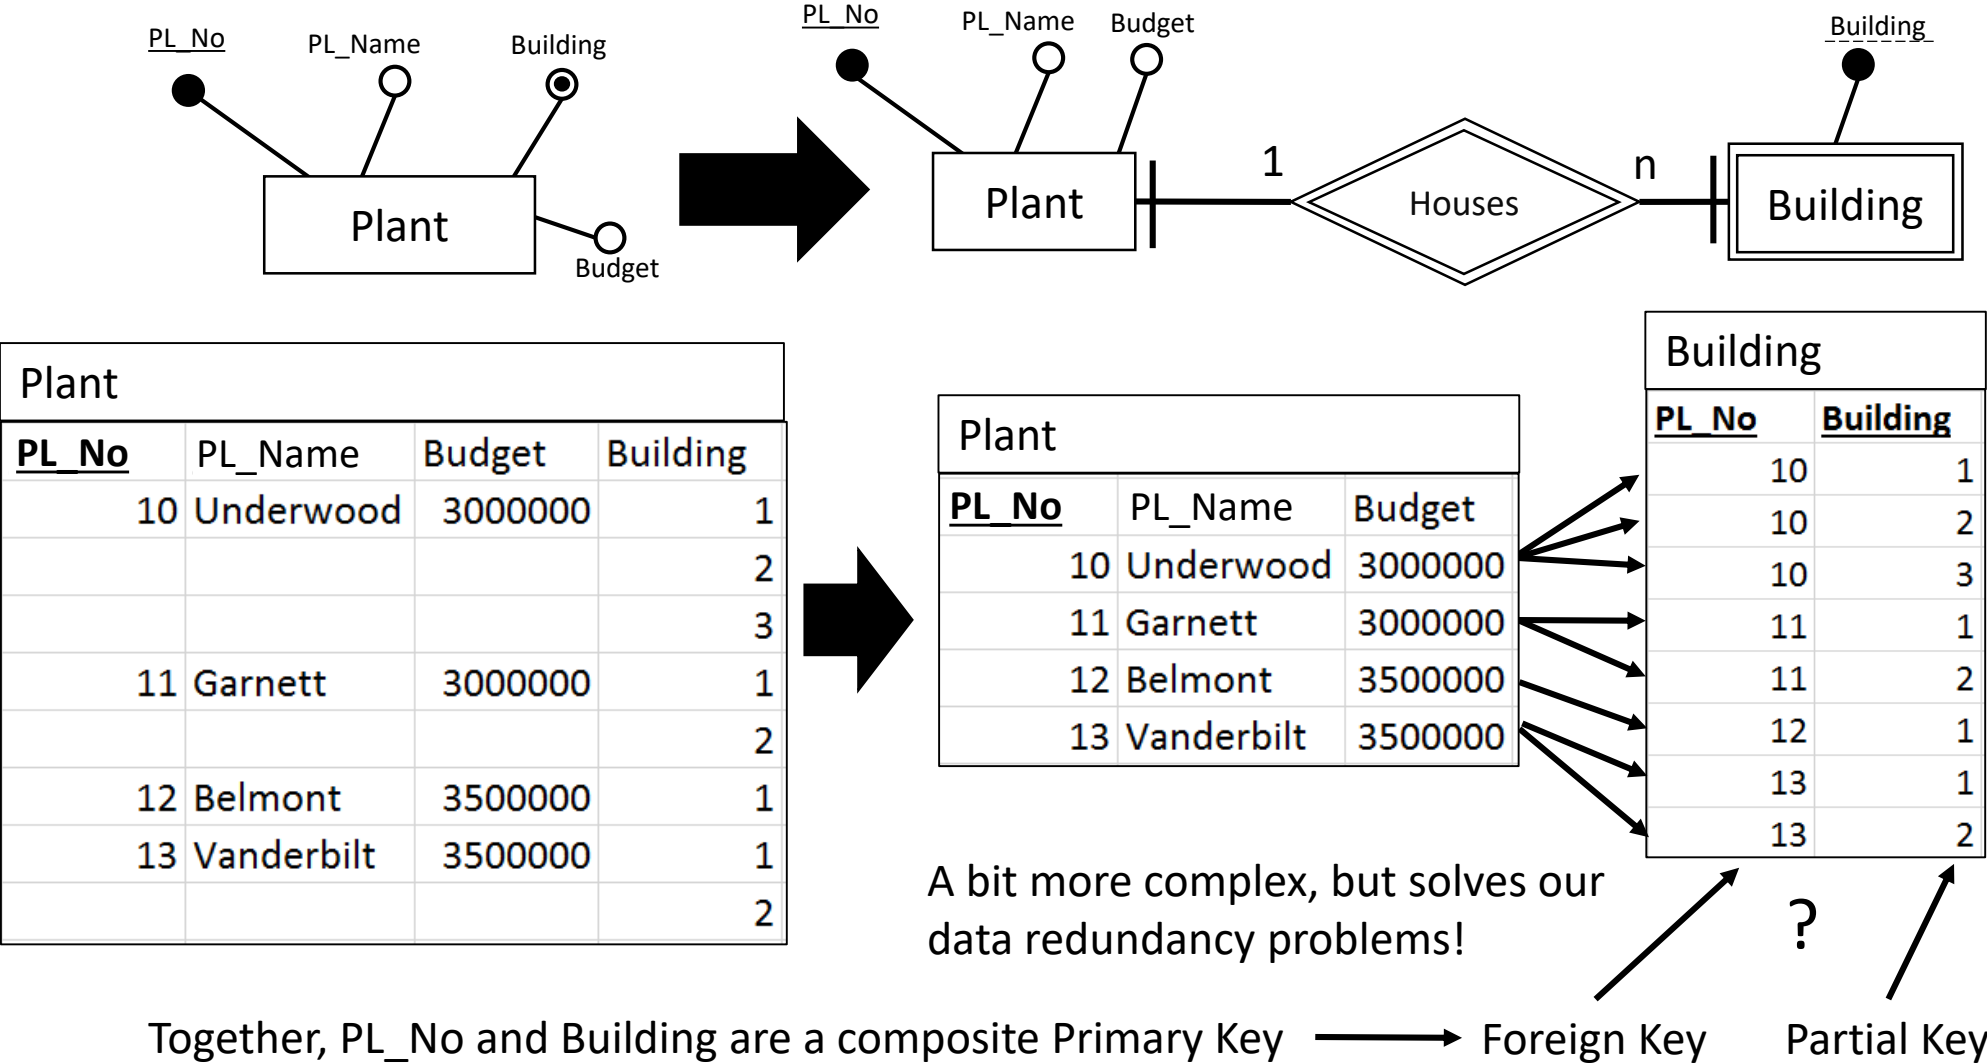
- This is basically what we have been saying, but this functionally does not work
- We have rows of data with required attributes that are not populated
- Note: We can ignore the derived attributes for now

Modeling a multi-value attribute – Option #1



Technically works, but not a good solution - Data Redundancy!

Modeling a multi-value attribute – Option #2 (correct!)



Important note

- While this is “correct” for most business → purposes as it reduces redundancy...
 - Data is “normalized”
- ...there is more “computational complexity” in the “correct” approach
- For data mining/data warehouse applications, this may be desirable →
 - Pro: Faster/Less complex
 - Con: Less efficient use of storage
 - Con: More difficult to update...
 - ...but we infrequently “update” in a data warehouse, so perhaps not an issue!
- What you choose depends on the business’s needs!

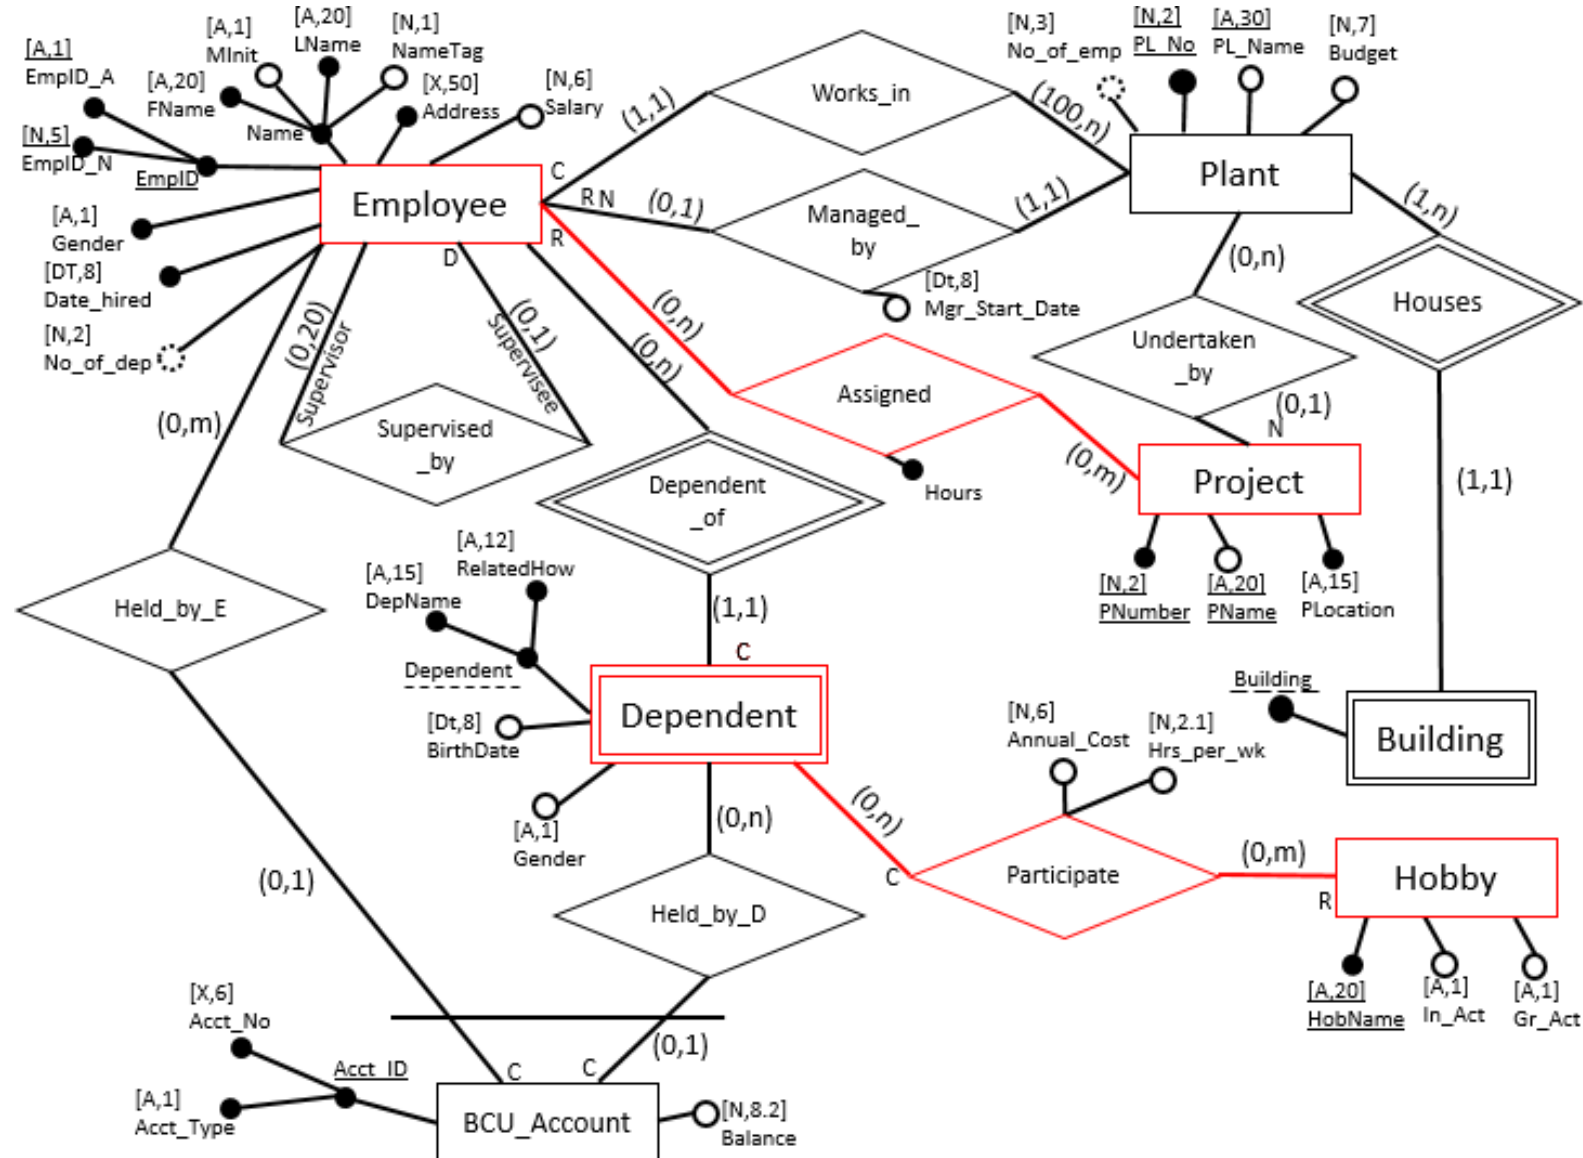
Plant			Building	
<u>PL No</u>	PL_Name	Budget	<u>PL No</u>	<u>Building</u>
10	Underwood	3000000	10	1
11	Garnett	3000000	10	2
12	Belmont	3500000	10	3
13	Vanderbilt	3500000	11	1
			11	2
			12	1
			13	1
			13	2

Plant			
<u>PL No</u>	<u>Building</u>	PL_Name	Budget
10	1	Underwood	3000000
10	2	Underwood	3000000
10	3	Underwood	3000000
11	1	Garnett	3000000
11	2	Garnett	3000000
12	1	Belmont	3500000
13	1	Vanderbilt	3500000
13	2	Vanderbilt	3500000

Modeling a many to many relationship


We have two

- Employee-Project
 - Many employees are assigned to a project
 - A project has many employees assigned to it
- Dependent-Hobby
 - A dependent may have many hobbies
 - A hobby may have many dependents that participate



Let's do dependent-hobby

- Ignore for the moment that dependent is a weak entity
 - Note: I am leaving out many attributes for simplicity of the example – they have not really disappeared



Dependents		
<u>Name</u>	Birthdate	Gender
Mark	5/11/1945	M
Jill	5/4/1976	F
Norm	3/1/1984	M
Mike	1/31/1992	M
Sue	9/2/1990	F

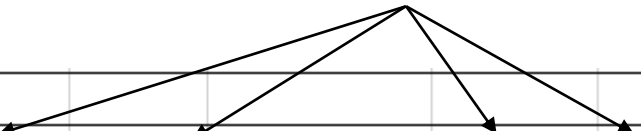
Hobby		
<u>HobbyName</u>	In_Act	Gr_Act
Soft Ball	O	G
Flag Football	O	G
Knitting	I	I
Cycling	O	I
Movies	I	G

- Do employee-project on your own for practice
 - It's in the book on page 116

Let's do dependent-hobby

- Name-HobbyName makes for a fine key value
- What's wrong here?

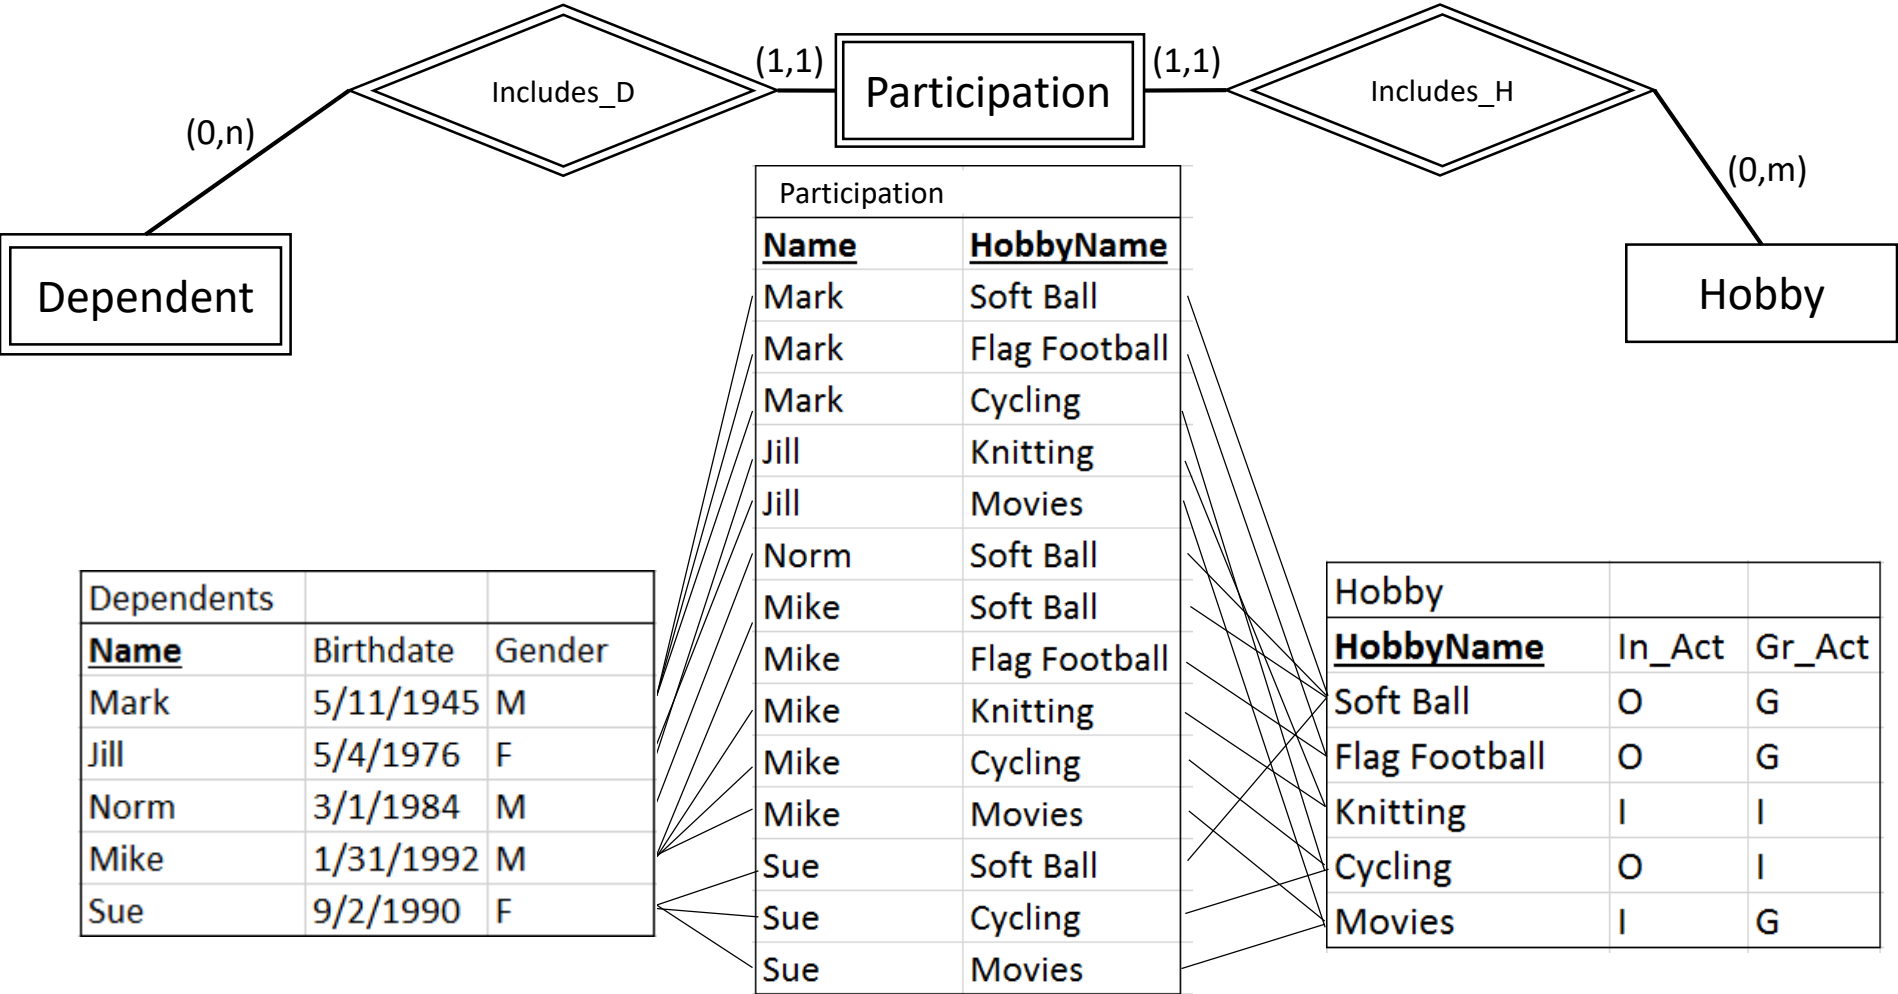
Lots of redundant data



Participates (relationship)					
<u>Name</u>	Birthdate	Gender	<u>HobbyName</u>	In_Act	Gr_Act
Mark	5/11/1945	M	Soft Ball	O	G
Mark	5/11/1945	M	Flag Football	O	G
Mark	5/11/1945	M	Cycling	O	I
Jill	5/4/1976	F	Knitting	I	I
Jill	5/4/1976	F	Movies	I	G
Norm	3/1/1984	M	Soft Ball	O	G
Mike	1/31/1992	M	Soft Ball	O	G
Mike	1/31/1992	M	Flag Football	O	G
Mike	1/31/1992	M	Knitting	I	I
Mike	1/31/1992	M	Cycling	O	I
Mike	1/31/1992	M	Movies	I	G
Sue	9/2/1990	F	Soft Ball	O	G
Sue	9/2/1990	F	Cycling	O	I
Sue	9/2/1990	F	Movies	I	G

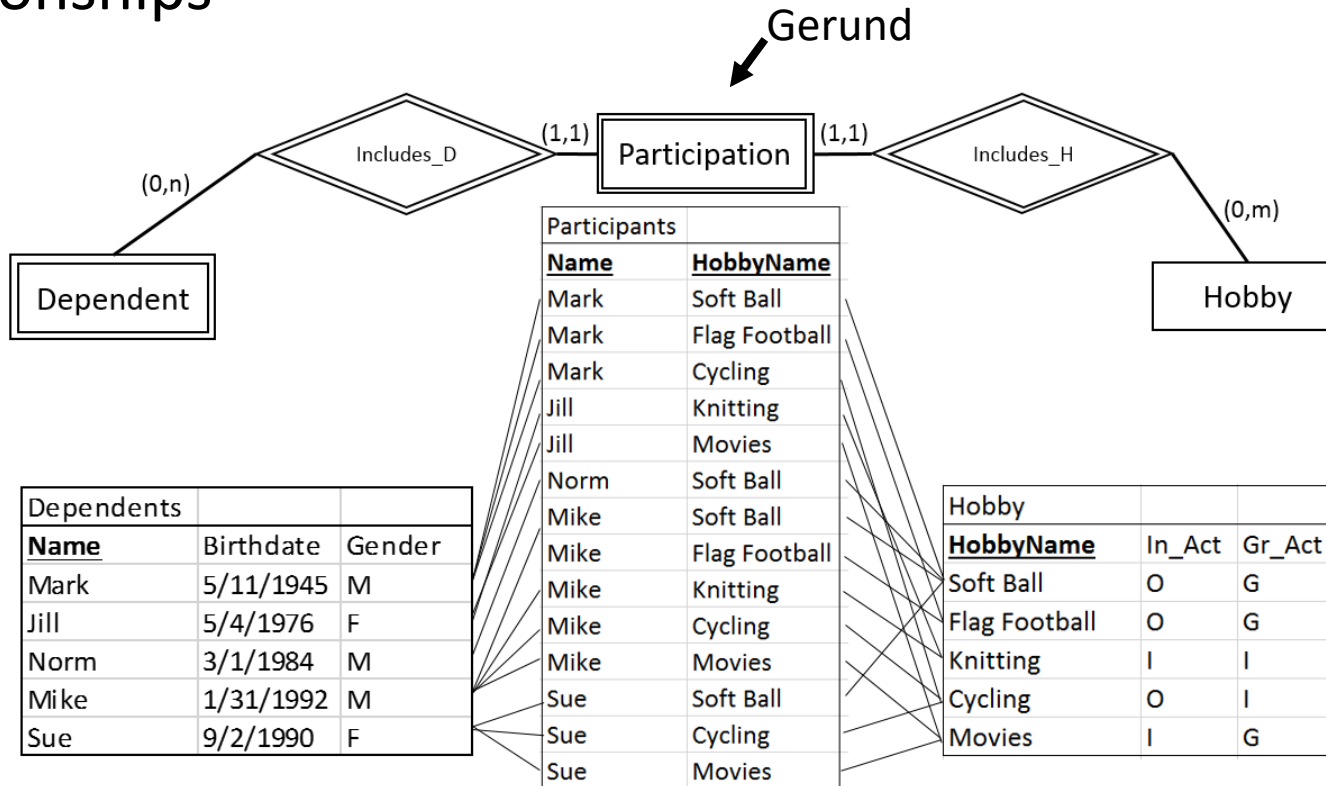
Let's do dependent-hobby

- Decomposed the M:N relationship into two 1:M relationships
- The Participation entity is a **gerund**



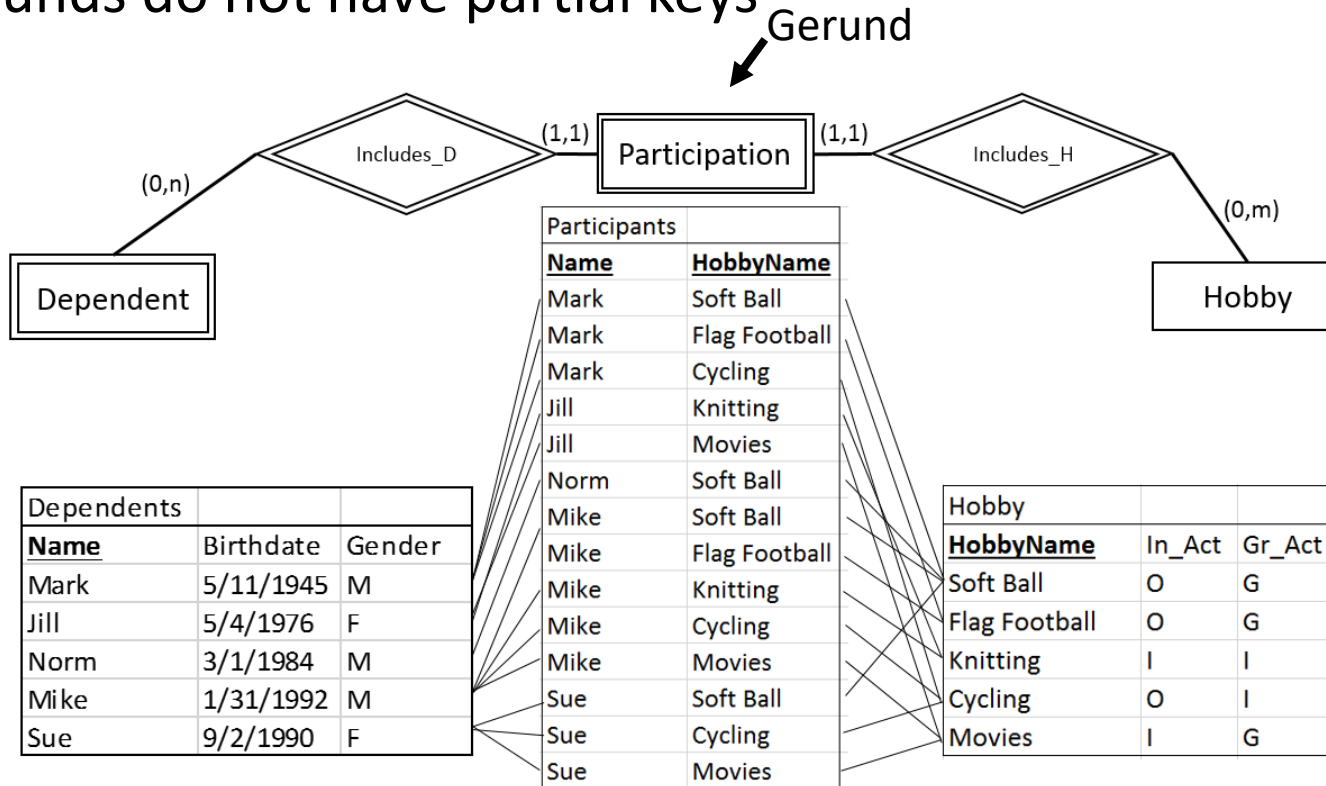
The Gerund

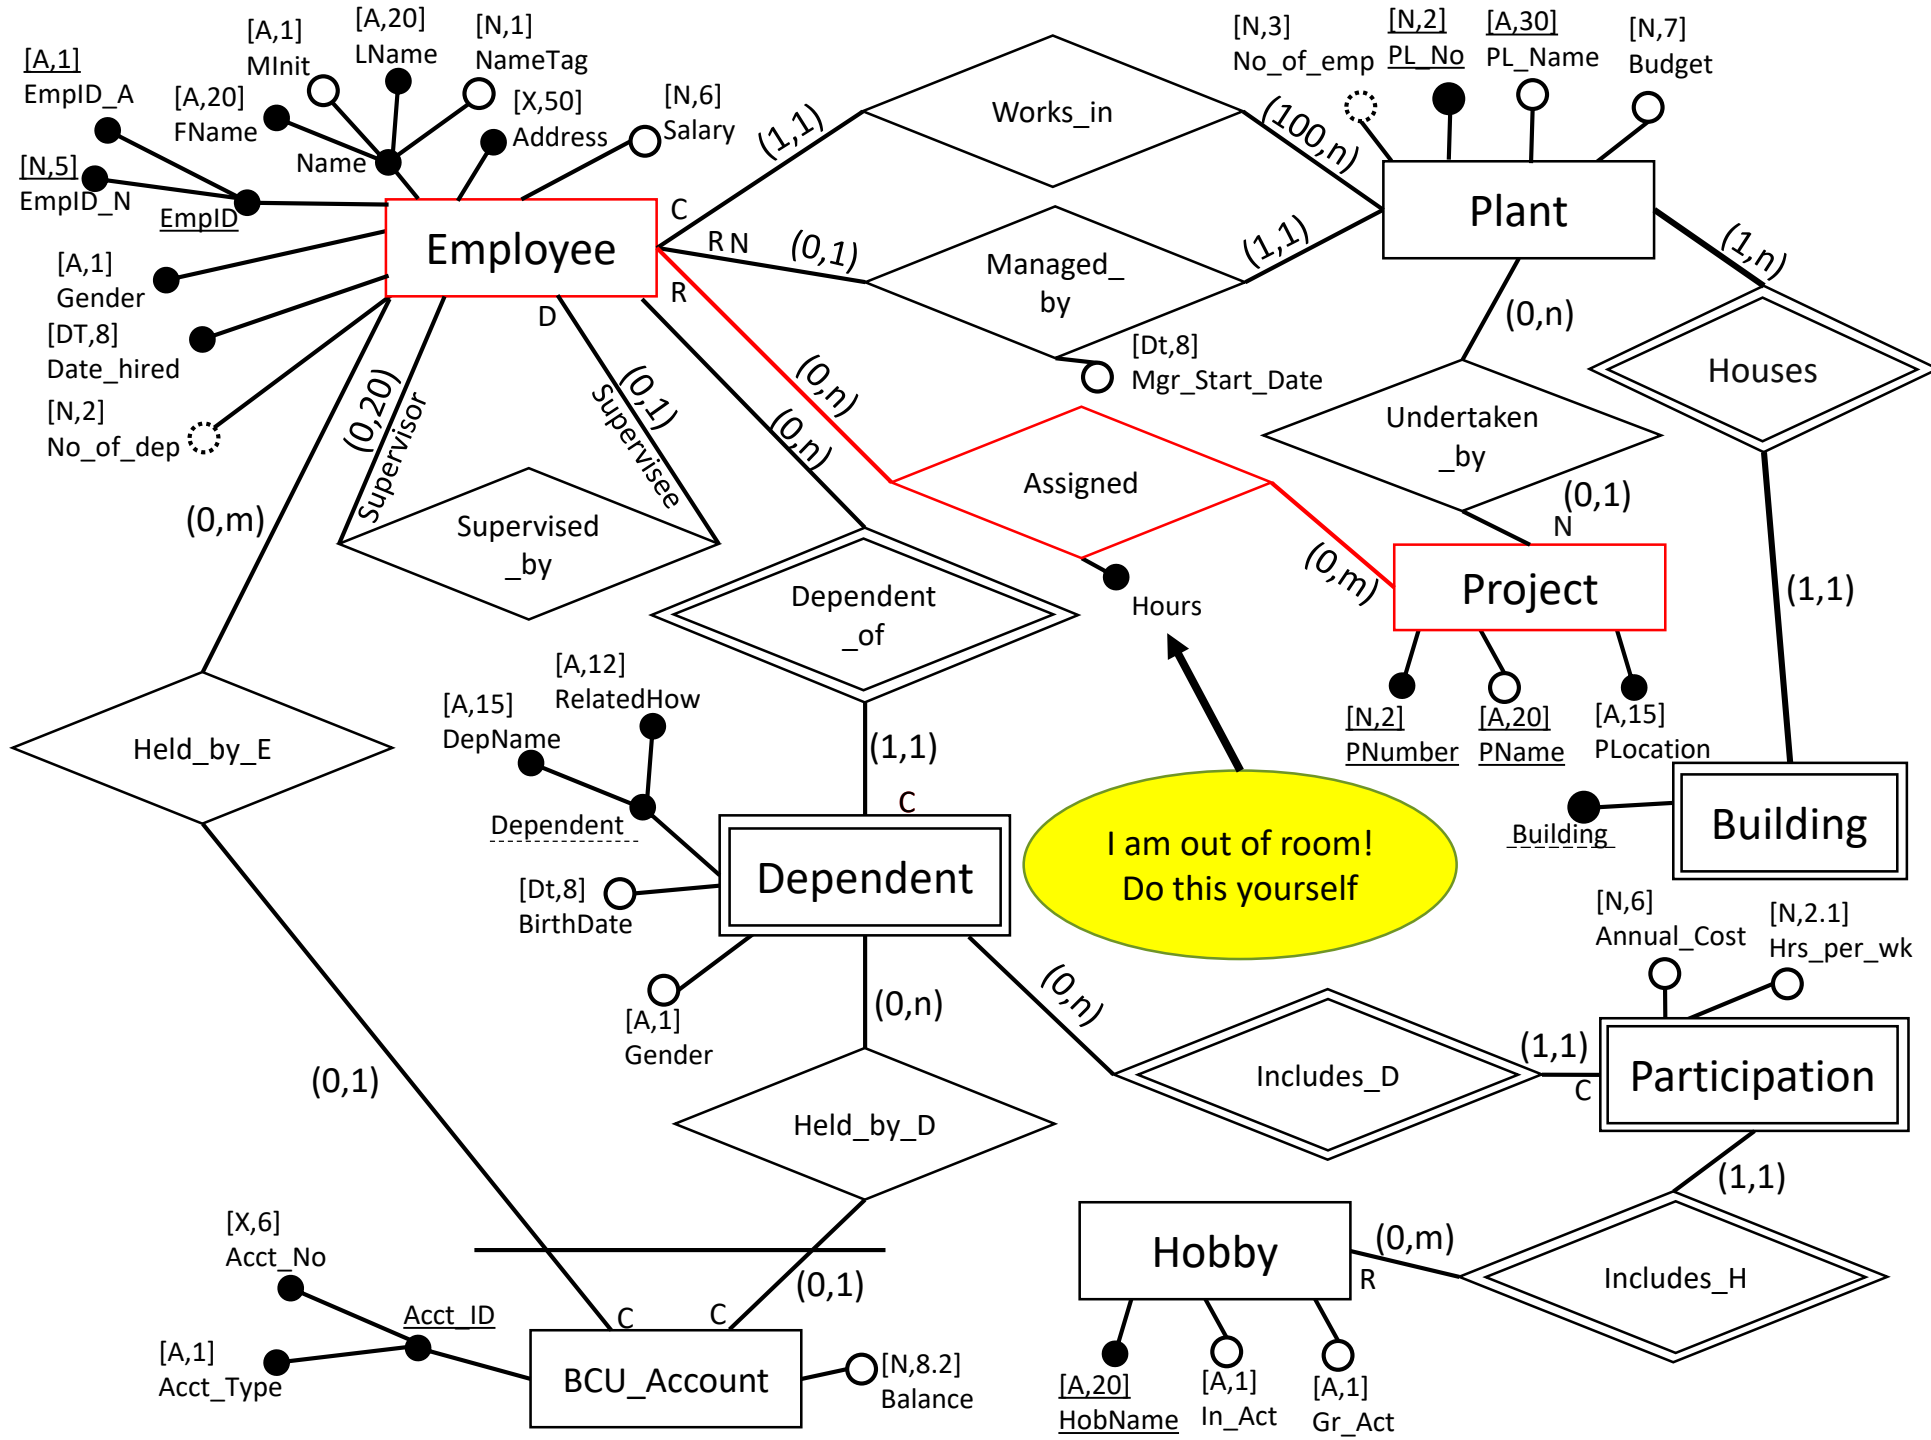
- Gerunds are a product of decomposing m:n relationships
 - Also called a “composite entity” or “bridge entity”
 - Takes the primary key from each participating entity to create a set of 1:n and m:1 relationships



The Gerund




- Looks and acts like a weak entity, but not the same
 - **Weak entities** are a product of the **business rules**
 - **Gerunds** are a product of **decomposing m:n relationships**
 - Gerunds do not have partial keys






Recall when we talked about this:

- Attributes may not only be part of an entity, but may also be part of the relationship
- Rent on a dorm room is the same no matter who is in it
- A student pays the same rent no matter what dorm room they have
- Rent varies based on student and room type



The gerund is also how we capture attributes of a relationship


Another case of attributes in a relationship



Note: Since the cardinality constraint of *Imports* is m:n, *Cost* cannot be an attribute of either VENDOR or PRODUCT – *Cost must remain as an attribute of Imports*


VENDOR		IMPORTS			PRODUCT	
<u>Vendor_id</u>	Vendor_name	<u>Vendor_id</u>	<u>Product_id</u>	Cost	Product_name	<u>Product_id</u>
V3	Buffet Inc.	V3	P11	7	Soup	P11
V5	Gates Inc.	V5	P11	10	Noodles	P13
V7	Jobs Inc.	V5	P17	17	Chocolate	P17
		V7	P17	19	Nuts	P19
		V7	P23	13	Coffee	P23

What if we moved cost to product?



VENDOR		IMPORTS		PRODUCT		
<u>Vendor_id</u>	Vendor_name	<u>Vendor_id</u>	<u>Product_id</u>	Product_name	<u>Product_id</u>	Cost
V3	Buffet Inc.	V3	P11	Soup	P11	7
V5	Gates Inc.	V5	P11	Noodles	P13	7
V7	Jobs Inc.	V5	P17	Chocolate	P17	17
		V7	P17	Nuts	P19	19
		V7	P23	Coffee	P23	13

What if we moved cost to product?



- Now products must cost the same regardless of the vendor – does this reflect reality?





VENDOR		IMPORTS		PRODUCT		
<u>Vendor_id</u>	Vendor_name	<u>Vendor_id</u>	<u>Product_id</u>	Product_name	<u>Product_id</u>	Cost
V3	Buffet Inc.	V3	P11	Soup	P11	7
V5	Gates Inc.	V5	P11	Noodles	P13	7
V7	Jobs Inc.	V5	P17	Chocolate	P17	17
		V7	P17	Nuts	P19	19
		V7	P23	Coffee	P23	13

Cost as part of product

- A box of girl scout cookies costs the same \$6 regardless of the vender (troop) you buy from
- In commodities markets, a barrel of oil or bushel of grain costs the same regardless of who you buy from



What if we moved cost to vendor?



VENDOR		
<u>Vendor_id</u>	Vendor_name	Cost
V3	Buffet Inc.	7
V5	Gates Inc.	10
V7	Jobs Inc.	19


IMPORTS	
<u>Vendor_id</u>	<u>Product_id</u>
V3	P11
V5	P11
V5	P17
V7	P17
V7	P23

PRODUCT	
Product_name	<u>Product_id</u>
Soup	P11
Noodles	P13
Chocolate	P17
Nuts	P19
Coffee	P23


Relationships are indicated by lines connecting the tables:

- V3 connects to V3 in IMPORTS.
- V5 connects to V5 in IMPORTS.
- V7 connects to V7 in IMPORTS.
- P11 connects to P11 in PRODUCT.
- P17 connects to P17 in PRODUCT.
- P23 connects to P23 in PRODUCT.

What if we moved cost to vendor?



- Now all products from a vendor cost the same regardless of the product – does this reflect reality?



VENDOR		
<u>Vendor_id</u>	Vendor_name	Cost
V3	Buffet Inc.	7
V5	Gates Inc.	10
V7	Jobs Inc.	19

IMPORTS	
<u>Vendor_id</u>	<u>Product_id</u>
V3	P11
V5	P11
V5	P17
V7	P17
V7	P23

PRODUCT	
Product_name	<u>Product_id</u>
Soup	P11
Noodles	P13
Chocolate	P17
Nuts	P19
Coffee	P23


Cost as part of vendor

- At a farmers market, it might be the case that all products from a vendor cost the same, regardless of what the product is
- From a vending machine, all items might cost the same, regardless of what the product is



The most normal case though...

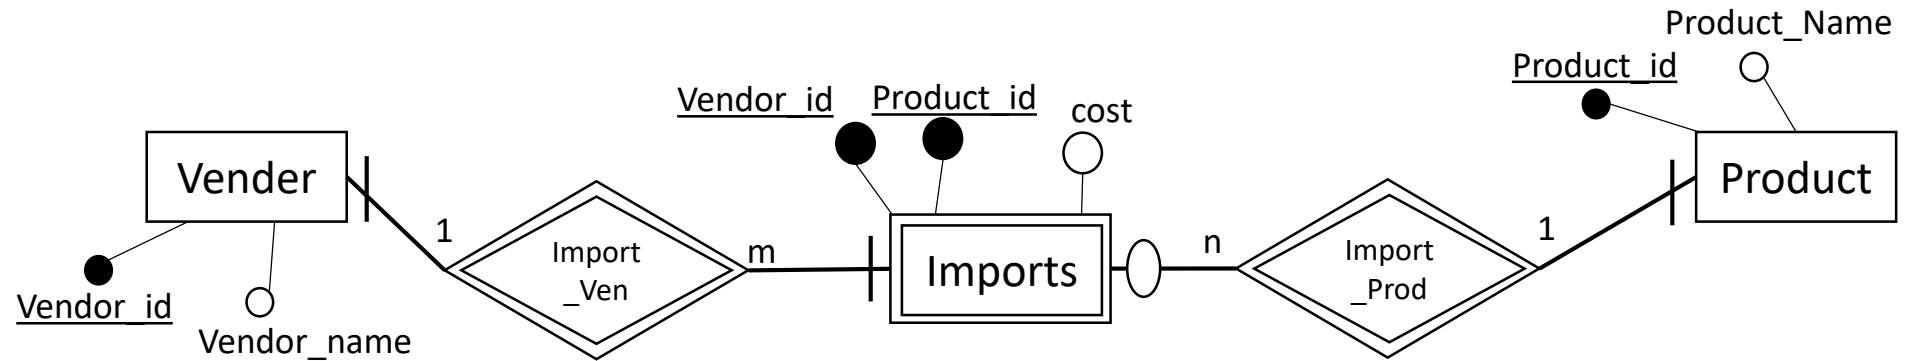
- Cost depends on both the vendor and the product, therefore



Note: Since the cardinality constraint of *Imports* is m:n, *Cost* cannot be an attribute of either *VENDOR* or *PRODUCT* – *Cost must remain as an attribute of Imports*

VENDOR		IMPORTS			PRODUCT	
<u>Vendor_id</u>	Vendor_name	<u>Vendor_id</u>	<u>Product_id</u>	Cost	Product_name	<u>Product_id</u>
V3	Buffet Inc.	V3	P11	7	Soup	P11
V5	Gates Inc.	V5	P11	10	Noodles	P13
V7	Jobs Inc.	V5	P17	17	Chocolate	P17
		V7	P17	19	Nuts	P19
		V7	P23	13	Coffee	P23

The decomposed model of this



Note: Since the cardinality constraint of *Imports* is m:n, Cost cannot be an attribute of either VENDOR or PRODUCT – Cost must remain as an attribute of Imports

VENDOR		IMPORTS			PRODUCT	
<u>Vendor_id</u>	Vendor_name	<u>Vendor_id</u>	<u>Product_id</u>	Cost	Product_name	<u>Product_id</u>
V3	Buffet Inc.	V3	P11	7	Soup	P11
V5	Gates Inc.	V5	P11	10	Noodles	P13
V7	Jobs Inc.	V5	P17	17	Chocolate	P17
		V7	P17	19	Nuts	P19
		V7	P23	13	Coffee	P23

Module 3.2.4

Decomposed Design-Specific ERD

- Attribute Placement
- Weak entities
- How do we handle multi-value attributes?
- How do we handle m:n relationships?

Module 3.3

Modeling Errors

- What are semantic errors?
- What are syntactical errors?

Modeling errors

- Errors may be in:
 - Syntax: using the symbols incorrectly
 - Relatively easy to spot
 - Semantics: not accurately reflecting the business rules
 - More difficult to spot, because these are often judgment calls
- How do we know what is an entity vs. an attribute vs. a value?
 - Often trial and error + experience
- It can be difficult to know exactly how something should be modeled until you start modeling it and fail
 - Was this anyone's assignment 1 experience?

Consider this story:

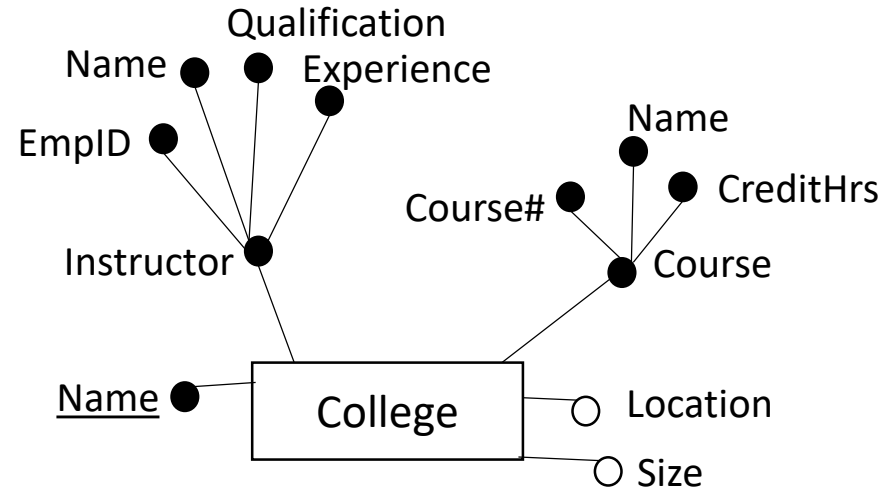
- There are several colleges in the university. Each college has a **Name**, **Location**, and **Size**. A college offers many **courses**, which are described by Course#, Name, and Credit Hours. The college also has several **instructors**. Instructors teach courses, but not all instructors are scheduled to teach during all terms. All courses must be taught by an instructor. Instructors are capable of teaching a multiple courses offered by the college. Instructors have a unique Employee ID and their Name, Qualification, and Experience are also recorded.
- Note: This is a simplified version of the vignette on page 120.

Working through the two vignettes presented on pages 120-133 is a good exercise/exam practice

Consider this story:

- College

- Name
- Location
- Size
- Courses
 - Course#
 - Name
 - CreditHrs
- Instructor
 - Name
 - EmpID
 - Qualification
 - Experience



This is syntactically OK

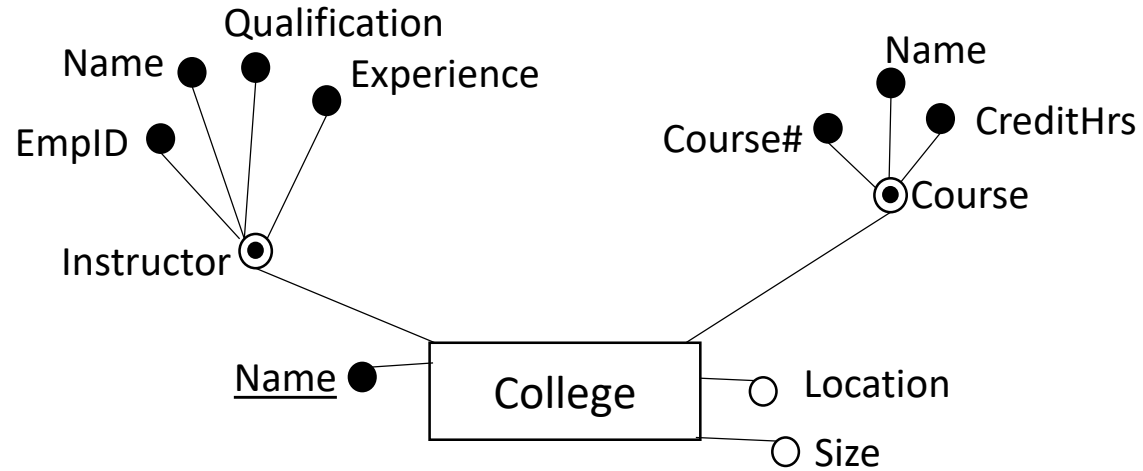
(But incorrect)

Semantic Problem: college can have Multiple Instructors and Courses

Consider this story:

- College

- Name
- Location
- Size
- Courses
 - Course#
 - Name
 - CreditHrs
- Instructor
 - Name
 - EmpID
 - Qualification
 - Experience



This is syntactically OK

(But incorrect)

Semantic Problem: Instructors teach courses

Consider this story:

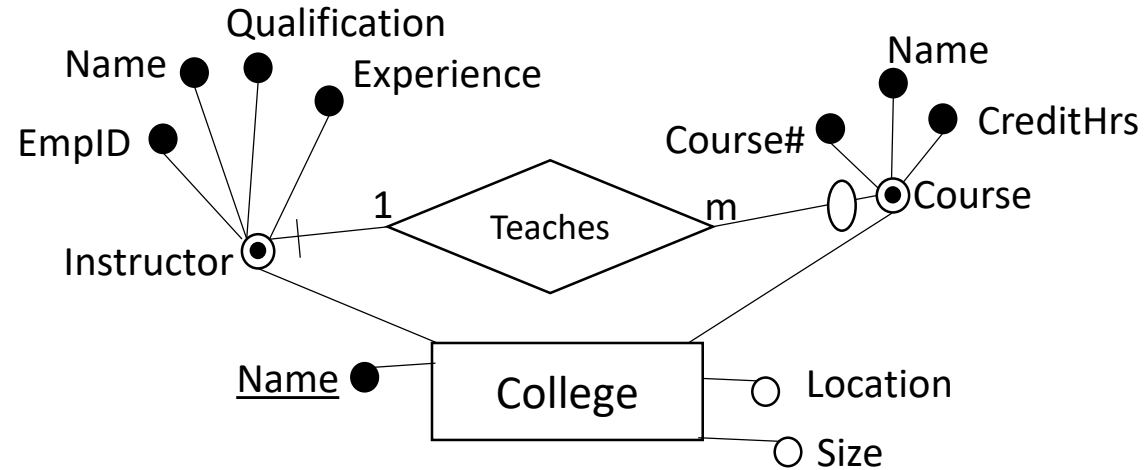
- College

- Name
- Location
- Size
- Courses

- Course#
- Name
- CreditHrs

- Instructor

- Name
- EmpID
- Qualification
- Experience



This is syntactically Incorrect

How do we properly model multi-value attributes?

Semantic Problems have been resolved

Consider this story:

- College

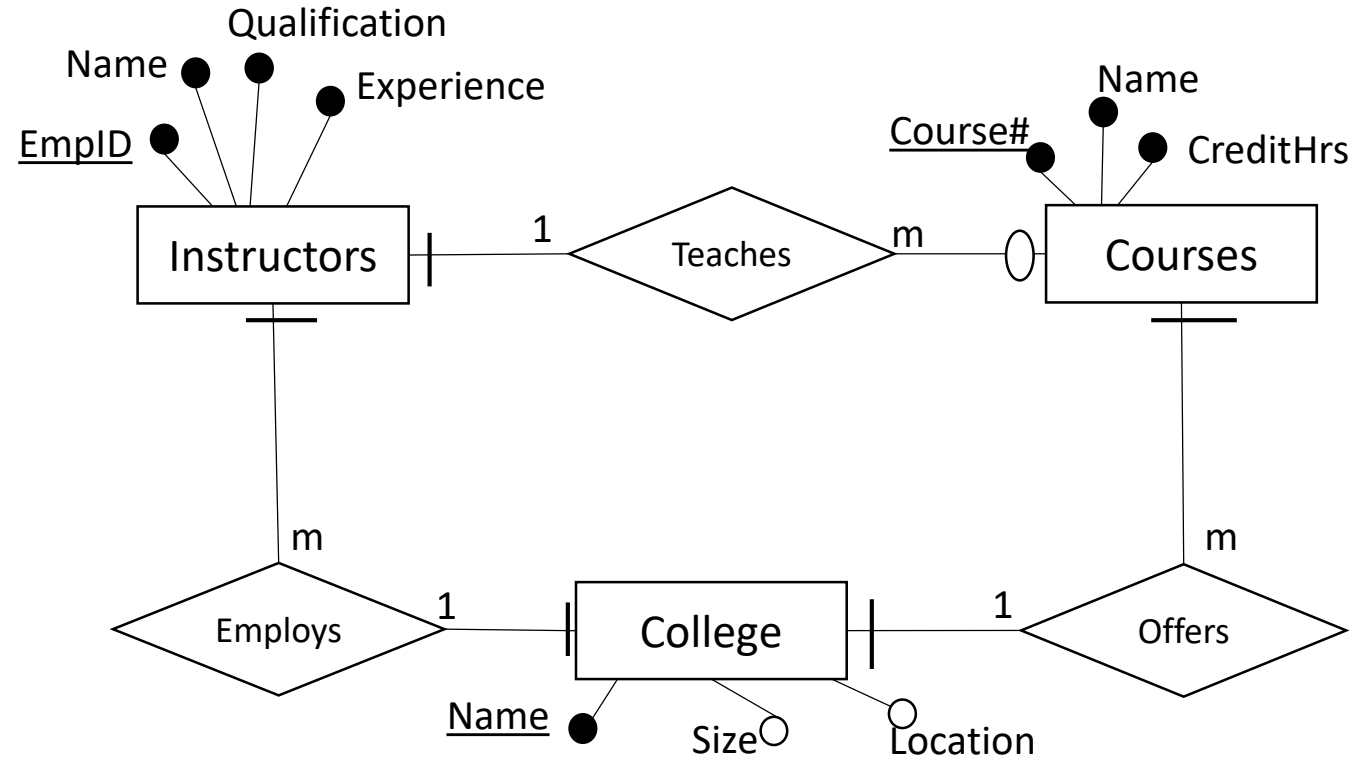
- Name
- Location
- Size

- Courses

- Course#
- Name
- CreditHrs

- Instructor


- Name
- EmpID
- Qualification
- Experience



**This is syntactically correct and
semantic Problems have been resolved**

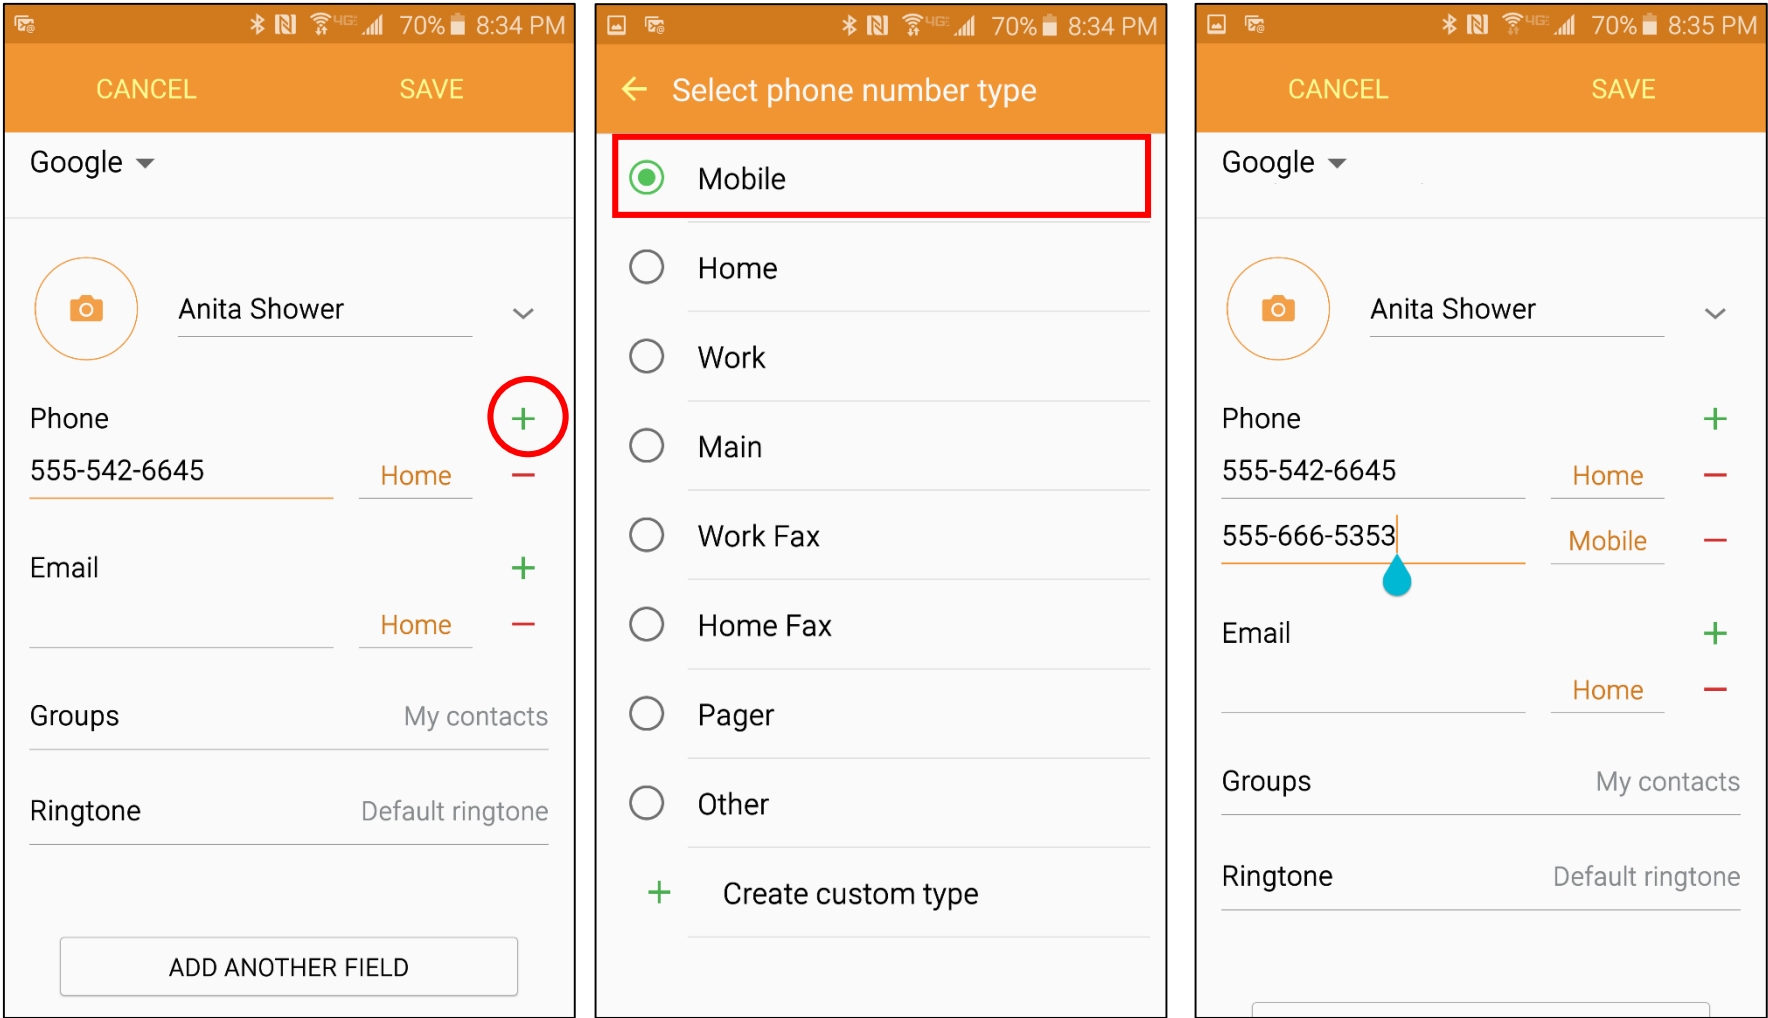
Consider this story: Contacts list on my phone

A screenshot of an Android contact editing interface. At the top, there's an orange header bar with 'CANCEL' and 'SAVE' buttons. Below it, a dropdown menu shows 'Google'. The contact's name 'Anita Shower' is displayed with a camera icon to its left and a dropdown arrow to its right. Below the name, there are fields for 'Phone' (555-542-6645) and 'Email'. Each field has a green '+' icon to its right and a 'Home' label with a red '-' icon below it. At the bottom, there are sections for 'Groups' (My contacts) and 'Ringtone' (Default ringtone). A button labeled 'ADD ANOTHER FIELD' is at the very bottom.

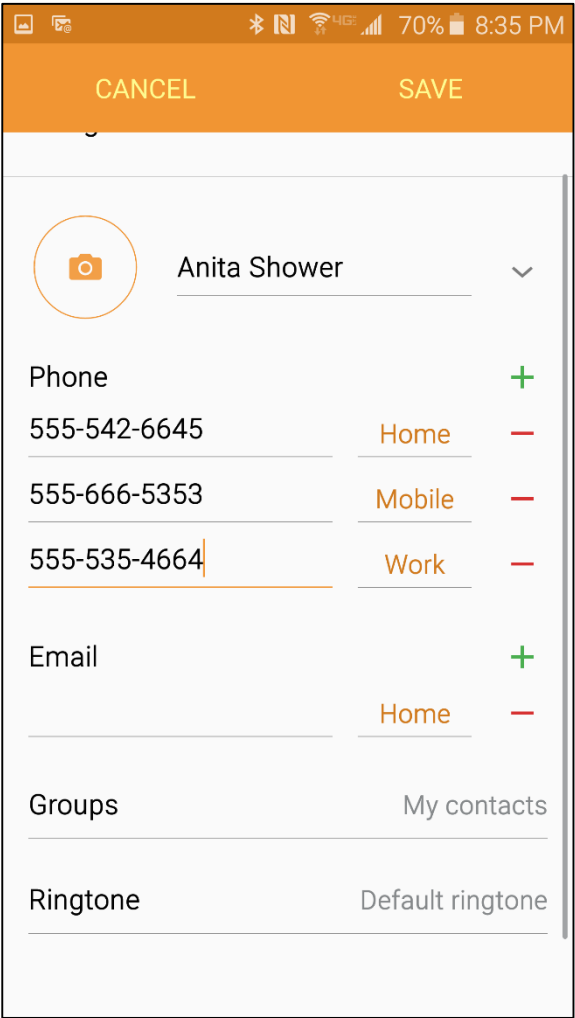
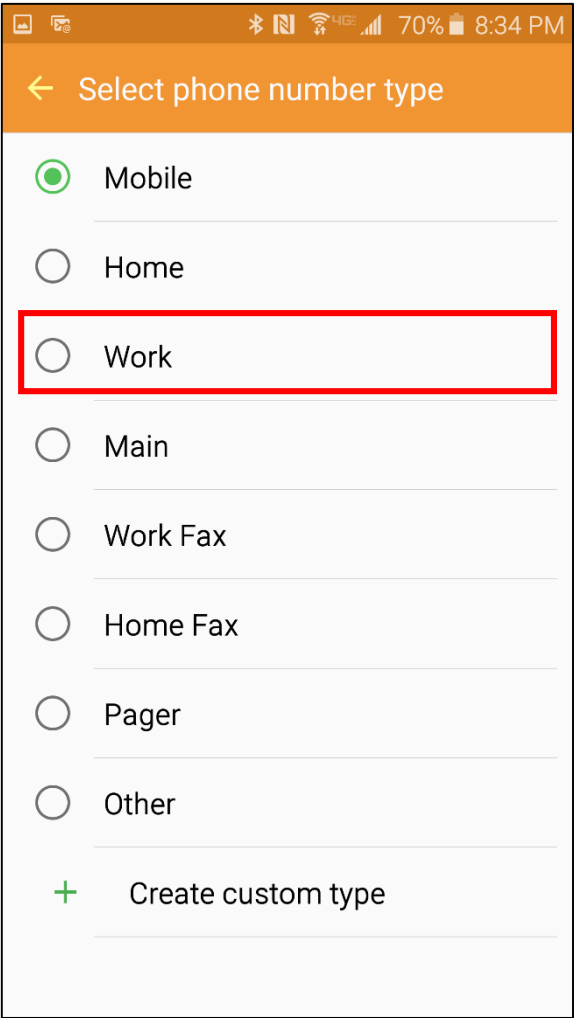



This seems reasonable and is syntactically correct, but does it capture all the business rules?

Consider this story: Contacts list on my phone

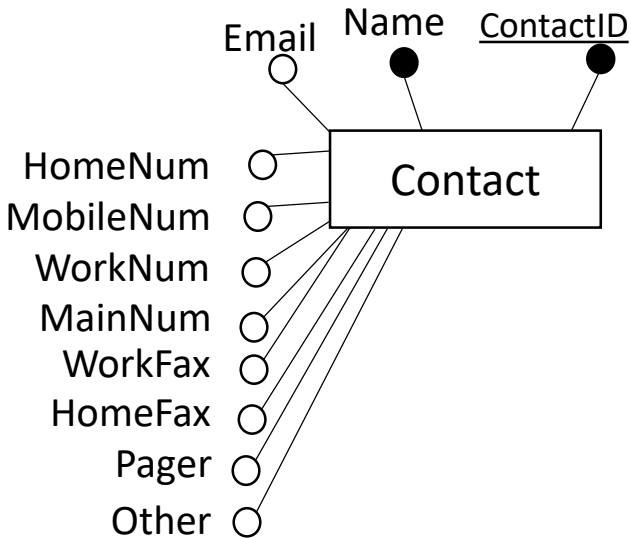


Contacts list on my phone



Perhaps this is the data structure...

<u>Name</u>	HomeNum	MobileNum	WorkNum	MainNum	WorkFax	HomeFax	Pager	Other
Anita Shower	555-123-2134							
Jane Smith	555-234-4344	555-234-3434	555-334-3434					
Tom Powers		555-345-3523						
Dave Rogers	555-234-3995		555-334-3434					



This seems reasonable and is syntactically correct, but does it capture all the business rules?



A screenshot of a mobile application interface. At the top, there is a status bar with icons for signal, Wi-Fi, and battery (70%), and the time 8:34 PM. Below the status bar is an orange header bar with a back arrow and the text "Select phone number type". The main area of the screen is white and contains a list of radio button options: "Mobile" (selected), "Home", "Work", "Main", "Work Fax", "Home Fax", "Pager", and "Other". At the bottom of the list is a green plus icon followed by the text "Create custom type".

I can create any type of number

70%8:34 PM

←

Select phone number type

Mobile

Home

Work

Main

Work Fax

Home Fax

Pager

Other

+

Create custom type

70%8:36 PM

CANCEL

SAVE

Google ▾

Anita Shower ▾

Phone

555-542-6645

Home

+

−

555-666-5353

Mobile

−

555-535-4664

Work

−

555-454-6456

Helicop...

−

Email

Home

+

−

Groups

My contacts

Ringtone

Default ringtone

70%8:36 PM

CANCEL

SAVE

Google ▾

Anita Shower ▾

Phone

555-542-6645

Home

+

−

555-666-5353

Mobile

−

555-535-4664

Work

−

555-454-6456

Helicop...

−

555-646-3346

Submar...

−

Email

Home

+

−

Groups


My contacts

How should this be modeled?

- Does this data model still make sense?


<u>Name</u>	HomeNum	MobileNum	WorkNum	MainNum	WorkFax	HomeFax	Pager	Other
Anita Shower	555-123-2134							
Jane Smith	555-234-4344	555-234-3434	555-334-3434					
Tom Powers		555-345-3523						
Dave Rogers	555-234-3995		555-334-3434					

- Perhaps we do not have multiple phone number attributes, but rather phone number is...a multi-value attribute!



...a multi-value attribute

- What type of entity is “PhoneNumbers”?
- What is “Number”?



Contacts		PhoneNumbers		
<u>ContactID</u>	Name	<u>ContactID</u>	<u>Number</u>	NumberType
101	Anita Shower	101	555-123-2134	Home
101	Jane Smith	101	555-343-5235	Mobile
103	Tom Powers	101	555-266-6463	Helicopter
104	Dave Rogers	101	555-234-2133	Submarine
		102	555-234-4344	Home
		102	555-234-5522	Mobile
		102	555-334-3434	Work
		103	555-345-3523	Mobile
		104	555-234-3995	Home
		104	555-334-3434	Work

Assignment 1

- This was one of the goals of assignment 1 – to look at a system and infer what the schema might look like.
- As an outsider looking in, you will likely be wrong. Why?
 - You only see the external schema – many of the details and business rules are hidden from you
- It can sometimes be difficult to determine if something should be an entity, attribute, or value
 - An entity is a group of related attributes
 - Imagine what an instance of the entity would be
 - Try to create values for each attribute for each instance you think of

Module 3.3

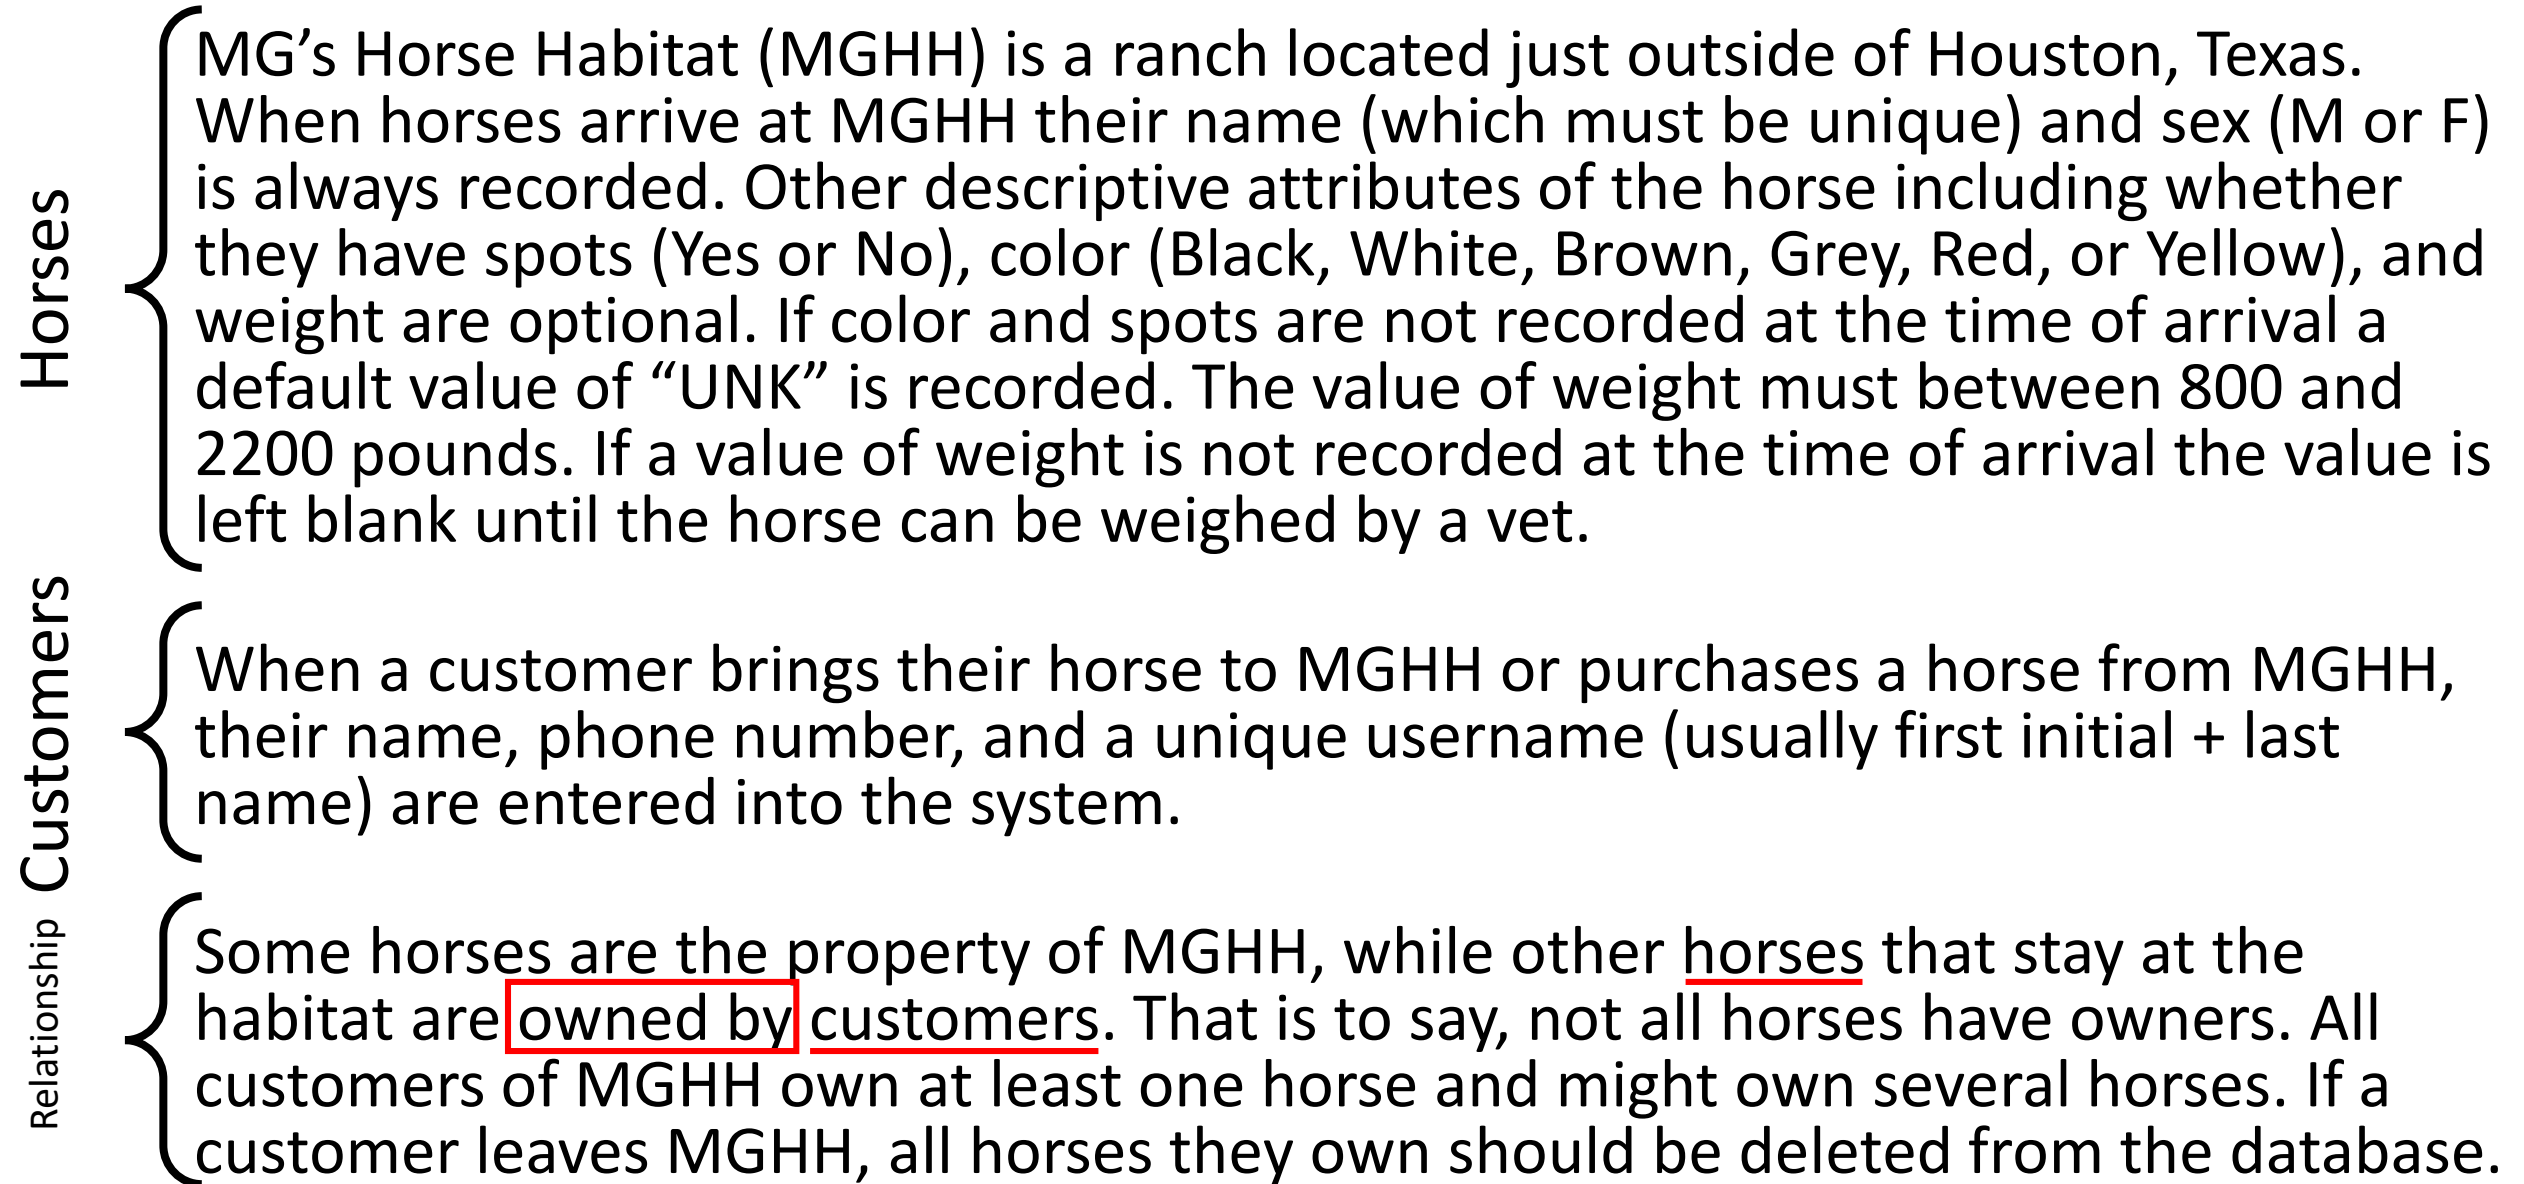
Modeling Errors

- What are semantic errors?
- What are syntactical errors?


Let's apply what we have learned to MGHH!



MG's Horse Habitat – Horses and Customers



MG's Horse Habitat – Horses and Customers



Semantic Integrity Constraints:

Color: {Black, White, Brown, Grey, Red, Yellow}, Default: 'UNK'

Spots: {Yes, No}, Default: 'UNK'

Sex: {M, F}

Weight: {800-2200}

Let's have a fresh start by recreating all we've done so far...

- The next few slides have the EXACT same SQL we've done so far, but in case you want everything in one easy to find spot, we can start over from here!
- Drop our existing tables:

```
DROP TABLE horses;  
DROP TABLE customers;
```

Recreate the Customers and Horses tables with all constraints

```
CREATE TABLE customers
```

```
(username varchar(50) CONSTRAINT pk_customers PRIMARY KEY,  
  Fname      varchar(50) CONSTRAINT nn_fname NOT NULL,  
  Lname      varchar(50) CONSTRAINT nn_lname NOT NULL,  
  Phone      varchar(14) CONSTRAINT nn_Phone NOT NULL  
);
```

```
CREATE TABLE horses
```

```
(Name      varchar(50) CONSTRAINT pk_horse PRIMARY KEY,  
  Color     varchar(50) DEFAULT 'UNK' CONSTRAINT chk_color CHECK (color IN  
    ('Black','White','Brown','Grey','Red','Yellow', 'UNK')),  
  Spots     varchar(3) DEFAULT 'UNK',  
  Sex       varchar(1) CONSTRAINT nn_sex NOT NULL,  
  Weight    integer,  
  owner     varchar(50),  
  CONSTRAINT chk_weight CHECK (weight >= 800 AND weight <=2200),  
  CONSTRAINT chk_sex CHECK (sex IN ('M','F')),  
  CONSTRAINT fk_cust FOREIGN KEY (owner) REFERENCES customers (username) ON DELETE CASCADE  
);
```

INSERTing all the data we've worked with so far:

```
INSERT INTO customers (username, fname, lname, phone) VALUES ('mgrimes', 'Marvin', 'Grimes', '(218) 330-8004');
INSERT INTO customers (username, fname, lname, phone) VALUES ('canderson', 'Christine', 'Anderson', '(555) 523-9989');
INSERT INTO customers (username, fname, lname, phone) VALUES ('tswift', 'Tina', 'Swift', '(555) 424-1313');
INSERT INTO customers (username, fname, lname, phone) VALUES ('jisbell', 'Jason', 'Isbell', '(615) 555-5555');
INSERT INTO customers (username, fname, lname, phone) VALUES ('ssimpson', 'Sam', 'Simpson', '(615) 387-9682');
```

```
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('Sam', 'Brown', 'No', 'F', 1500, 'mgrimes');
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('Erica', 'Yellow', 'Yes', 'F', 920, 'canderson');
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('John', 'Grey', 'No', 'M', 1800, 'mgrimes');
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('Trotty', 'Brown', 'Yes', 'M', 1300, 'mgrimes');
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('Rio', 'Grey', 'No', 'F', 1700, 'tswift');
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('Robin', 'Yellow', 'No', 'M', 1100, 'jisbell');
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('Katy', 'Brown', 'No', 'F', 1200, 'jisbell');
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('Pegasus', 'Brown', 'No', 'M', 1750, 'mgrimes');
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('Sammy', 'Black', 'Yes', 'M', 2200, 'mgrimes');
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('Pinky', 'Red', 'No', 'M', 1050, 'tswift');
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('Hulk', 'Grey', 'No', 'M', 2050, 'mgrimes');
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('Pat', 'White', 'No', 'F', 1400, 'mgrimes');
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('Betty', 'White', 'Yes', 'F', 1250, 'tswift');
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('Shamrock', 'Black', 'No', 'M', 1400, 'ssimpson');
```

Now on to the new stuff!

MG's Horse Habitat – Horses and Medications

While at MGHH, medications may be ordered for horses to treat a variety of illnesses, or to ensure their continued health.

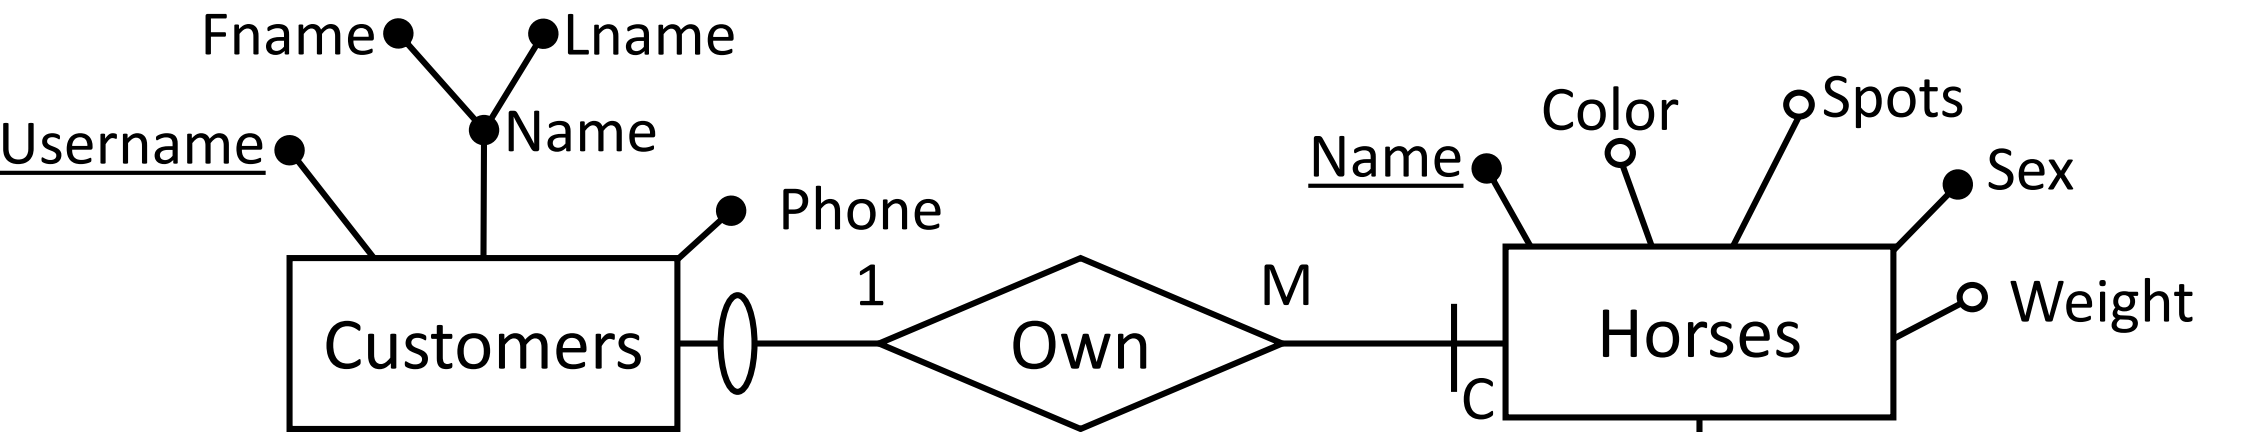
Medications are identified by a medication code (MedCode) but also have a commonly used name (for example Ivermectin or Ibuprofen). Each medication has a cost per dose that ranges from \$1 to \$200 and a classification that describes its use (NSAID, Antibiotic, Sedative, etc...).

MGHH keeps track of quantity on hand (QoH) as well as quantity on order (QoO). The value of QoH and QoO together must be between 10 and 100—that is, if QoH gets down to 10, then some units should be ordered to replenish the stock, but not to a quantity greater than 100.


When medications are ordered for horses, a dose and frequency should be recorded. If dose and frequency are not provided, a default dose of “1” should be used, with a frequency of “2” times per day. Orders are uniquely identified by a prescription number (RxNum).

If a horse is deleted from the database, all records of their medication orders should be deleted. A medication that has been ordered for a horse cannot be deleted from the database.

MG's Horse Habitat – Horses and Medications

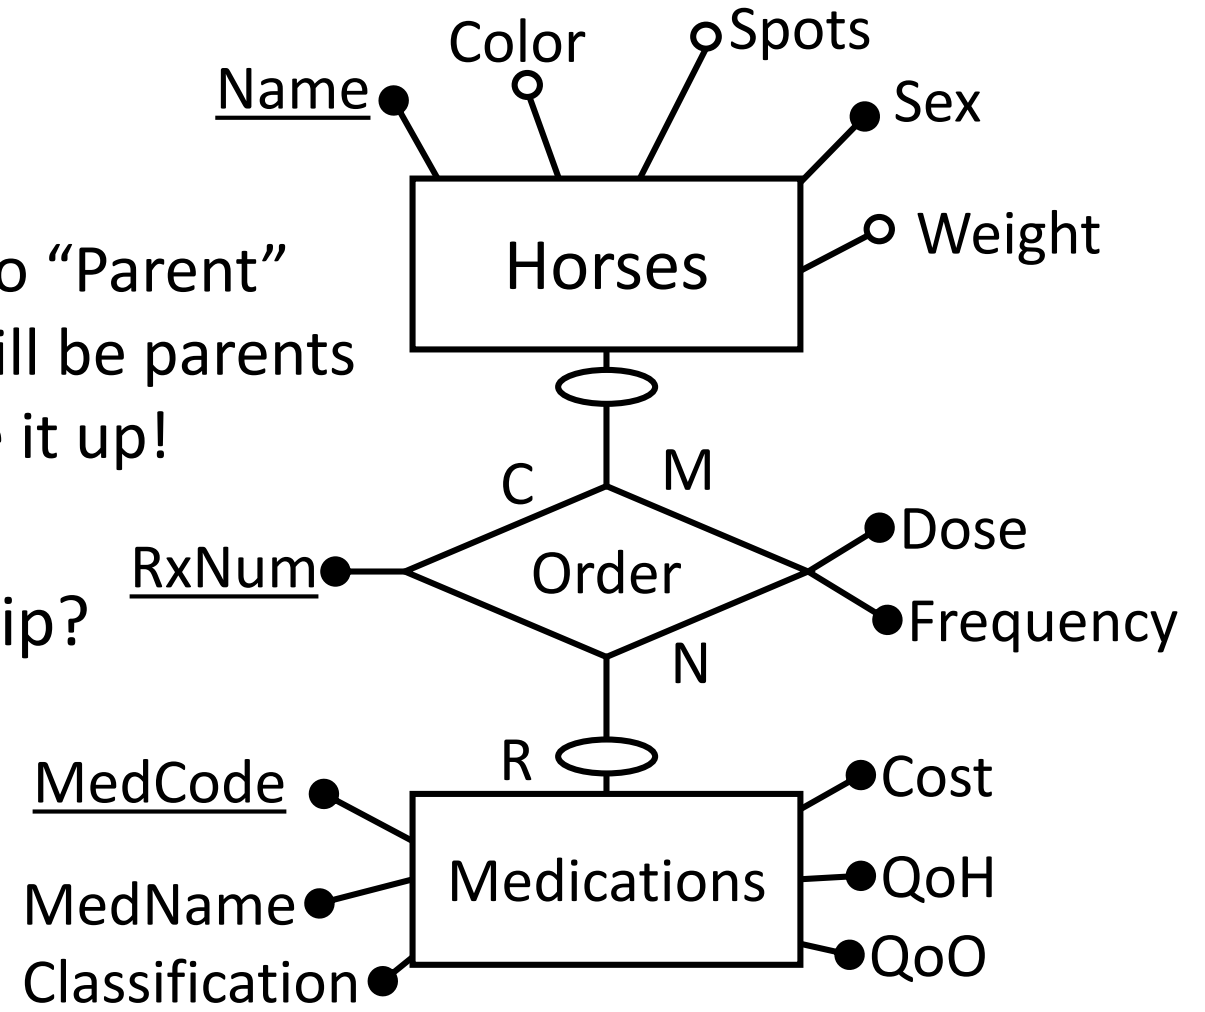


Semantic Integrity Constraints:
Color: {Black, White, Brown, Grey, Red, Yellow}, Default: 'UNK'
Spots: {Yes, No}, Default: 'UNK'
Sex: {M, F}
Weight: {800-2200}
Cost: {1-200}
QoH + QoO: {10-100}
Dose: Default: 1
Frequently: Default 2




MG's Horse Habitat – Horses and Medications


- Which entity is the “Parent” in this relationship?
 - Either One?
 - Neither One?
 - Both?
 - All of the above are correct – there is no “Parent” in a M:N relationship, because BOTH will be parents in the two 1:M Relationships that make it up!
- How do we model this M:N Relationship?
 - Using a Gerund to decompose into two 1:M relationships



MG's Horse Habitat – Horses and Medications



- The M:N relationship becomes two 1:M relationships!
- Horses is the parent in “Ordered For”
- Medications is the parent in “Contains”
- **Name** and **MedCode** will individually be Foreign Keys in the Order table, and together will be a composite primary key

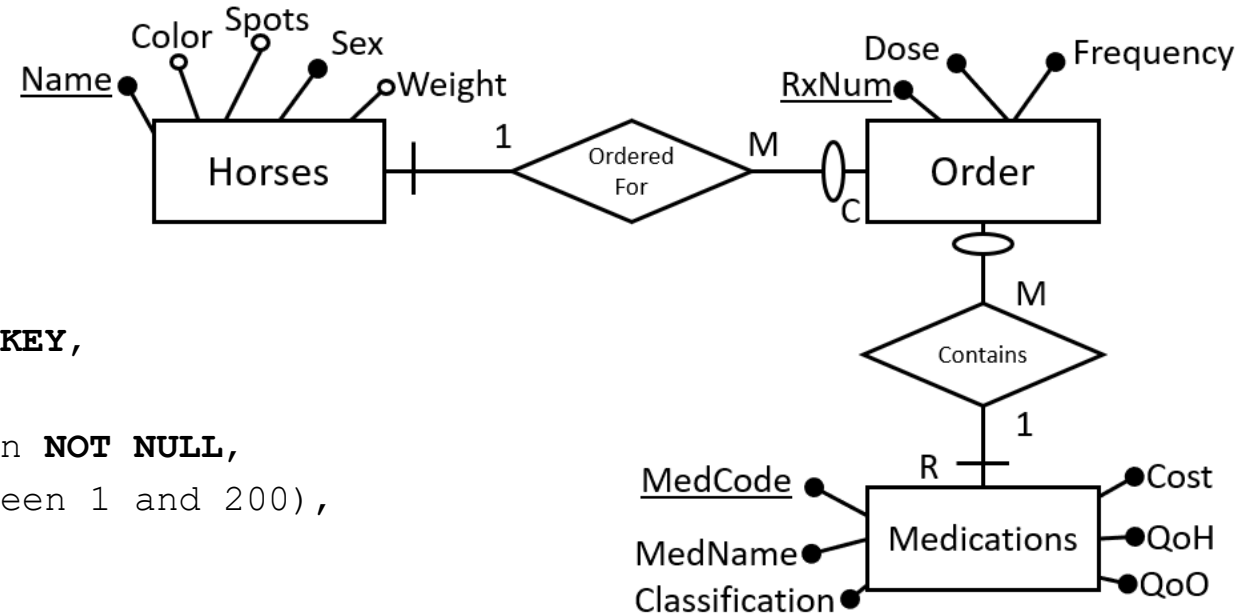


MG's Horse Habitat – Horses and Medications

- First we need to CREATE the Medications table:

```
CREATE TABLE medications
(
  medcode varchar(50) CONSTRAINT pk_medications PRIMARY KEY,
  medname varchar(50) CONSTRAINT nn_medname NOT NULL,
  classification varchar(50) CONSTRAINT nn_classification NOT NULL,
  cost float(32) CONSTRAINT chk_cost CHECK (cost between 1 and 200),
  QoH integer,
  QoO integer,
  CONSTRAINT chk_qty CHECK ((QoH + QoO) BETWEEN 10 AND 100)
);
```

- Note that we must create the Medications table before the Order table, since Order will reference Medications



MG's Horse Habitat – Horses and Medications

- Let's INSERT some horse medications:

```
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Flux','Flunixin Meglumine','NSAID',27,43,0);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Bute','Phenylbutazone','NSAID',19,84,0);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Oxi','Oxibuzone','NSAID',12,55,0);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Mel','Meloxicam','NSAID',6,72,0);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Pen','Penicillin','Antibiotic',38,73,0);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Doxy','Doxycycline','Antibiotic',81,89,0);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('TMPS','Trimethoprim Sulfa','Antibiotic',27,50,0);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Enro','Enrofloxacin','Antibiotic',36,78,0);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Xyl','Xylazine ','Sedative',22,28,0);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Acep','Acepromazine','Sedative',33,50,0);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Fen','Fenbendazole','Dewormer',52,15,25);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Meb','Mebendazole','Dewormer',81,25,0);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Pyr','Pyrantel Embonate','Dewormer',11,29,0);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Iver','Ivermectin','Dewormer',39,21,0);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Clen','Clenbuterol','Respiratory',132,11,20);
```

MG's Horse Habitat – Horses and Medications

- Note that we cannot insert a combined value greater than 100

```
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Pres','Prednisone','Corticosteroid',47,48,70);
```



SQL Error [2290] [23000]: ORA-02290:
check constraint
(GMGRIMES.CHK_QTY) violated

- ...But we can if we fix this error by reducing the Quantity on Order (QoO):

```
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Pres','Prednisone','Corticosteroid',47,48,7);
```

- We also cannot create a medication with a cost greater than 200

```
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Dexa','Dexamethasone ','Corticosteroid',247,18,82);
```



SQL Error [2290] [23000]: ORA-02290:
check constraint
(GMGRIMES.CHK_COST) violated

- Fixed:

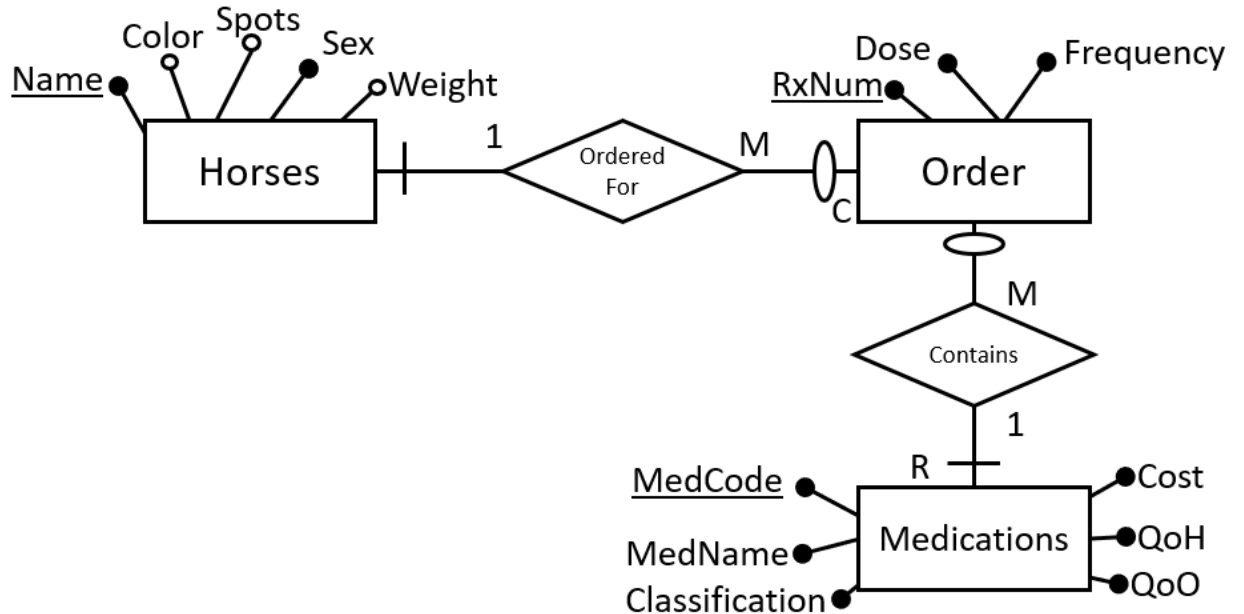
```
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Dexa','Dexamethasone ','Corticosteroid',47,18,82);
```

MG's Horse Habitat – Horses and Medications

- Now let's CREATE the Order table

For your notes – there is an error in this DDL

```
CREATE TABLE order
(RxNum integer CONSTRAINT pk_order PRIMARY KEY,
ORD_Medcode varchar (50),
ORD_Horsename varchar(50),
Dose integer DEFAULT 1,
Frequency integer DEFAULT 2,
CONSTRAINT fk_meds FOREIGN KEY (ORD_Medcode) REFERENCES medications (medcode),
CONSTRAINT fk_horses FOREIGN KEY (ORD_Horsename) REFERENCES horses (name) ON DELETE CASCADE
);
```

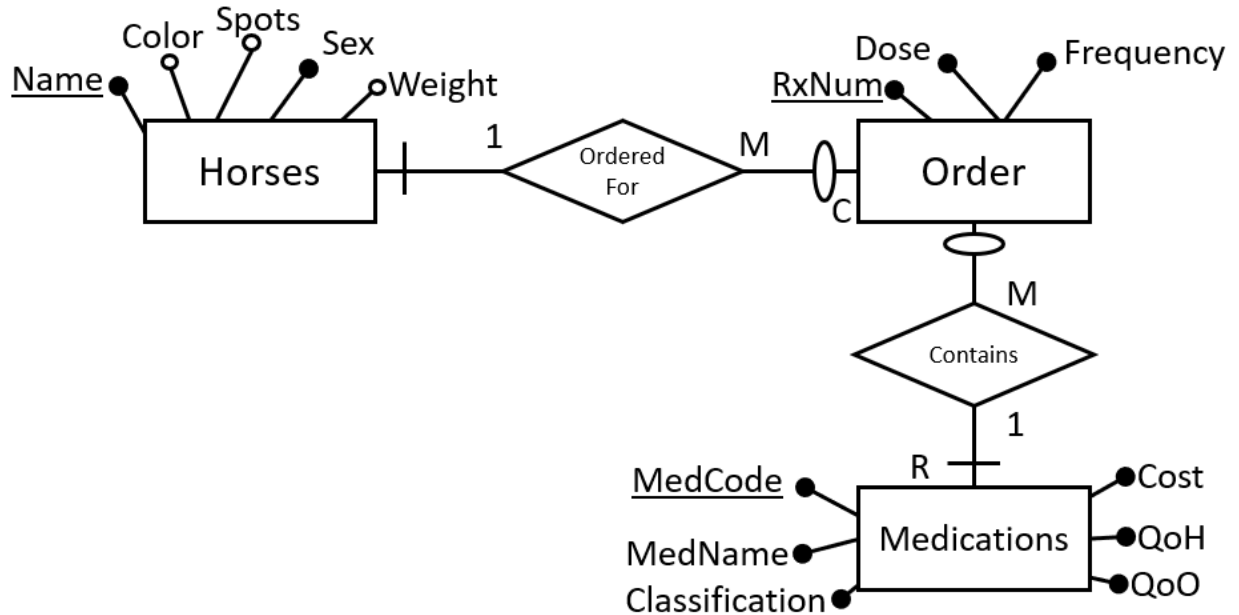


SQL Error [903] [42000]: ORA-00903:
invalid table name

MG's Horse Habitat – Horses and Medications

- You cannot call the table “order” because “order” is a reserved keyword in SQL
 - We can just rename this table to be “orders” instead:

```
CREATE TABLE orders
(RxNum integer CONSTRAINT pk_orders PRIMARY KEY,
ORD_Medcode varchar (50),
ORD_Horsename varchar(50),
Dose integer DEFAULT 1,
Frequency integer DEFAULT 2,
CONSTRAINT fk_meds FOREIGN KEY (ORD_Medcode) REFERENCES medications (medcode),
CONSTRAINT fk_horses FOREIGN KEY (ORD_Horsename) REFERENCES horses (name) ON DELETE CASCADE
);
```



MG's Horse Habitat – Horses and Medications

- Let's INSERT some orders for medications:

```
INSERT INTO ORDERS (RxNum, ord_medcode, ord_horsename, dose, frequency) values (101, 'Flux', 'Sam', 1, 2);
INSERT INTO ORDERS (RxNum, ord_medcode, ord_horsename, dose, frequency) values (102, 'Pen', 'Sam', 1, 1);
INSERT INTO ORDERS (RxNum, ord_medcode, ord_horsename, dose, frequency) values (103, 'Iver', 'John', 2, 1);
INSERT INTO ORDERS (RxNum, ord_medcode, ord_horsename, dose, frequency) values (104, 'Iver', 'Trotty', 2, 1);
INSERT INTO ORDERS (RxNum, ord_medcode, ord_horsename, dose, frequency) values (105, 'Doxy', 'Trotty', 1, 3);
INSERT INTO ORDERS (RxNum, ord_medcode, ord_horsename, dose, frequency) values (106, 'Pres', 'Shamrock', 1, 1);
INSERT INTO ORDERS (RxNum, ord_medcode, ord_horsename, dose, frequency) values (107, 'Dexa', 'Shamrock', 2, 1);
INSERT INTO ORDERS (RxNum, ord_medcode, ord_horsename, dose, frequency) values (108, 'Dexa', 'Betty', 2, 1);
INSERT INTO ORDERS (RxNum, ord_medcode, ord_horsename, dose, frequency) values (109, 'Xyl', 'Hulk', 1, 1);
INSERT INTO ORDERS (RxNum, ord_medcode, ord_horsename, dose, frequency) values (110, 'Enro', 'Erica', 3, 2);
```

MG's Horse Habitat – Horses and Medications

- You cannot INSERT an order for a medication that does not exist:

```
INSERT INTO ORDERS (RxNum, ord_medcode, ord_horsename, dose, frequency) values (999, 'NottaMed', 'Sam', 1, 2);
```



SQL Error [2291] [23000]: ORA-02291:
integrity constraint
(GMGRIMES.FK_MEDS) violated -
parent key not found

- Likewise you cannot INSERT an order for a horse that does not exist:

```
INSERT INTO ORDERS (RxNum, ord_medcode, ord_horsename, dose, frequency) values (888, 'Iver', 'NottaHorse', 1, 2);
```



SQL Error [2291] [23000]: ORA-02291:
integrity constraint
(GMGRIMES.FK_HORSES) violated -
parent key not found

- Remember: All values of a **Foreign Key** must be in the domain of values of the **Candidate Key** to which it refers

MG's Horse Habitat – Horses and Medications

- You cannot DELETE a medication that has been ordered for a horse:

```
DELETE FROM Medications WHERE medcode = 'Iver';
```



SQL Error [2292] [23000]: ORA-02292:
integrity constraint
(GMGRIMES.FK_MEDS) violated -
child record found

- But you CAN delete a medication that has NOT been ordered for a horse

```
DELETE FROM Medications WHERE medcode = 'Bute';
```

MG's Horse Habitat – Horses and Medications

- If you delete a horse (Parent), all records of that horses orders (Child) are also deleted (because of the CASCADE constraint applied to the foreign key)

	ABC NAME ▼	ABC COLOR ▼	ABC SPOTS ▼	ABC SEX ▼	123 WEIGHT ▼	ABC OWNER ▼
1	Sam	Brown	No	F	1,500	mgrimes
2	Erica	Yellow	Yes	F	920	canderson
3	John	Grey	No	M	1,800	mgrimes
4	Trotty	Brown	Yes	M	1,300	mgrimes
5	Rio	Grey	No	F	1,700	tswift
6	Robin	Yellow	No	M	1,100	jisbell
7	Katy	Brown	No	F	1,200	jisbell
8	Pegasus	Brown	No	M	1,750	mgrimes
9	Sammy	Black	Yes	M	2,200	mgrimes
10	Pinky	Red	No	M	1,050	tswift
11	Hulk	Grey	No	M	2,050	mgrimes
12	Pat	White	No	F	1,400	mgrimes
13	Betty	White	Yes	F	1,250	tswift
14	Shamrock	Black	No	M	1,400	ssimpson

	123 RXNUM ▼	ABC ORD_MEDCODE ▼	ABC ORD_HORSENAME ▼	123 DOSE ▼	123 FREQUENCY ▼
1	101	Flux	Sam	1	2
2	102	Pen	Sam	1	1
3	103	Iver	John	2	1
4	104	Iver	Trotty	2	1
5	105	Doxy	Trotty	1	3
6	106	Pres	Shamrock	1	1
7	107	Dexa	Shamrock	2	1
8	108	Dexa	Betty	2	1
9	109	Xyl	Hulk	1	1
10	110	Enro	Erica	3	2

DELETE FROM Horses WHERE name = 'Shamrock';

	ABC NAME ▼	ABC COLOR ▼	ABC SPOTS ▼	ABC SEX ▼	123 WEIGHT ▼	ABC OWNER ▼
1	Sam	Brown	No	F	1,500	mgrimes
2	Erica	Yellow	Yes	F	920	canderson
3	John	Grey	No	M	1,800	mgrimes
4	Trotty	Brown	Yes	M	1,300	mgrimes
5	Rio	Grey	No	F	1,700	tswift
6	Robin	Yellow	No	M	1,100	jisbell
7	Katy	Brown	No	F	1,200	jisbell
8	Pegasus	Brown	No	M	1,750	mgrimes
9	Sammy	Black	Yes	M	2,200	mgrimes
10	Pinky	Red	No	M	1,050	tswift
11	Hulk	Grey	No	M	2,050	mgrimes
12	Pat	White	No	F	1,400	mgrimes
13	Betty	White	Yes	F	1,250	tswift

	123 RXNUM ▼	ABC ORD_MEDCODE ▼	ABC ORD_HORSENAME ▼	123 DOSE ▼	123 FREQUENCY ▼
1	101	Flux	Sam	1	2
2	102	Pen	Sam	1	1
3	103	Iver	John	2	1
4	104	Iver	Trotty	2	1
5	105	Doxy	Trotty	1	3
6	108	Dexa	Betty	2	1
7	109	Xyl	Hulk	1	1
8	110	Enro	Erica	3	2

- This deletes ONE row from the Horses table and TWO rows from the Orders tables

MG's Horse Habitat – Horses and Medications

- If we want to see all our orders for medications, we can use the INNER JOIN introduced in Lecture 3 (we will talk much more about this after the exam)
 - Remember, all joins are binary (only two tables), so we first join Horses to Orders, then join that resulting table to Medications

```
SELECT * FROM Horses INNER JOIN orders ON name=ord_horsename  
INNER JOIN medications ON ord_medcode = medcode;
```

	ABC NAME	ABC COLOR	ABC SPOTS	ABC SEX	123 WEIG	ABC OWNER	123 RXNUM	ABC ORI	ABC ORD_	123 DOSI	123 FREQI	ABC MEDCODE	ABC MEDNAME	ABC CLASSIFICATION	123 COST	123 QOH	123 QOO
1	Sam	Brown	No	F	1,500	mgrimes	101	Flux	Sam	1	2	Flux	Flunixin Meglumine	NSAID	27	43	0
2	Sam	Brown	No	F	1,500	mgrimes	102	Pen	Sam	1	1	Pen	Penicillin	Antibiotic	38	73	0
3	Trotty	Brown	Yes	M	1,300	mgrimes	105	Doxy	Trotty	1	3	Doxy	Doxycycline	Antibiotic	81	89	0
4	Erica	Yellow	Yes	F	920	canderson	110	Enro	Erica	3	2	Enro	Enrofloxacin	Antibiotic	36	78	0
5	Hulk	Grey	No	M	2,050	mgrimes	109	Xyl	Hulk	1	1	Xyl	Xylazine	Sedative	22	28	0
6	Trotty	Brown	Yes	M	1,300	mgrimes	104	Iver	Trotty	2	1	Iver	Ivermectin	Dewormer	39	21	0
7	John	Grey	No	M	1,800	mgrimes	103	Iver	John	2	1	Iver	Ivermectin	Dewormer	39	21	0
8	Shamrock	Black	No	M	1,400	ssimpson	106	Pres	Shamrock	1	1	Pres	Prednisone	Corticosteroid	47	48	7
9	Shamrock	Black	No	M	1,400	ssimpson	107	Dexa	Shamrock	2	1	Dexa	Dexamethasone	Corticosteroid	47	18	82
10	Betty	White	Yes	F	1,250	tswift	108	Dexa	Betty	2	1	Dexa	Dexamethasone	Corticosteroid	47	18	82

MG's Horse Habitat – Horses and Medications


	ABC NAME	ABC COLOR	ABC SPOTS	ABC SEX	123 WEIG	ABC OWNER	123 RXNUM	ABC ORI	ABC ORD_	123 DOSI	123 FREQI	ABC MEDCODE	ABC MEDNAME	ABC CLASSIFICATION	123 COST	123 QOH	123 QOO
1	Sam	Brown	No	F	1,500	mgrimes	101	Flux	Sam	1	2	Flux	Flunixin Meglumine	NSAID	27	43	0
2	Sam	Brown	No	F	1,500	mgrimes	102	Pen	Sam	1	1	Pen	Penicillin	Antibiotic	38	73	0
3	Trotty	Brown	Yes	M	1,300	mgrimes	105	Doxy	Trotty	1	3	Doxy	Doxycycline	Antibiotic	81	89	0
4	Erica	Yellow	Yes	F	920	canderson	110	Enro	Erica	3	2	Enro	Enrofloxacin	Antibiotic	36	78	0
5	Hulk	Grey	No	M	2,050	mgrimes	109	Xyl	Hulk	1	1	Xyl	Xylazine	Sedative	22	28	0
6	Trotty	Brown	Yes	M	1,300	mgrimes	104	Iver	Trotty	2	1	Iver	Ivermectin	Dewormer	39	21	0
7	John	Grey	No	M	1,800	mgrimes	103	Iver	John	2	1	Iver	Ivermectin	Dewormer	39	21	0
8	Shamrock	Black	No	M	1,400	ssimpson	106	Pres	Shamrock	1	1	Pres	Prednisone	Corticosteroid	47	48	7
9	Shamrock	Black	No	M	1,400	ssimpson	107	Dexa	Shamrock	2	1	Dexa	Dexamethasone	Corticosteroid	47	18	82
10	Betty	White	Yes	F	1,250	tswift	108	Dexa	Betty	2	1	Dexa	Dexamethasone	Corticosteroid	47	18	82

- ...There is a lot of extra information here, so we can “project” a subset of attributes we are interested in (again, we will talk much more about this after the exam) :


```
SELECT Horses.Name, medications.MEDNAME, orders.dose, orders.FREQUENCY FROM Horses
INNER JOIN orders ON name=ord_horsename
INNER JOIN medications ON ord_medcode = medcode;
```

	ABC NAME	ABC MEDNAME	123 DOSE	123 FREQUENCY
1	Sam	Flunixin Meglumine	1	2
2	Sam	Penicillin	1	1
3	Trotty	Doxycycline	1	3
4	Erica	Enrofloxacin	3	2
5	Hulk	Xylazine	1	1
6	Trotty	Ivermectin	2	1
7	John	Ivermectin	2	1
8	Shamrock	Prednisone	1	1
9	Betty	Dexamethasone	2	1
10	Shamrock	Dexamethasone	2	1


Our current Conceptual ERD (and semantic Integrity Constraints)



Semantic Integrity Constraints:
Color: {Black, White, Brown, Grey, Red, Yellow}, Default: 'UNK'
Spots: {Yes, No}, Default: 'UNK'
Sex: {M, F}
Weight: {800-2200}
Cost: {1-200}
QoH + QoO: {10-100}
Dose: Default: 1
Frequently: Default 2



Our current decomposed ERD



Same overall meaning, just different levels of specificity!

Final DDL for now (Creating tables)

```
CREATE TABLE customers
(username varchar(50) CONSTRAINT pk_customers PRIMARY KEY,
Fname    varchar(50) CONSTRAINT nn_fname NOT NULL,
Lname    varchar(50) CONSTRAINT nn_lname NOT NULL,
Phone    varchar(14) CONSTRAINT nn_Phone NOT NULL
);

CREATE TABLE horses
(Name     varchar(50) CONSTRAINT pk_horse PRIMARY KEY,
Color    varchar(50) DEFAULT 'UNK' CONSTRAINT chk_color CHECK (color IN
('Black','White','Brown','Grey','Red','Yellow', 'UNK')),
Spots    varchar(3)  DEFAULT 'UNK',
Sex      varchar(1)  CONSTRAINT nn_sex NOT NULL,
Weight   integer,
owner    varchar(50),
CONSTRAINT chk_weight CHECK (weight >= 800 AND weight <=2200),
CONSTRAINT chk_sex CHECK (sex IN ('M','F')),
CONSTRAINT fk_cust FOREIGN KEY (owner) REFERENCES customers (username) ON DELETE CASCADE
);

CREATE TABLE medications
(medcode varchar(50) CONSTRAINT pk_medications PRIMARY KEY,
medname  varchar(50) CONSTRAINT nn_medname NOT NULL,
classification varchar(50) CONSTRAINT nn_classification NOT NULL,
cost     float(32) CONSTRAINT chk_cost CHECK (cost between 1 and 200),
QoH      integer,
QoO      integer,
CONSTRAINT chk_qty CHECK ((QoH + QoO) BETWEEN 10 AND 100)
);

CREATE TABLE orders
(RxNum integer CONSTRAINT pk_orders PRIMARY KEY,
ORD_Medcode varchar (50),
ORD_Horsename varchar(50),
Dose integer DEFAULT 1,
Frequency integer DEFAULT 2,
CONSTRAINT fk_meds FOREIGN KEY (ORD_Medcode) REFERENCES medications (medcode),
CONSTRAINT fk_horses FOREIGN KEY (ORD_Horsename) REFERENCES horses (name) ON DELETE CASCADE
);
```

Final DDL for now (Inserting data)

```
INSERT INTO customers (username, fname, lname, phone) VALUES ('mgrimes', 'Marvin', 'Grimes', '(218) 330-8004');
INSERT INTO customers (username, fname, lname, phone) VALUES ('canderson', 'Christine', 'Anderson', '(555) 523-9989');
INSERT INTO customers (username, fname, lname, phone) VALUES ('tswift', 'Tina', 'Swift', '(555) 424-1313');
INSERT INTO customers (username, fname, lname, phone) VALUES ('jisbell', 'Jason', 'Isbell', '(615) 555-5555');
INSERT INTO customers (username, fname, lname, phone) VALUES ('ssimpson', 'Sam', 'Simpson', '(615) 387-9682');
```

```
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('Sam', 'Brown', 'No', 'F', 1500, 'mgrimes');
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('Erica', 'Yellow', 'Yes', 'F', 920, 'canderson');
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('John', 'Grey', 'No', 'M', 1800, 'mgrimes');
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('Trotty', 'Brown', 'Yes', 'M', 1300, 'mgrimes');
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('Rio', 'Grey', 'No', 'F', 1700, 'tswift');
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('Robin', 'Yellow', 'No', 'M', 1100, 'jisbell');
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('Katy', 'Brown', 'No', 'F', 1200, 'jisbell');
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('Pegasus', 'Brown', 'No', 'M', 1750, 'mgrimes');
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('Sammy', 'Black', 'Yes', 'M', 2200, 'mgrimes');
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('Pinky', 'Red', 'No', 'M', 1050, 'tswift');
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('Hulk', 'Grey', 'No', 'M', 2050, 'mgrimes');
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('Pat', 'White', 'No', 'F', 1400, 'mgrimes');
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('Betty', 'White', 'Yes', 'F', 1250, 'tswift');
INSERT INTO horses (name, color, spots, sex, weight, owner) VALUES ('Shamrock', 'Black', 'No', 'M', 1400, 'ssimpson');
```

```
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Flux', 'Flunixin Meglumine', 'NSAID', 27, 43, 0);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Bute', 'Phenylbutazone', 'NSAID', 19, 84, 0);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Oxi', 'Oxibuzone', 'NSAID', 12, 55, 0);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Mel', 'Meloxicam', 'NSAID', 6, 72, 0);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Pen', 'Penicillin', 'Antibiotic', 38, 73, 0);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Doxy', 'Doxycycline', 'Antibiotic', 81, 89, 0);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('TMPS', 'Trimethoprim Sulfa', 'Antibiotic', 27, 50, 0);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Enro', 'Enrofloxacin', 'Antibiotic', 36, 78, 0);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Xyl', 'Xylazine ', 'Sedative', 22, 28, 0);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Acep', 'Acepromazine', 'Sedative', 33, 50, 0);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Fen', 'Fenbendazole', 'Dewormer', 52, 15, 25);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Meb', 'Mebendazole', 'Dewormer', 81, 25, 0);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Pyr', 'Pyrantel Embonate', 'Dewormer', 11, 29, 0);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Iver', 'Ivermectin', 'Dewormer', 39, 21, 0);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Clen', 'Clenbuterol', 'Respiratory', 132, 11, 20);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Pres', 'Prednisone', 'Corticosteroid', 47, 48, 7);
INSERT INTO Medications (medcode, medname, classification, cost, qoh, qoo) values ('Dexa', 'Dexamethasone ', 'Corticosteroid', 47, 18, 82);
```

```
INSERT INTO ORDERS (RxNum, ord_medcode, ord_horsename, dose, frequency) values (101, 'Flux', 'Sam', 1, 2);
INSERT INTO ORDERS (RxNum, ord_medcode, ord_horsename, dose, frequency) values (102, 'Pen', 'Sam', 1, 1);
INSERT INTO ORDERS (RxNum, ord_medcode, ord_horsename, dose, frequency) values (103, 'Iver', 'John', 2, 1);
INSERT INTO ORDERS (RxNum, ord_medcode, ord_horsename, dose, frequency) values (104, 'Iver', 'Trotty', 2, 1);
INSERT INTO ORDERS (RxNum, ord_medcode, ord_horsename, dose, frequency) values (105, 'Doxy', 'Trotty', 1, 3);
INSERT INTO ORDERS (RxNum, ord_medcode, ord_horsename, dose, frequency) values (106, 'Pres', 'Shamrock', 1, 1);
INSERT INTO ORDERS (RxNum, ord_medcode, ord_horsename, dose, frequency) values (107, 'Dexa', 'Shamrock', 2, 1);
INSERT INTO ORDERS (RxNum, ord_medcode, ord_horsename, dose, frequency) values (108, 'Dexa', 'Betty', 2, 1);
INSERT INTO ORDERS (RxNum, ord_medcode, ord_horsename, dose, frequency) values (109, 'Xyl', 'Hulk', 1, 1);
INSERT INTO ORDERS (RxNum, ord_medcode, ord_horsename, dose, frequency) values (110, 'Enro', 'Erica', 3, 2);
```

Break


Module 6.1

The Relational Data Model

- Why do we care about logical data modeling?
- Be able to explain terms as used in Relational Data Models
 - Relation
 - Tuple
 - Cardinality
 - Degrees

The Relational Data Model

- Edgar F. Codd in 1970 used the concept of mathematical relations to define the relational data model
- Database: collection of relations
- Relation: two-dimensional table
- Tuple: row of related data values in the table
- Attribute: column in the table
- Domain: set of possible atomic values of an attribute
- Degree of the relation: number of attributes in a relation
- Cardinality of the relation: number of tuples in a relation



A bold, and true, statement:

- The foundation of modern database technology is without question the relational model; it is that foundation that makes the field a science. Thus any book on the fundamentals of database technology that does not include a thorough coverage of the relational model is by definition shallow. Likewise any claim to expertise in the database field can hardly be justified if the claimant does not understand the relational model in depth.

Dale, C.J., Darwin, H., Foundation for Object/Relational Databases: The Third Manifesto

The Motivation for Logical Data Modeling

- Completion of conceptual modeling phase results in a picture of data requirements at high level of abstraction
- During conceptual modeling, we are not constrained by technology limitations that will be used for implementation
 - We got to use things like multi-value attributes and m:n relationships
- Conceptual schema may contain constructs not directly compatible with technology intended for implementation
 - Like multi-value attributes and m:n relationships!

The Motivation for Logical Data Modeling

- Further refinement may be required to eliminate data redundancy in design
 - Getting rid of multi-value attributes
 - Decomposing m:n relationships into 1:m using gerunds
- **Transforming conceptual schema to something that is more compatible with implementation technology of choice** is achieved via logical data modeling
- Logical data modeling phase serves as transition from technology-independent conceptual schema to technology-dependent design that can actually be implemented
- The next step will be generating the actual SQL code to create the entities, attributes, and relationships

Warning



- This may be confusing at first!
- Conceptual modeling uses many of the same concepts/terms as logical modeling, but in a slightly different (but really the same) way
 - Cardinality
 - Degrees of relationships
 - “Relations” are different than “relationships”
 - But a relationship is a relation...

Relation: two-dimensional table

- Two components of a Relation
 - Heading – a single tuple listing the attributes (Relation Schema)
 - Body – collection of data tuples

Attributes

Heading
(Single tuple of Attributes
"relation schema")

Body
(set of
data tuples)

Plants	PI No	PI_Name	PI_Budget
	11	Black Horse	2500000
	13	Mayde Creek	1930000
	12	Whitefield	2910000
	17	River Oaks	1930000
	19	King's Island	2500000

Relation
Degrees: 3
Cardinality: 5

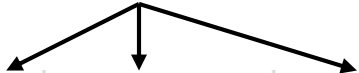
Plants	PI No	PI_Name	PI_Budget
	11	Black Horse	2500000
	12	Whitefield	2910000
	17	River Oaks	1930000
	19	King's Island	2500000

Relation
Degrees: 3
Cardinality: 4

Degree of the relation: number of attributes

- Degree is the number of attributes (columns) in a relation
 - But why?

Attributes





Plants	PI No	PI_Name	PI_Budget
	11	Black Horse	2500000
	13	Mayde Creek	1930000
	12	Whitefield	2910000
	17	River Oaks	1930000
	19	King's Island	2500000

Relation
Degrees: 3
Cardinality: 5

Degree of the relation: number of attributes


- Degree is the number of attributes (columns) in a relation
 - But why?
- Recall that an entity is a collection of related attributes
- Imagine this: each attribute of an entity can be considered an entity that is associated with the other attributes...



- We would never model it this way
 - ...but this is how degree describes the number of attributes

Degree of the relation: number of attributes


- A working example - remember the relationship between hobbies and dependents?



- How many degrees in the participation relationship?
- Since we cannot have $m:n$ relationships in our logical model we decomposed this...


Degree of the relation: number of attributes

- When we decompose the m:n relationship, how many degrees are in the participants relation?



Degree of the relation: number of attributes

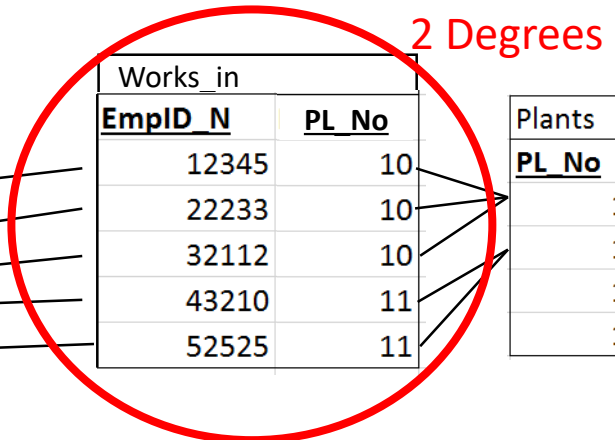
- Note: while it is not necessary to do so (and should generally not be done unless there are attributes of the relationship), 1:n and 1:1 relationships can also be decomposed using a gerund.



Employees					
<u>EmpID_N</u>	Fname	Minit	Lname	NameTag	PL_No
12345	Adam	B	Christie		10
22233	Danny	E	Francisco		10
32112	Greg		Hernandez	1	10
43210	Ivana	J	Klink		11
52525	Greg		Hernandez	2	11

Plants		
<u>PL_No</u>	PL_Name	Budget
10	Underwood	3000000
11	Garnett	3000000
12	Belmont	3500000
13	Vanderbilt	3500000

Employees				
<u>EmpID_N</u>	Fname	Minit	Lname	NameTag
12345	Adam	B	Christie	
22233	Danny	E	Francisco	
32112	Greg		Hernandez	1
43210	Ivana	J	Klink	
52525	Greg		Hernandez	2



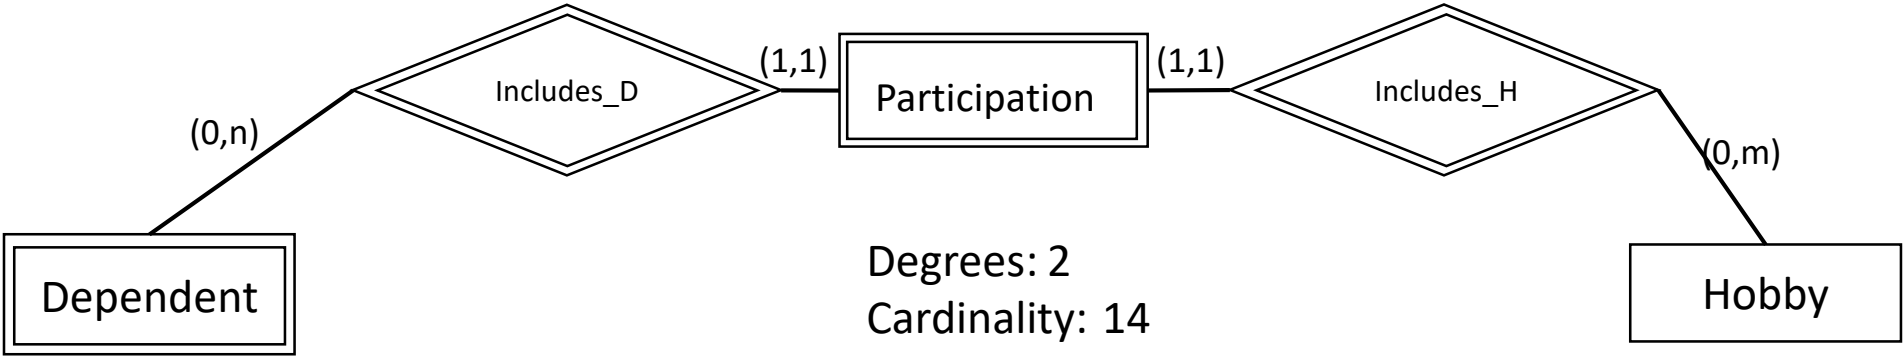
Works_in	
<u>EmpID_N</u>	<u>PL_No</u>
12345	10
22233	10
32112	10
43210	11
52525	11

Plants		
<u>PL_No</u>	PL_Name	Budget
10	Underwood	3000000
11	Garnett	3000000
12	Belmont	3500000
13	Vanderbilt	3500000

Cardinality of the relation: number of tuples

- Cardinality is the number of tuples (rows) in a relation
- When we say things like:
 - “many students may take a class”
- We really mean
 - “Any number of students up to number of students we have enrolled at the university may take a class”
 - ...which is the cardinality of the students relation
(Obviously other business rules may limit the number of student that may take a class – like the size of the classroom!)

Cardinality of the relation: number of tuples



Degrees: 2
Cardinality: 14

Degrees: 4
Cardinality: 5

Dependents			
<u>DepID</u>	Name	Birthdate	Gender
100	Mark	5/11/1975	M
101	Jill	5/4/1976	F
102	Norm	3/1/1984	M
103	Mike	1/31/1992	M
104	Sue	9/2/1990	F

Participation	
<u>DepID</u>	<u>HobbyID</u>
100	1
100	2
100	4
101	3
101	5
102	1
103	1
103	2
103	3
103	4
103	5
104	1
104	4
104	5


Degrees: 4
Cardinality: 5

Hobby			
<u>HobbyID</u>	HobbyName	In_Act	Gr_Act
1	Soft Ball	0	G
2	Flag Football	0	G
3	Knitting	I	I
4	Cycling	0	I
5	Movies	I	G

Cardinality of the relation: number of tuples


- Knowing the cardinality of a relation is important for understanding how much data is present...however, even more important is that it allows us to understand how large/complex join operations will be
- What is the maximum cardinality of the participants relation?

$5 * 5 = \text{Max cardinality of } 25$



Cardinality of the relation: number of tuples

- What if there were 5,000 dependents and 1,000 hobbies?
 - Max cardinality of 5,000,000
- ...and each hobby met once per month (12 times)?
 - Maximum cardinality of 60,000,000 (5,000 x 1,000 x 12)!
- We can reduce cardinality prior to joining in order to reduce computational complexity!



Module 6.1

The Relational Data Model

- Why do we care about logical data modeling?
- Be able to explain terms as used in Relational Data Models
 - Relation
 - Tuple
 - Cardinality
 - Degrees

Module 6.2

Characteristics of a Relation

- Be able to describe characteristics of a relation

Characteristics of a Relation

- A relation is a mathematical term that resembles a two-dimensional table
- Has a heading, which is a tuple of attributes, also known as the relation schema
- Has a body, which is made up of many tuples of data containing the same attributes
- Attributes of relation schema have unique names
- Values of an attribute in a relation come from same domain

Characteristics of a Relation (continued)

- Order of arrangement of tuples does not matter
- Order of attributes does not matter
- Each attribute value in tuple is atomic; hence, composite and multi-valued attributes are not allowed in a relation
- Derived attributes are not captured in relation schema
- All tuples in relation must be distinct (that is to say every relation schema must have unique identifier)
 - **If every tuple is distinct, then at a minimum the combination of the values of every attribute would be a unique identifier, right? We'll come back to this idea soon.**

Module 6.2

Characteristics of a Relation

- Be able to describe characteristics of a relation
- Describe three set operators

Module 6.3

Data Integrity Constraints

- Be able to define: Super keys, candidate keys, key attributes, non-key attributes, primary keys, and alternate keys
- What is the entity integrity constraint?
- What is the referential integrity constraint?
- What is the foreign key constraint?

A “Key” topic in your Database education as we talk about

- Super Keys
- Candidate Keys
- Primary Keys
- Alternate Keys
- Key Attributes
- Non-Key Attributes

Not to mention those we’ve already talked about:

- Partial Keys
- Foreign Keys

Data Integrity Constraints

- Rules that govern the behavior of data in a database
- Are technical expressions of business rules that emerge from user requirement specifications for database application
- Prevail across all tiers of data modeling – conceptual, logical, and physical
- Are considered to be part of the schema
 - Declared along with structural design of data model and hold for all valid states of a database and at all levels (conceptual, external, and internal)

The Concept of Unique Identifiers

- Superkey:
 - A set of one or more attributes, which taken collectively, uniquely identifies a tuple of a relation {uniqueness property}
- Candidate Key:
 - A superkey with no proper subset that uniquely identifies a tuple of a relation {uniqueness property + irreducibility}
- Primary Key:
 - A candidate key with no missing values for the constituent attributes {uniqueness property + irreducibility + entity integrity constraint}
- Alternate Key:
 - Any candidate key that is not serving the role of the primary key

A Sample Relation Instance: Prescription A

Rx_rx#	Rx_pat#	Rx_medcode	Rx-dosage
A100	7642	PCN	3
A103	4678	TYL	2
A102	4772	CLR	2
A101	6742	ASP	2
A104	4772	ZAN	3
A105	7456	CLR	2
A107	2222	TYL	2
A106	4772	VAL	2
A108	7384	CLR	3
A109	7384	ZAN	2
A110	7642	VAL	2

Prescription A

Superkey: one or more attributes which together can uniquely identify a tuple


- Rx_rx#
- (Rx_rx#, Rx_pat#)
- (Rx_rx#, Rx_medcode)
- (Rx_rx#, Rx_dosage)
- (Rx_pat#, Rx_medcode)
- (Rx_pat#, Rx_medcode, Rx_dosage)
- (Rx_pat#, Rx_medcode, Rx_rx#)
- (Rx_rx#, Rx_pat#, Rx_dosage)
- (Rx_rx#, Rx_medcode, Rx_dosage)
- (Rx_rx#, Rx_pat#, Rx_medcode, Rx_dosage)

Rx_rx#	Rx_pat#	Rx_medcode	Rx-dosage
A100	7642	PCN	3
A103	4678	TYL	2
A102	4772	CLR	2
A101	6742	ASP	2
A104	4772	ZAN	3
A105	7456	CLR	2
A107	2222	TYL	2
A106	4772	VAL	2
A108	7384	CLR	3
A109	7384	ZAN	2
A110	7642	VAL	2

Prescription A

Superkey: one or more attributes which together can uniquely identify a tuple

- In general if K is a superkey, then any superset of K (i.e., any set that contains K) is also a superkey
- Superkeys uniquely define a tuple, but there are lots of them – so let's do better



The diagram shows two arrows labeled 'Unique'. The first arrow points down to the Rx_pat# column header. The second arrow is a double-headed arrow spanning the Rx_pat# and Rx_medcode columns, indicating a unique constraint on the combination of these two attributes.

Rx_rx#	Rx_pat#	Rx_medcode	Rx-dosage
A100	7642	PCN	3
A103	4678	TYL	2
A102	4772	CLR	2
A101	6742	ASP	2
A104	4772	ZAN	3
A105	7456	CLR	2
A107	2222	TYL	2
A106	4772	VAL	2
A108	7384	CLR	3
A109	7384	ZAN	2
A110	7642	VAL	2

Prescription A

Candidate key (uniqueness + irreducibility)

- A superkey with no proper subset that uniquely identifies a tuple
- Rx_rx# ← Yes, Candidate key
- (Rx_rx#, Rx_pat#) ← No, the proper subset Rx_rx# uniquely identifies a tuple
- (Rx_rx#, Rx_medcode) ← No, the proper subset Rx_rx# uniquely identifies a tuple
- (Rx_rx#, Rx_dosage) ← No, the proper subset Rx_rx# uniquely identifies a tuple
- (Rx_pat#, Rx_medcode) ← Yes, Candidate key
- (Rx_pat#, Rx_medcode, Rx_dosage) ← No, why?
- (Rx_pat#, Rx_medcode, Rx_rx#) ← No, why?
- (Rx_rx#, Rx_pat#, Rx_dosage) ← No, why?
- (Rx_rx#, Rx_medcode, Rx_dosage) ← No, why?
- (Rx_rx#, Rx_pat#, Rx_medcode, Rx_dosage) ← No, why?

Key/Non-Key Attributes

- An attribute is always a key attribute, non-key attribute, or candidate key
- Key attribute:
 - Any attribute that is a proper subset of a candidate key
 - Note: A set is a subset of itself, but is not a **proper** subset
- Non-key attribute:
 - Any attribute that is not a subset of a candidate key
- Example:
 - Rx_rx# is not a key attribute of PRESCRIPTION-A since it is not a proper subset of a candidate key.
 - Rx_rx# is not a non-key attribute of PRESCRIPTION-A since it is a subset of a candidate key.
 - Rx_rx# is a candidate key of PRESCRIPTION-A since it is an irreducible superkey of PRESCRIPTION-A.

Key, Non-Key, or candidate key attribute?

- We identified two candidate keys for the prescription-A relation

- Rx_rx# Candidate Key
- (Rx_pat#, Rx_medcode) Candidate Key

- There are four atomic attributes in the prescription-A relation.

- Rx_rx# Candidate Key
- Rx_pat# Key Attribute
- Rx_medcode Key Attribute
- Rx_dosage Non-key Attribute

- An attribute is one and only one of the 3

- Candidate keys themselves are not key attributes

- A set is not a proper subset of itself

Rx_rx#	Rx_pat#	Rx_medcode	Rx-dosage
A100	7642	PCN	3
A103	4678	TYL	2
A102	4772	CLR	2
A101	6742	ASP	2
A104	4772	ZAN	3
A105	7456	CLR	2
A107	2222	TYL	2
A106	4772	VAL	2
A108	7384	CLR	3
A109	7384	ZAN	2
A110	7642	VAL	2

Primary Key

- A primary Key is a candidate key with one additional constraint: It must not be missing (NULL)
 - Entity integrity constraint
- Which to choose?
 - Technically does not matter, but choose the easy one (and one that will never be NULL)
 - Candidate keys not chosen as PK are alternate keys



Rx_rx#	Unique		
	Rx_pat#	Rx_medcode	Rx-dosage
A100	7642	PCN	3
A103	4678	TYL	2
A102	4772	CLR	2
A101	6742	ASP	2
A104	4772	ZAN	3
A105	7456	CLR	2
A107	2222	TYL	2
A106	4772	VAL	2
A108	7384	CLR	3
A109	7384	ZAN	2
A110	7642	VAL	2

Another example – Prescription B

Diagram illustrating the uniqueness of columns in Prescription B:

- Not Unique:** Points to the Rx_rx# column.
- Unique:** Points to the Rx_pat# and Rx_medcode columns.

Rx_rx#	Rx_pat#	Rx_medcode	Rx_dosage
B100	7642	PCN	3
B103	4678	TYL	2
B102	4772	CLR	2
B101	6742	ASP	2
B102	4772	ZAN	2
B105	7456	CLR	2
B107	2222	TYL	2
B106	4772	VAL	2
B108	7384	CLR	3
B109	7384	ZAN	2
B100	7642	VAL	2

Prescription B

What are the superkeys?

- (Rx_pat#, Rx_medcode)
- (Rx_pat#, Rx_medcode, Rx_rx#)
- (Rx_pat#, Rx_medcode, Rx_dosage#)
- (Rx_rx#, Rx_medcode)
- (Rx_rx#, Rx_medcode, Rx_dosage#)
- (Rx_rx#, Rx_medcode, Rx_pat#, Rx_dosage#)

Rx_rx#	Rx_pat#	Rx_medcode	Rx_dosage
B100	7642	PCN	3
B103	4678	TYL	2
B102	4772	CLR	2
B101	6742	ASP	2
B102	4772	ZAN	2
B105	7456	CLR	2
B107	2222	TYL	2
B106	4772	VAL	2
B108	7384	CLR	3
B109	7384	ZAN	2
B100	7642	VAL	2

Prescription B

What are the candidate keys?

- (Rx_pat#, Rx_medcode) 😊
- (Rx_pat#, Rx_medcode, Rx_rx#) ❌
- (Rx_pat#, Rx_medcode, Rx_dosage#) ❌
- (Rx_rx#, Rx_medcode) 😊
- (Rx_rx#, Rx_medcode, Rx_dosage#) ❌
- (Rx_rx#, Rx_medcode, Rx_pat#, Rx_dosage#) ❌

Diagram illustrating candidate keys for the Prescription B table:

- Not Unique** (points to Rx_rx#)
- Unique** (points to Rx_pat#)
- Unique** (points to Rx_medcode)

Rx_rx#	Rx_pat#	Rx_medcode	Rx_dosage
B100	7642	PCN	3
B103	4678	TYL	2
B102	4772	CLR	2
B101	6742	ASP	2
B102	4772	ZAN	2
B105	7456	CLR	2
B107	2222	TYL	2
B106	4772	VAL	2
B108	7384	CLR	3
B109	7384	ZAN	2
B100	7642	VAL	2

Prescription B

What attributes are keys vs. non-keys?

- We identified two candidate keys for the prescription-B relation

- (Rx_pat#, Rx_medcode)

Candidate Key

- (Rx_rx#, Rx_medcode)

Candidate Key

- There are four atomic attributes in the prescription-B relation

- Rx_rx#

Key Attribute

- Rx_pat#

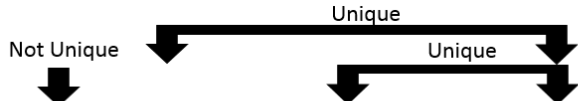
Key Attribute

- Rx_medcode

Key Attribute

- Rx_dosage

Non-key Attribute



Rx_rx#	Rx_pat#	Rx_medcode	Rx_dosage
B100	7642	PCN	3
B103	4678	TYL	2
B102	4772	CLR	2
B101	6742	ASP	2
B102	4772	ZAN	2
B105	7456	CLR	2
B107	2222	TYL	2
B106	4772	VAL	2
B108	7384	CLR	3
B109	7384	ZAN	2
B100	7642	VAL	2

Prescription B

Superkey and Candidate Key (page 287)

PRESCRIPTION (Rx_rx#, Rx_pat#, Rx_medcode, Rx_dosage)				
	PRESCRIPTION-A		PRESCRIPTION-B	
	Superkey	Candidate Key	Superkey	Candidate Key
Rx_rx#	Yes	Yes	No	No
Rx_pat#	No	No	No	No
Rx_medcode	No	No	No	No
Rx_dosage	No	No	No	No
Rx_rx#, Rx_pat#	Yes	No	No	No
Rx_rx#, Rx_medcode	Yes	No	Yes	Yes
Rx_rx#, Rx_dosage	Yes	No	No	No
Rx_pat#, Rx_medcode	Yes	Yes	Yes	Yes
Rx_pat#, Rx_dosage	No	No	No	No
Rx_medcode, Rx_dosage	No	No	No	No
Rx_rx#, Rx_pat#, Rx_medcode	Yes	No	Yes	No
Rx_rx#, Rx_pat#, Rx_dosage	Yes	No	No	No
Rx_rx#, Rx_medcode, Rx_dosage	Yes	No	Yes	No
Rx_pat#, Rx_medcode, Rx_dosage	Yes	No	Yes	No
Rx_rx#, Rx_pat#, Rx_medcode, Rx_dosage	Yes	No	Yes	No

Table 6.2 Superkeys and candidate keys in the PRESCRIPTION-A and PRESCRIPTION-B relations

Referential Integrity Constraint


- Key constraints (superkey and candidate key) and entity integrity constraint (primary key) pertain to individual relation schemas
- Referential integrity constraints are specified between two relation schemas (i.e., R1 and R2)
- Specifically, a referential integrity constraint is specified between two relations in order to maintain consistency across tuples of the two relations
- Informal definition: A tuple in one relation that refers to another relation must refer to an existing tuple in that relation.

Foreign Key Constraint

- A special form of referential integrity constraint specification
- Establishes an explicit association between two relation schemas and maintains the integrity of such an association
- Foreign key: An attribute(s) set, A2, in a relation schema R2 that shares the same domain with a **candidate key** (A1) of another relation schema R1; A2 is said to reference or refer to the relation schema R1.
 - Note: R2 is known as the referencing relation and R1 is called the referenced relation. The attribute(s) doing the referencing (A2 in R2) is the foreign key, while the candidate key (**advisably, the primary key**) being referenced (A1 in R1) is the referenced attribute(s).
- Referred to as Inclusion Dependency, this constraint is algebraically expressed as:
 $R2.\{A2\} \subseteq R1.\{A1\}$
- Meaning: Child.foreignkey (R2.{A2}) is inclusion dependent on parent.primarykey (R1.{A1})

Example – Source Schema

- Bearcat Incorporated is a manufacturing company that has several plants in the northeastern part of the United States. These plants are responsible for leading different projects that the company might undertake, depending on a plants' function. A certain plant might even be associated with several projects but a project is always under the control of just one plant. Some plants do not undertake any projects at all.




FK in the 1:n relationship

PLANT	Pl_name	<u>Pl_p#</u>	Pl_budget
	Black Horse	11	1230000
	Mayde Creek	13	1930000
	Whitefield	12	2910000
	River Oaks	17	1930000
	King's Island	19	2500000
	Ashton	15	2500000

PROJECT	<u>Prj_name</u>	Prj_n#	Prj_pl_p#
	Solar Heating	41	11
	Lunar Cooling	17	17
	Synthetic Fuel	29	17
	Nitro-Cooling	23	12
	Robot Sweeping	31	11
	Robot Painting	37	19
	Ozone Control	13	19

Foreign Key



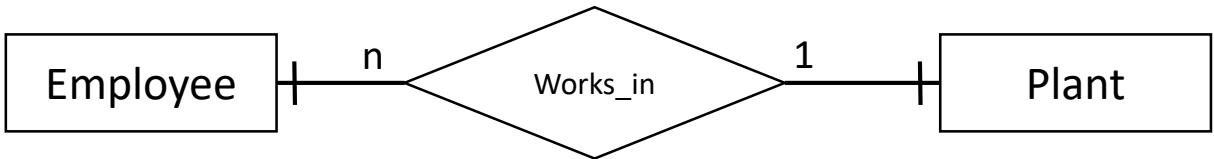
$$R2.\{A2\} \subseteq R1.\{A1\}$$

$$\text{Project.Prj_pl_P\#} \subseteq \text{Plant.Pl_p\#}$$

Note: PROJECT.Prj_pl_p# is the foreign key referencing PLANT.Pl_p#, the primary key of PLANT.

Note: Prj_n# would be a better primary key for the Project entity. We'll talk about why later, but this works for this example

Foreign key placement



- What is the cardinality of the relationship? 1:n
- Which entity is the parent? Plant
- In a 1:M relationship, Foreign key (FK) goes with the “child” side
 - The FK is an attribute in the child entity (employees in this case) that refers to the primary key of the entity on the other side of the relationship

Employees						
<u>EmpID_A</u>	<u>EmpID_N</u>	Fname	Minit	Lname	NameTag	PL_No
E	12345	Adam	B	Christie		10
E	22233	Danny	E	Francisco		10
C	32112	Greg		Hernandez	1	10
E	43210	Ivana	J	Klink		11
E	52525	Greg		Hernandez	2	11

Plants		
<u>PL_No</u>	PL_Name	Budget
10	Underwood	3000000
11	Garnett	3000000
12	Belmont	3500000
13	Vanderbilt	3500000

- Would this even make sense the other way?

Note this is dataset is truncated. In reality every entity in the plant table must have at least 100 employees (as specified in the business rule – the ERD above merely specifies that it must participate)

Note: Employees.PL_No is a FK that refers to Plants.PL_No. All values of Employees.PL_No must be found in the domain of values for Plants.PL_No (i.e., must be a valid plant)

In 1:m The FK must go with the child entity

- If we put the FK with the parent it is INCORRECT:
 - Lots of data redundancy in the Plants table if the FK is there

Employees					
EmpID_A	EmpID_N	Fname	Minit	Lname	NameTag
E	12345	Adam	B	Christie	
E	22233	Danny	E	Francisco	
C	32112	Greg		Hernande	1
E	43210	Ivana	J	Klink	
E	52525	Greg		Hernande	2

Plants			
PL_No	PL_Name	Budget	EmpID
10	Underwood	3000000	E12345
10	Underwood	3000000	E22233
10	Underwood	3000000	C32112
11	Garnett	3000000	E43210
11	Garnett	3000000	E52525
12	Belmont	3500000	NULL
13	Vanderbilt	3500000	NULL

Note: In this case, Plants.EmpID is a FK that refers to the composite attribute Employees.EmpID_A + Employees.EmpID_N.

There is now redundancy for the Plants.PL_Name and Plants.Budget attributes, and PL_No is no longer unique!

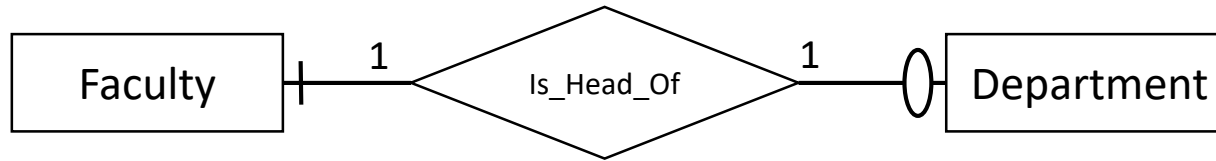
- No redundancy when FK is with Employees (CORRECT):

Employees					
EmpID_A	EmpID_N	Fname	Minit	Lname	NameTag
E	12345	Adam	B	Christie	
E	22233	Danny	E	Francisco	
C	32112	Greg		Hernandez	1
E	43210	Ivana	J	Klink	
E	52525	Greg		Hernandez	2

Plants		
PL_No	PL_Name	Budget
10	Underwood	3000000
11	Garnett	3000000
12	Belmont	3500000
13	Vanderbilt	3500000

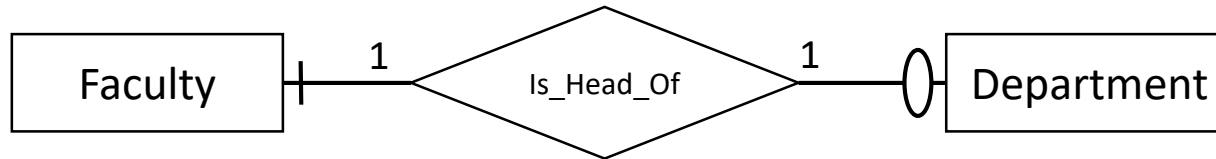
Where do we put the foreign key in a 1:1 relationship?

- There is no clear parent....



In 1:1 we place the FK based on participation

- FK goes with the entity that has total (mandatory) participation



Faculty		
<u>EmpID</u>	Name	Title
112	Chad Larson	Assistant Professor
108	Dusya Vera	Professor
109	Edward Blair	Professor
111	Gerald Lobo	Professor
106	Kris Jacobs	Professor
102	Mark Grimes	Assistant Professor
110	Melanie Rudd	Assistant Professor
103	Michael Parks	Associate Professor
101	Norm Johnson	Professor
104	Praveen Kumar	Professor
107	Steve Werner	Professor
105	Thomas George	Professor

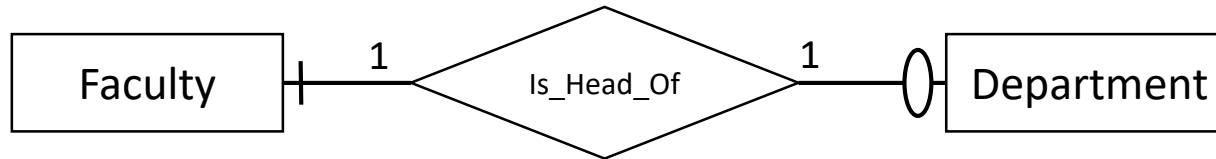
Departments		
<u>DeptID</u>	Deptname	Head
1	Accounting	111
2	DISC	101
3	Finanace	104
4	Management	107
5	Marketing	109

Note: In this case, Departments.Head is a FK that refers to Faculty.EmpID.

**CORRECT
WAY!**

In 1:1 we place the FK based on participation

- Imagine if we did it the other way, it still works, but lots of NULL values!
 - NULL values are difficult to work with – avoid them as much as possible!



Faculty			
<u>EmpID</u>	Name	Title	Head of
112	Chad Larson	Assistant Professor	
108	Dusya Vera	Professor	
109	Edward Blair	Professor	5
111	Gerald Lobo	Professor	1
106	Kris Jacobs	Professor	
102	Mark Grimes	Assistant Professor	
110	Melanie Rudd	Assistant Professor	
103	Michael Parks	Associate Professor	
101	Norm Johnson	Professor	2
104	Praveen Kumar	Professor	3
107	Steve Werner	Professor	4
105	Thomas George	Professor	

Departments	
<u>DeptID</u>	Deptname
1	Accounting
2	DISC
3	Finanace
4	Management
5	Marketing

Note: In this case, Faculty.HeadOf is a FK that refers to Departments.DeptID.

No data redundancy, but lots of nasty NULL values


**WRONG
WAY!**

In 1:1 we place the FK based on participation


- There are some fringe cases...
- If both sides of a 1:1 have mandatory participation
 - It doesn't matter where you place the FK
 - This would be a pretty rare thing to find
- If both sides of a 1:1 have optional participation
 - No hard and fast rule, but generally would be better to put it with the entity that will have fewer instances
 - Alternatively, you can use a gerund to ensure you will have no NULL values!

Watch out!

- The “is head of” relationships is different than the “member of” relationship:



- Where should the FK go for the “Member of” relationship? **Faculty (The Child)**



Note: In this case, Faculty.Department is a FK that refers to Departments.DeptID - this is the “Member_of” relationship.

Departments.Head is a FK that refers to Faculty.EmpID - this is the “Is_Head_of” relationship.

We have no redundancy and no NULL values with this design

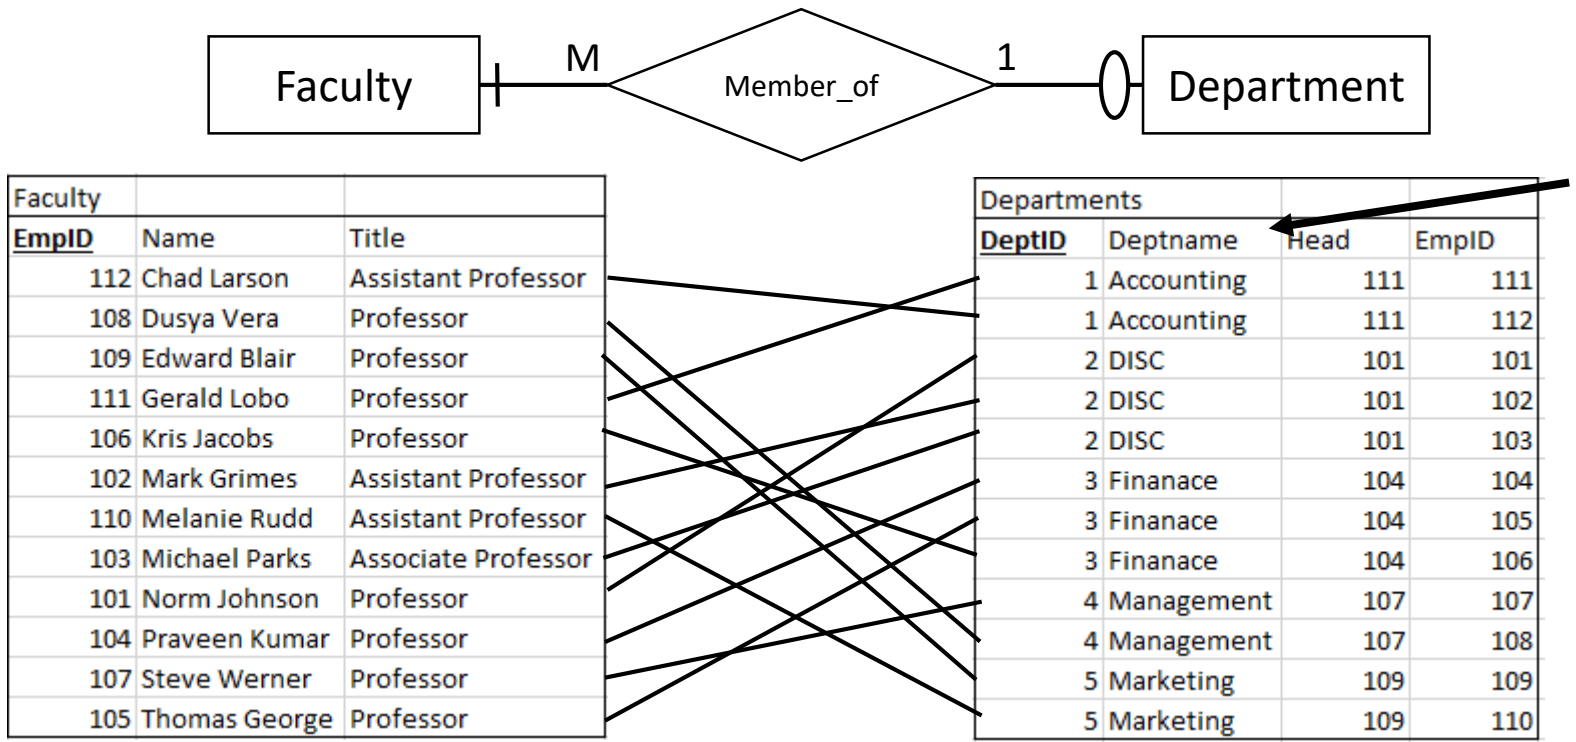
Faculty			
EmpID	Name	Title	Department
112	Chad Larson	Assistant Professor	1
108	Dusya Vera	Professor	4
109	Edward Blair	Professor	5
111	Gerald Lobo	Professor	1
106	Kris Jacobs	Professor	3
102	Mark Grimes	Assistant Professor	2
110	Melanie Rudd	Assistant Professor	5
103	Michael Parks	Associate Professor	2
101	Norm Johnson	Professor	2
104	Praveen Kumar	Professor	3
107	Steve Werner	Professor	4
105	Thomas George	Professor	3

Departments		
DeptID	Deptname	Head
1	Accounting	111
2	DISC	101
3	Finanace	104
4	Management	107
5	Marketing	109

CORRECT
WAY!

Watch out!

- What if we did it the wrong way?



Redundant data for the Deptname and Head attributes **AND** DeptID is no longer unique!

WRONG WAY!

Module 6.3

Data Integrity Constraints

- Be able to define: Super keys, candidate keys, key attributes, non-key attributes, primary keys, and alternate keys
- What is the entity integrity constraint?
- What is the referential integrity constraint?
- What is the foreign key constraint?

Progress Quiz Time!

- The Progress Quiz is available in Canvas
 - You MUST complete the quiz on Canvas by 5:00 on Friday – This in-class activity does not count for points!
 - Each week we will discuss the questions, so for those of you that are in class and keeping up with things, you'll have an extra easy time with it!
- Go to <http://kahoot.it> and we'll get started momentarily!

Go forth and do great things

- Get to work on assignment 2
- We will continue relational data modeling next week
- Exam 1 is coming up in two weeks – March 4

BZAN 6354

Live Lecture 5

February 19, 2024

Dr. Mark Grimes, Ph.D.
gmgrimes@bauer.uh.edu

UNIVERSITY of
HOUSTON

C. T. BAUER COLLEGE of BUSINESS
Department of Decision & Information Sciences