

# BZAN 6354

# Lecture 6

# February 26, 2024

Dr. Mark Grimes, Ph.D.  
gmgrimes@bauer.uh.edu

UNIVERSITY of  
**HOUSTON**

---

C. T. BAUER COLLEGE of BUSINESS  
Department of Decision & Information Sciences

# Agenda

- Quick Review
- Module 6.3 – Data Integrity Constraints (Continued...)
- Module 6.4 – Relational Algebra and sets
- Module 6.5 – Views and Materialized Views
- Break
- Module 6.7 – Logical Modeling
- Exam Discussion

# Last time we talked about...

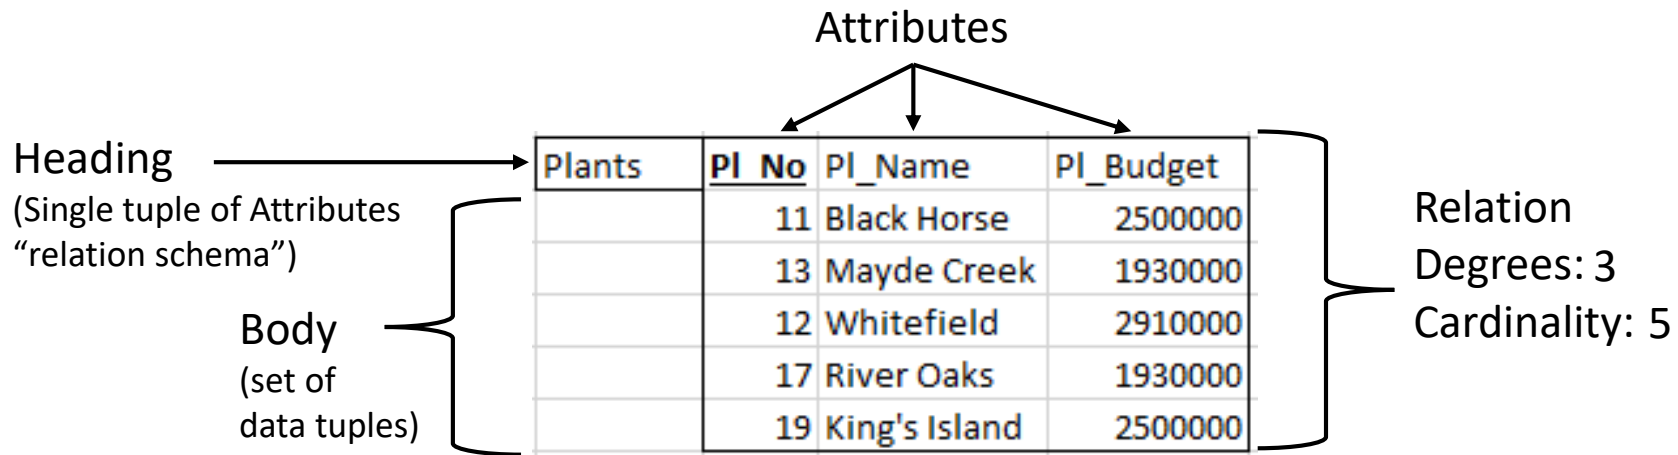
- Attribute Placement
  - Attributes mean different things depending on what entity they are part of, OR if they are an attribute of a relationship!
- Weak Entities / Partial Keys / Identifying relationships
  - Weak entities do not exist without the strong/base entity that “identifies” them
  - Weak entities have a partial key that is NOT unique by itself
- Decomposing M:N Relationships with the Gerund
  - Primary key from each participating relation become a foreign key in the gerund
- Data modeling errors
  - Syntactical Errors: Using the grammar incorrectly, relatively easy to spot
  - Semantic Errors: Not modeling the business case correctly, you must understand the case to spot these

# Last time we talked about...

- Relational modeling
  - Degree of a relation: Number of attributes
  - Cardinality of a relation: Number of tuples
  - Characteristics of a relation
- The first part of “Data Integrity Constraints”
  - Super Keys: Unique
  - Candidate Keys: Unique and Irreducible
  - Primary Key: Unique, Irreducible, and cannot be NULL (Entity Integrity Constraint)
  - Alternate Keys: Candidate keys not selected as primary key
  - Key attribute: Proper subset of a candidate key
  - Non-key attribute: Not a subset of a candidate key
- ...And we ran out of time before finishing integrity constraints and talking about foreign key placement!

# Review: What is a relation?

- A mathematical terms approximated by a two dimensional table:
  - Heading – a single tuple listing the attributes (Relation Schema)
  - Body – collection of data tuples



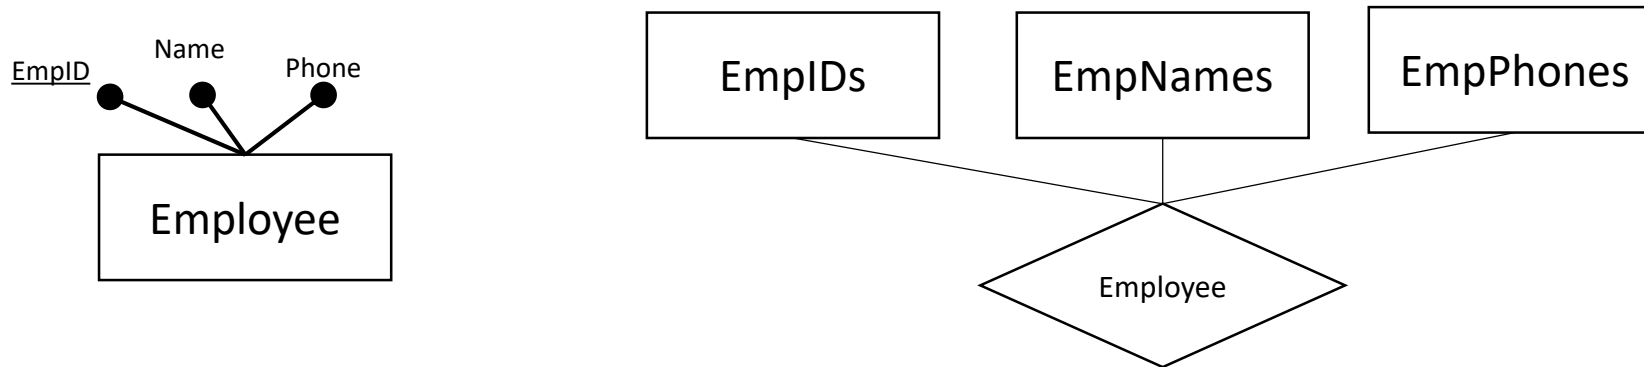
| Plants | PI No | PI_Name       | PI_Budget |
|--------|-------|---------------|-----------|
|        | 11    | Black Horse   | 2500000   |
|        | 12    | Whitefield    | 2910000   |
|        | 17    | River Oaks    | 1930000   |
|        | 19    | King's Island | 2500000   |

Relation  
Degrees: 3  
Cardinality: 4

The diagram shows a table with 4 columns: 'Plants', 'PI No', 'PI\_Name', and 'PI\_Budget'. The first column is labeled 'Plants' and is the heading. The other three columns are labeled 'PI No', 'PI\_Name', and 'PI\_Budget' and are the attributes. The table has 4 data tuples. The relation has 3 degrees and a cardinality of 4.

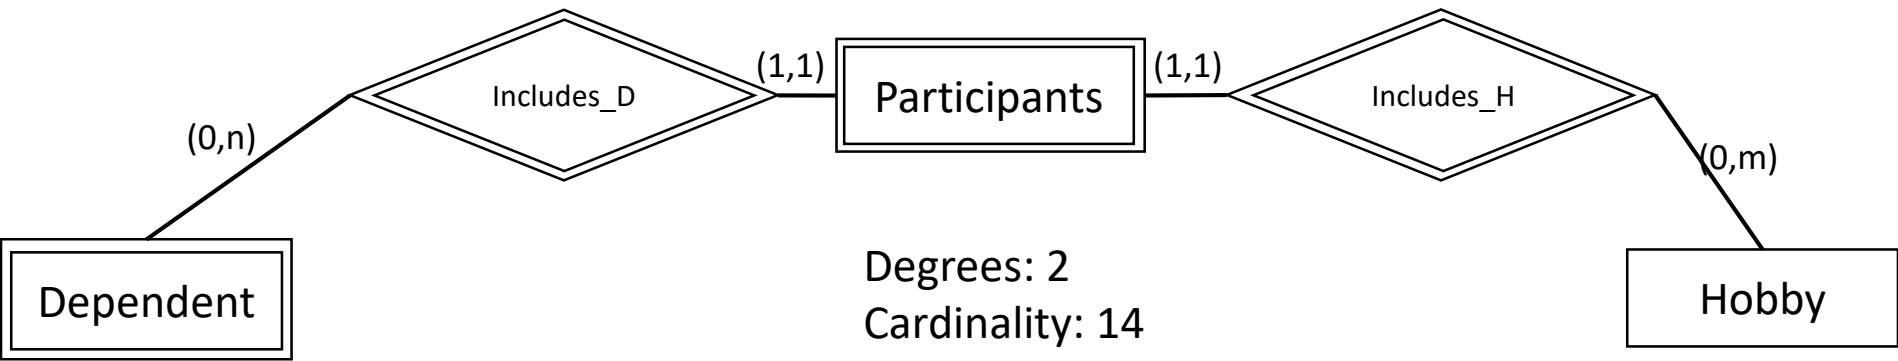
# Review: What is the degree of a relation?

- Degree is the number of attributes (columns) in a relation
  - But why?
- Imagine this: each attribute of an entity can be considered an entity that is associated with the other attributes...



- We would never model it this way
  - ...but this is how degree describes the number of attributes

# Review: What is the cardinality of a relation?



Degrees: 2  
Cardinality: 14

Degrees: 4  
Cardinality: 5

| Dependents   |      |           |        |
|--------------|------|-----------|--------|
| <u>EmpID</u> | Name | Birthdate | Gender |
| 100          | Mark | 5/11/1975 | M      |
| 101          | Jill | 5/4/1976  | F      |
| 102          | Norm | 3/1/1984  | M      |
| 103          | Mike | 1/31/1992 | M      |
| 104          | Sue  | 9/2/1990  | F      |

| Participants |                |
|--------------|----------------|
| <u>EmpID</u> | <u>HobbyID</u> |
| 100          | 1              |
| 100          | 2              |
| 100          | 4              |
| 101          | 3              |
| 101          | 5              |
| 102          | 1              |
| 103          | 1              |
| 103          | 2              |
| 103          | 3              |
| 103          | 4              |
| 103          | 5              |
| 104          | 1              |
| 104          | 4              |
| 104          | 5              |

Degrees: 4  
Cardinality: 5

| Hobby          |               |        |        |
|----------------|---------------|--------|--------|
| <u>HobbyID</u> | HobbyName     | In_Act | Gr_Act |
| 1              | Soft Ball     | O      | G      |
| 2              | Flag Football | O      | G      |
| 3              | Knitting      | I      | I      |
| 4              | Cycling       | O      | I      |
| 5              | Movies        | I      | G      |

# Exam 1 – Next week

- March 4
  - In class on paper
  - 75 minutes
  - Closed Book / Individual Effort
- Will be approximately one third each of:
  - Multiple choice
  - Short answer / Matching / Fill in the Blank / etc.
  - Creating ERDs



# Module 6.3

## Data Integrity Constraints

- Be able to define: Super keys, candidate keys, key attributes, non-key attributes, primary keys, and alternate keys (Already Done)
- What is the entity integrity constraint? (Already done)
- What is the referential integrity constraint?
- What is the foreign key constraint?

# Referential Integrity Constraint

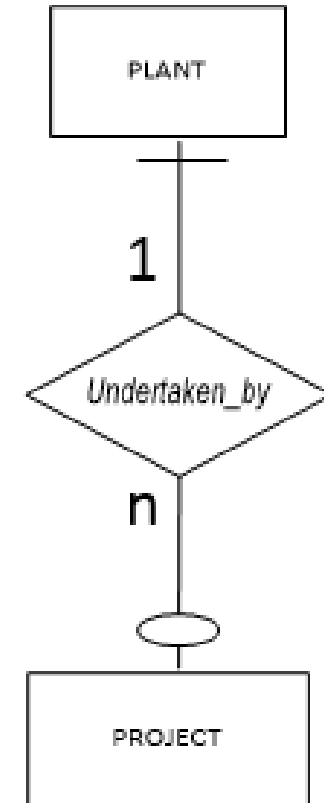
- Key constraints (superkey and candidate key) and entity integrity constraint (primary key) pertain to individual relation schemas
- Referential integrity constraints are specified between two relation schemas (i.e., R1 and R2)
- Specifically, a referential integrity constraint is specified between two relations in order to maintain consistency across tuples of the two relations
- Informal definition: A tuple in one relation that refers to another relation must refer to an existing tuple in that relation.

# Foreign Key Constraint

- A special form of referential integrity constraint specification
- Establishes an explicit association between two relation schemas and maintains the integrity of such an association
- Foreign key: An attribute(s) set, A2, in a relation schema R2 that shares the same domain with a **candidate key** (A1) of another relation schema R1; A2 is said to reference or refer to the relation schema R1.
  - Note: R2 is known as the referencing relation and R1 is called the referenced relation. The attribute(s) doing the referencing (A2 in R2) is the foreign key, while the candidate key (**advisably, the primary key**) being referenced (A1 in R1) is the referenced attribute(s).
- Referred to as Inclusion Dependency, this constraint is algebraically expressed as:  
 **$R2.\{A2\} \subseteq R1.\{A1\}$**
- Meaning: Child.foreignkey (R2.{A2}) is inclusion dependent on parent.primarykey (R1.{A1})

# Example – Source Schema

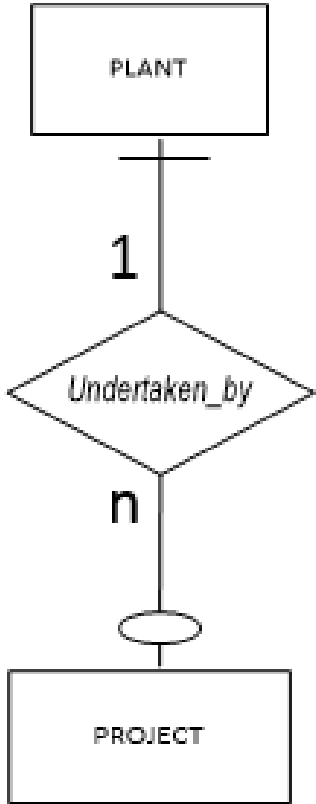
- Bearcat Incorporated is a manufacturing company that has several plants in the northeastern part of the United States. These plants are responsible for leading different projects that the company might undertake, depending on a plants' function. A certain plant might even be associated with several projects but a project is always under the control of just one plant. Some plants do not undertake any projects at all.



# FK in the 1:n relationship

| PLANT | Pl_name       | <u>Pl_p#</u> | Pl_budget |
|-------|---------------|--------------|-----------|
|       | Black Horse   | 11           | 1230000   |
|       | Mayde Creek   | 13           | 1930000   |
|       | Whitefield    | 12           | 2910000   |
|       | River Oaks    | 17           | 1930000   |
|       | King's Island | 19           | 2500000   |
|       | Ashton        | 15           | 2500000   |

| PROJECT | <u>Prj_name</u> | Prj_n# | Prj_pl_p# |
|---------|-----------------|--------|-----------|
|         | Solar Heating   | 41     | 11        |
|         | Lunar Cooling   | 17     | 17        |
|         | Synthetic Fuel  | 29     | 17        |
|         | Nitro-Cooling   | 23     | 12        |
|         | Robot Sweeping  | 31     | 11        |
|         | Robot Painting  | 37     | 19        |
|         | Ozone Control   | 13     | 19        |



Foreign Key

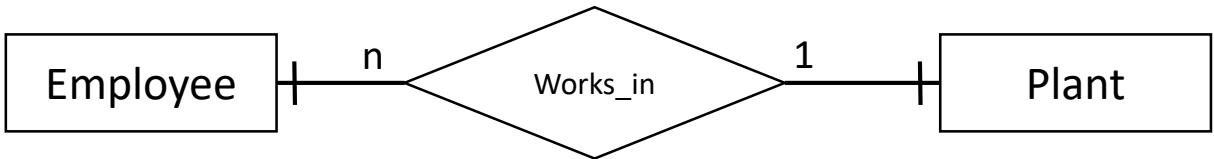
$$R2.\{A2\} \subseteq R1.\{A1\}$$

$$\text{Project.Prj\_pl\_P\#} \subseteq \text{Plant.Pl\_p\#}$$

*Note:* PROJECT.Prj\_pl\_p# is the foreign key referencing PLANT.Pl\_p#, the primary key of PLANT.

**Note: Prj\_n# would be a better primary key for the Project entity. We'll talk about why later, but this works for this example**

# Foreign key placement



- What is the cardinality of the relationship? 1:n
- Which entity is the parent? Plant
- In a 1:M relationship, Foreign key (FK) goes with the “child” side
  - The FK is an attribute in the child entity (employees in this case) that refers to the primary key of the entity on the other side of the relationship

| Employees      |                |       |       |           |         |       |
|----------------|----------------|-------|-------|-----------|---------|-------|
| <u>EmpID_A</u> | <u>EmpID_N</u> | Fname | Minit | Lname     | NameTag | PL_No |
| E              | 12345          | Adam  | B     | Christie  |         | 10    |
| E              | 22233          | Danny | E     | Francisco |         | 10    |
| C              | 32112          | Greg  |       | Hernandez | 1       | 10    |
| E              | 43210          | Ivana | J     | Klink     |         | 11    |
| E              | 52525          | Greg  |       | Hernandez | 2       | 11    |

| Plants       |            |         |
|--------------|------------|---------|
| <u>PL_No</u> | PL_Name    | Budget  |
| 10           | Underwood  | 3000000 |
| 11           | Garnett    | 3000000 |
| 12           | Belmont    | 3500000 |
| 13           | Vanderbilt | 3500000 |

- Would this even make sense the other way?

Note this is dataset is truncated. In reality every entity in the plant table must have at least 100 employees (as specified in the business rule – the ERD above merely specifies that it must participate)

Note: Employees.PL\_No is a FK that refers to Plants.PL\_No. All values of Employees.PL\_No must be found in the domain of values for Plants.PL\_No (i.e., must be a valid plant)

# In 1:m The FK must go with the child entity

- If we put the FK with the parent it is INCORRECT:
  - Lots of data redundancy in the Plants table if the FK is there

| Employees |         |       |       |           |         |
|-----------|---------|-------|-------|-----------|---------|
| EmpID_A   | EmpID_N | Fname | Minit | Lname     | NameTag |
| E         | 12345   | Adam  | B     | Christie  |         |
| E         | 22233   | Danny | E     | Francisco |         |
| C         | 32112   | Greg  |       | Hernande  | 1       |
| E         | 43210   | Ivana | J     | Klink     |         |
| E         | 52525   | Greg  |       | Hernande  | 2       |

| Plants |            |         |        |
|--------|------------|---------|--------|
| PL_No  | PL_Name    | Budget  | EmpID  |
| 10     | Underwood  | 3000000 | E12345 |
| 10     | Underwood  | 3000000 | E22233 |
| 10     | Underwood  | 3000000 | C32112 |
| 11     | Garnett    | 3000000 | E43210 |
| 11     | Garnett    | 3000000 | E52525 |
| 12     | Belmont    | 3500000 | NULL   |
| 13     | Vanderbilt | 3500000 | NULL   |

Note: In this case, Plants.EmpID is a FK that refers to the composite attribute Employees.EmpID\_A + Employees.EmpID\_N.

There is now redundancy for the Plants.PL\_Name and Plants.Budget attributes, and PL\_No is no longer unique!

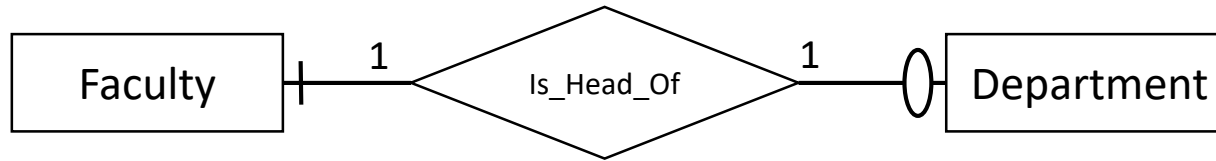
- No redundancy when FK is with Employees (CORRECT):

| Employees |         |       |       |           |         |
|-----------|---------|-------|-------|-----------|---------|
| EmpID_A   | EmpID_N | Fname | Minit | Lname     | NameTag |
| E         | 12345   | Adam  | B     | Christie  |         |
| E         | 22233   | Danny | E     | Francisco |         |
| C         | 32112   | Greg  |       | Hernandez | 1       |
| E         | 43210   | Ivana | J     | Klink     |         |
| E         | 52525   | Greg  |       | Hernandez | 2       |

| Plants |            |         |
|--------|------------|---------|
| PL_No  | PL_Name    | Budget  |
| 10     | Underwood  | 3000000 |
| 11     | Garnett    | 3000000 |
| 12     | Belmont    | 3500000 |
| 13     | Vanderbilt | 3500000 |

# Where do we put the foreign key in a 1:1 relationship?

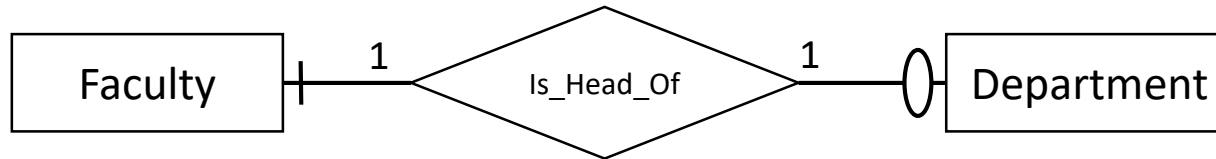
- There is no clear parent....





# In 1:1 we place the FK based on participation

- FK goes with the entity that has total (mandatory) participation



| Faculty |               |                     |
|---------|---------------|---------------------|
| EmpID   | Name          | Title               |
| 112     | Chad Larson   | Assistant Professor |
| 108     | Dusya Vera    | Professor           |
| 109     | Edward Blair  | Professor           |
| 111     | Gerald Lobo   | Professor           |
| 106     | Kris Jacobs   | Professor           |
| 102     | Mark Grimes   | Assistant Professor |
| 110     | Melanie Rudd  | Assistant Professor |
| 103     | Michael Parks | Associate Professor |
| 101     | Norm Johnson  | Professor           |
| 104     | Praveen Kumar | Professor           |
| 107     | Steve Werner  | Professor           |
| 105     | Thomas George | Professor           |

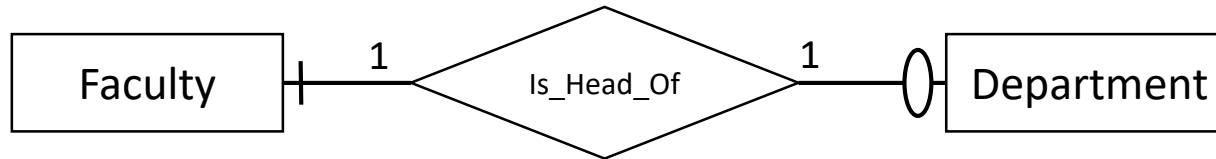
| Departments |            |      |
|-------------|------------|------|
| DeptID      | Deptname   | Head |
| 1           | Accounting | 111  |
| 2           | DISC       | 101  |
| 3           | Finanace   | 104  |
| 4           | Management | 107  |
| 5           | Marketing  | 109  |

Note: In this case, Departments.Head is a FK that refers to Faculty.EmpID.

**CORRECT  
WAY!**

# In 1:1 we place the FK based on participation

- Imagine if we did it the other way, it still works, but lots of NULL values!
  - NULL values are difficult to work with – avoid them as much as possible!



| Faculty      |               |                     |         |
|--------------|---------------|---------------------|---------|
| <u>EmpID</u> | Name          | Title               | Head of |
| 112          | Chad Larson   | Assistant Professor |         |
| 108          | Dusya Vera    | Professor           |         |
| 109          | Edward Blair  | Professor           | 5       |
| 111          | Gerald Lobo   | Professor           | 1       |
| 106          | Kris Jacobs   | Professor           |         |
| 102          | Mark Grimes   | Assistant Professor |         |
| 110          | Melanie Rudd  | Assistant Professor |         |
| 103          | Michael Parks | Associate Professor |         |
| 101          | Norm Johnson  | Professor           | 2       |
| 104          | Praveen Kumar | Professor           | 3       |
| 107          | Steve Werner  | Professor           | 4       |
| 105          | Thomas George | Professor           |         |

| Departments   |            |
|---------------|------------|
| <u>DeptID</u> | Deptname   |
| 1             | Accounting |
| 2             | DISC       |
| 3             | Finanace   |
| 4             | Management |
| 5             | Marketing  |

Note: In this case,  
Faculty.HeadOf is a FK that  
refers to Departments.DeptID.

No data redundancy, but lots of  
nasty NULL values

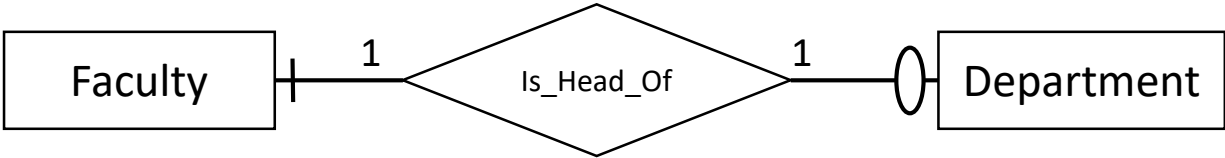
**WRONG  
WAY!**

# In 1:1 we place the FK based on participation

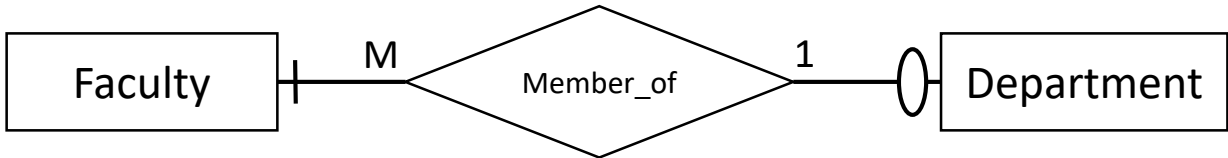
- There are some fringe cases...
- If both sides of a 1:1 have mandatory participation
  - It doesn't matter where you place the FK
  - This would be a pretty rare thing to find
- If both sides of a 1:1 have optional participation
  - No hard and fast rule, but generally would be better to put it with the entity that will have fewer instances
  - Alternatively, you can use a gerund to ensure you will have no NULL values!

# Watch out!

- The “is head of” relationships is different than the “member of” relationship:



- Where should the FK go for the “Member of” relationship? **Faculty (The Child)**



Note: In this case, Faculty.Department is a FK that refers to Departments.DeptID - this is the “Member\_of” relationship.

Departments.Head is a FK that refers to Faculty.EmpID - this is the “Is\_Head\_of” relationship.

We have no redundancy and no NULL values with this design

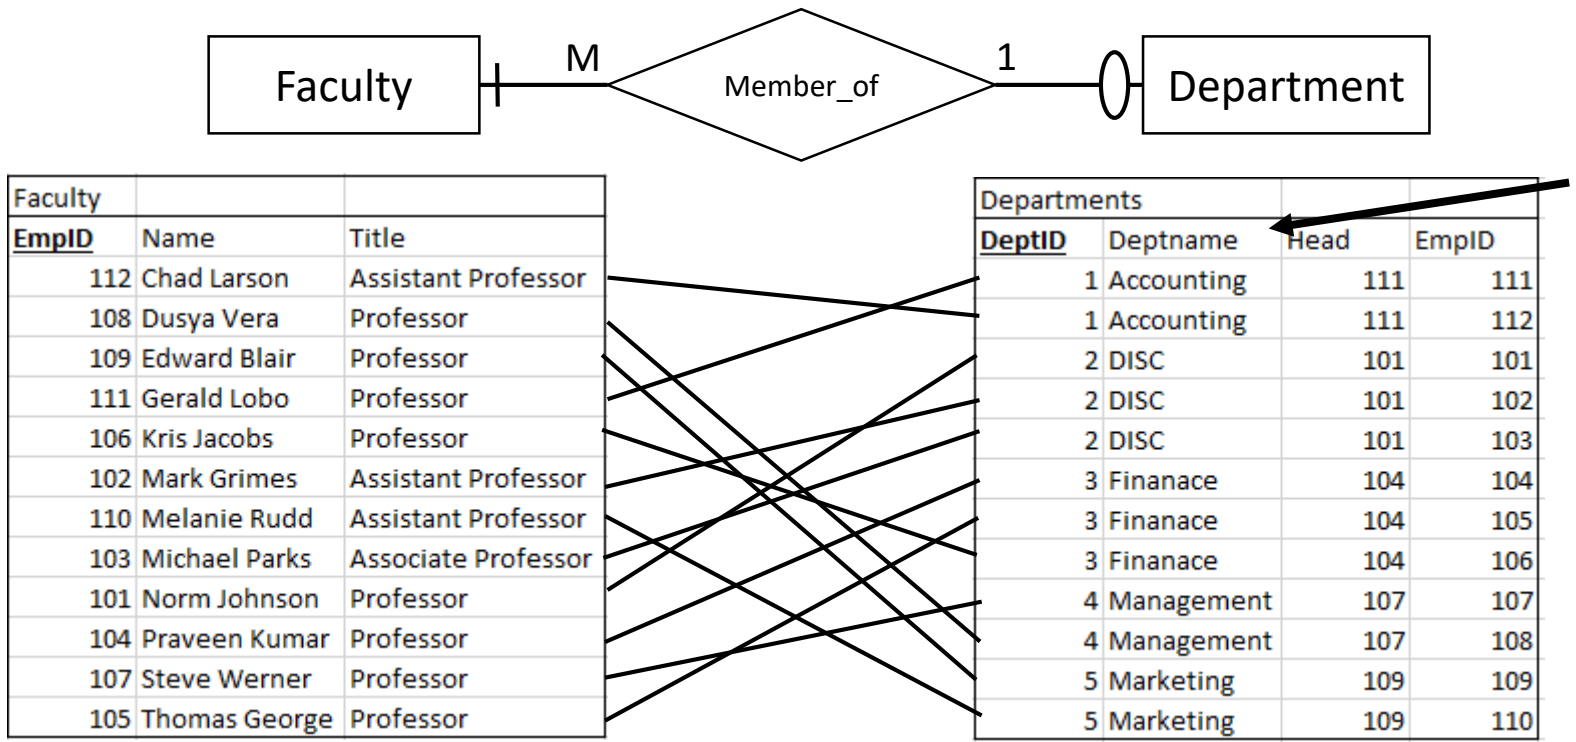
| Faculty |               |                     |            |
|---------|---------------|---------------------|------------|
| EmpID   | Name          | Title               | Department |
| 112     | Chad Larson   | Assistant Professor | 1          |
| 108     | Dusya Vera    | Professor           | 4          |
| 109     | Edward Blair  | Professor           | 5          |
| 111     | Gerald Lobo   | Professor           | 1          |
| 106     | Kris Jacobs   | Professor           | 3          |
| 102     | Mark Grimes   | Assistant Professor | 2          |
| 110     | Melanie Rudd  | Assistant Professor | 5          |
| 103     | Michael Parks | Associate Professor | 2          |
| 101     | Norm Johnson  | Professor           | 2          |
| 104     | Praveen Kumar | Professor           | 3          |
| 107     | Steve Werner  | Professor           | 4          |
| 105     | Thomas George | Professor           | 3          |

| Departments |            |      |
|-------------|------------|------|
| DeptID      | Deptname   | Head |
| 1           | Accounting | 111  |
| 2           | DISC       | 101  |
| 3           | Finanace   | 104  |
| 4           | Management | 107  |
| 5           | Marketing  | 109  |

CORRECT  
WAY!

# Watch out!

- What if we did it the wrong way?



Redundant data for the Deptname and Head attributes **AND** DeptID is no longer unique!

**WRONG WAY!**

# Module 6.3

## Data Integrity Constraints

- Be able to define: Super keys, candidate keys, key attributes, non-key attributes, primary keys, and alternate keys
- What is the entity integrity constraint?
- What is the referential integrity constraint?
- What is the foreign key constraint?

# Module 6.4

## Relational Algebra

- What are the eight operations in relational algebra?
- Describe the difference in select and project
- Describe the difference in the Union, Intersection, and Difference operations
- What is a natural join?

# Relational algebra primer

- Relational algebra just lets us abstractly talk about relations and data (like algebra lets us abstractly talk about numbers)
- We can substitute symbols to represent data and do abstract calculations just like in traditional algebra
  - Area of a room is  $L \times W$ 
    - $L = 10 \quad W = 20 \quad A = L \times W = ? \rightarrow A = 10 \times 20 = 200$
  - If we want a room that is 4 times as large, we can nest our equations
    - $4(A) \rightarrow 4(L \times W) \rightarrow 4(10 \times 20) \rightarrow 4(200) \rightarrow 800$
- For databases
  - We consider a relation (R) as being made up of attributes  $A_1, A_2, A_3, \dots, A_n$
  - The relation schema is  $R(A_1, A_2, A_3, \dots, A_n)$



# Relational algebra primer

- An attribute is defined as an ordered set  $(N, D)$ 
  - $N$  is the name of the attribute
  - $D$  is the domain of the attribute
- A set of attributes can be represented as a vector  $(C)$ 
  - $C$  is the set  $\{(N_1, D_1), (N_2, D_2), (N_3, D_3), \dots, (N_n, D_n)\}$

# Relational algebra primer

- So.... A relation is a set  $(R,C)$  where  $R$  is the name of the relation schema and  $C$  is the list of attributes that make up  $R$
- Your database is made up of multiple relations  $(R,C)$
- Each  $(R,C)$  is equivalent to...
  - $R(A_1, A_2, A_3, \dots, A_n)$  which is equivalent to ...
  - $R(\{(N_1, D_1), (N_2, D_2), (N_3, D_3), \dots, (N_n, D_n)\})$

# Relational algebra primer

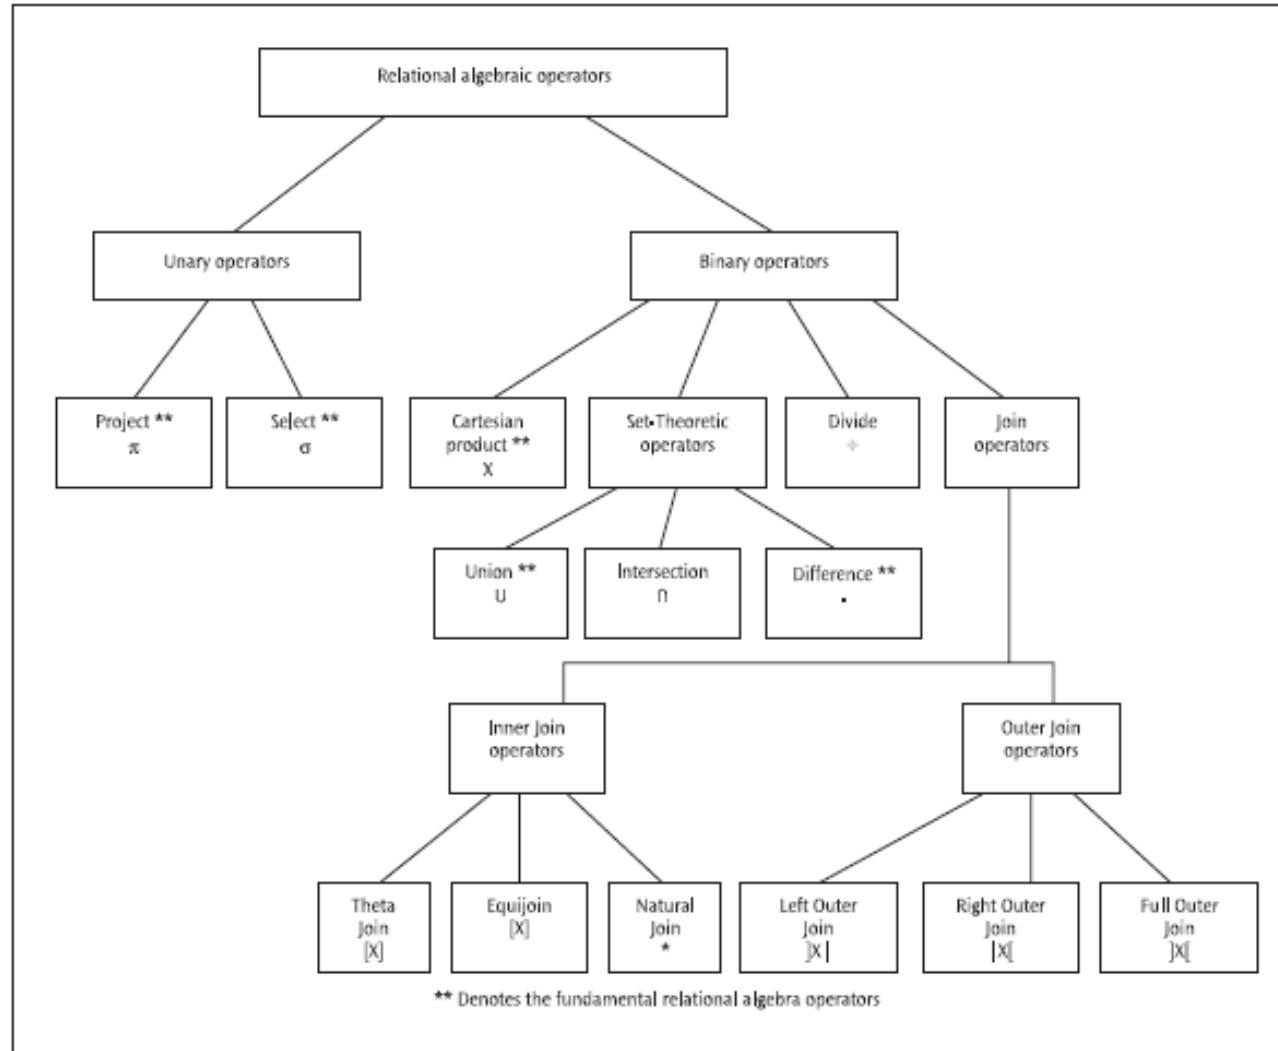


Figure 11.2 Classification of relational algebra operators

# Eight basic functions of relational algebra

- Two unary operations
  - Selection ( $\sigma$ )
  - Projection ( $\pi$ )
- Six binary operators
  - Union ( $\cup$ )
  - Intersection ( $\cap$ )
  - Difference ( $-$ )
  - Join ( $\bowtie$ )
  - Cartesian product ( $\times$ )
  - Division ( $\div$ )

(We'll talk about Cartesian product and division after spring break)

# Imagine these three relations

- Award Winning Plants

| AW_PLANT        |               |              |
|-----------------|---------------|--------------|
| <u>Aw_pl No</u> | Aw_pl_Name    | Aw_pl_Budget |
| 11              | Black Horse   | 2500000      |
| 13              | Mayde Creek   | 1930000      |
| 12              | Whitefield    | 2910000      |
| 17              | River Oaks    | 1930000      |
| 19              | King's Island | 2500000      |
| 15              | Ashton        | 2500000      |

- Texas Plants

| TX_PLANT        |               |              |
|-----------------|---------------|--------------|
| <u>Tx_pl No</u> | Tx_pl_Name    | Tx_pl_Budget |
| 16              | Southern Oaks | 1930000      |
| 17              | River Oaks    | 1930000      |
| 18              | Kingwood      | 1930000      |

- Projects

| PROJECTS        |        |              |              |
|-----------------|--------|--------------|--------------|
| <u>Prj_name</u> | Prj_no | Prj_location | Prj_aw_pl_no |
| Solar Heating   | 41     | Sealy        | 11           |
| Lunar Cooling   | 17     | Yoakum       | 17           |
| Synthetic Fuel  | 29     | Salem        | 17           |
| Nitro-Cooling   | 23     | Parthi       | 12           |
| Robot Sweeping  | 31     | Ponca City   | 11           |
| Robot Painting  | 37     | Poakum       | 19           |
| Ozone Control   | 13     | Parthi       | 19           |

# Unary Operations

## Selection ( $\sigma$ )

|  |  |  |  |  |  |
|--|--|--|--|--|--|
|  |  |  |  |  |  |
|  |  |  |  |  |  |
|  |  |  |  |  |  |
|  |  |  |  |  |  |
|  |  |  |  |  |  |

## Projection ( $\pi$ )

[illegible]

# Select Operator

- Selects a horizontal subset of tuples that satisfy a selection condition from the relation

$$\sigma_{\langle \text{selection condition} \rangle} R$$

- Lower case sigma ( $\sigma$ ) designates “select”
- $\langle \text{selection condition} \rangle$  is a Boolean expression specified on the attributes of relation R

# Selection (Unary Operation)

- Creates a second relation by extracting a subset of tuples
- Question: Which award winning plants have a budget over \$2,000,000
- $\sigma_{Aw\_pl\_Budget > 2000000} AW\_Plant$
- `SELECT * FROM AW_Plant WHERE Aw_pl_Budget > 2000000`

| AW_PLANT |               |              |
|----------|---------------|--------------|
| Aw_pl_No | Aw_pl_Name    | Aw_pl_Budget |
| 11       | Black Horse   | 2500000      |
| 13       | Mayde Creek   | 1930000      |
| 12       | Whitefield    | 2910000      |
| 17       | River Oaks    | 1930000      |
| 19       | King's Island | 2500000      |
| 15       | Ashton        | 2500000      |



| AW_PLANT |               |              |
|----------|---------------|--------------|
| Aw_pl_No | Aw_pl_Name    | Aw_pl_Budget |
| 11       | Black Horse   | 2500000      |
| 12       | Whitefield    | 2910000      |
| 19       | King's Island | 2500000      |
| 15       | Ashton        | 2500000      |



# Projection Operator

- Selects a vertical subset of attributes from a relation

$$\pi_{\langle \text{attribute list} \rangle} R$$

- Lower case pi ( $\pi$ ) designates “project”
- $\langle \text{attribute list} \rangle$  is a subset of attributes of relation R

# Projection (Unary Operation)

- Creates a second relation by extracting a subset of columns
- Question: What is the plant number and budget for each of the award winning plants?
- $\pi_{Aw\_pl\_No, Aw\_pl\_Budget} AW\_Plant$
- `SELECT Aw_pl_No, Aw_pl_Budget FROM AW_Plant`

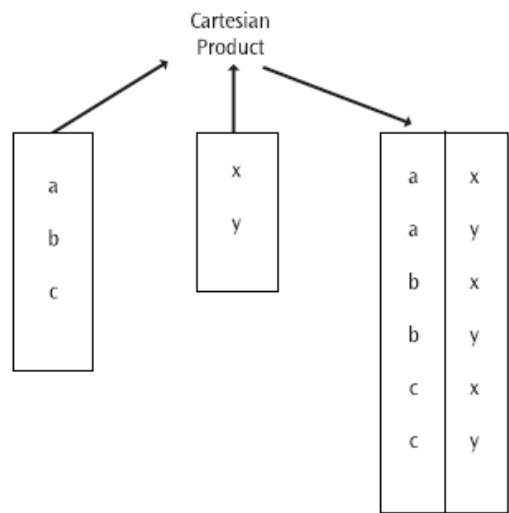
| AW_PLANT |               |              |
|----------|---------------|--------------|
| Aw pl No | Aw_pl_Name    | Aw_pl_Budget |
| 11       | Black Horse   | 2500000      |
| 13       | Mayde Creek   | 1930000      |
| 12       | Whitefield    | 2910000      |
| 17       | River Oaks    | 1930000      |
| 19       | King's Island | 2500000      |
| 15       | Ashton        | 2500000      |



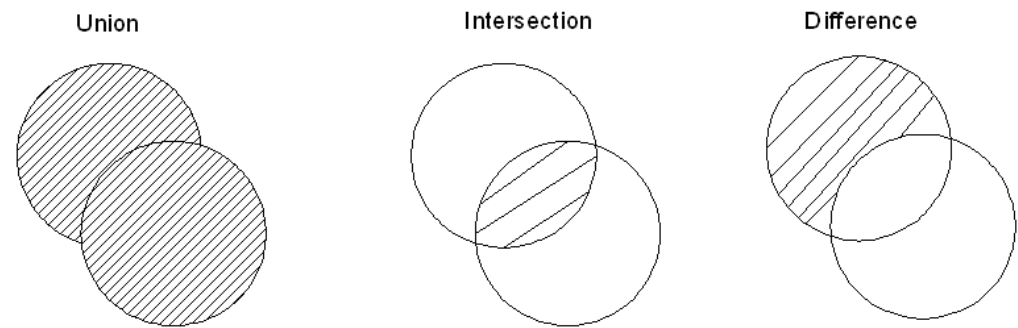
| AW_PLANT |              |
|----------|--------------|
| Aw pl No | Aw_pl_Budget |
| 11       | 2500000      |
| 13       | 1930000      |
| 12       | 2910000      |
| 17       | 1930000      |
| 19       | 2500000      |
| 15       | 2500000      |

# Binary Operators

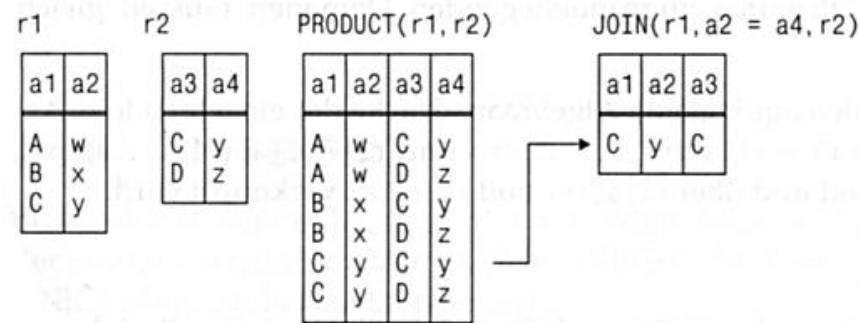
## Cartesian product



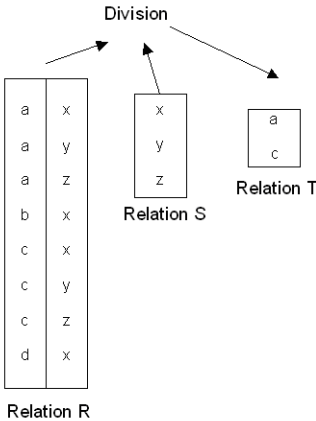
## Union ( $\cup$ ), Intersection ( $\cap$ ), and Difference ( $-$ )



## Join

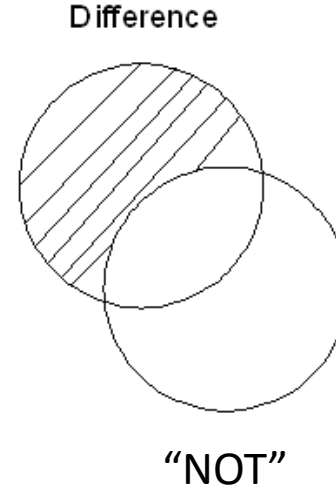
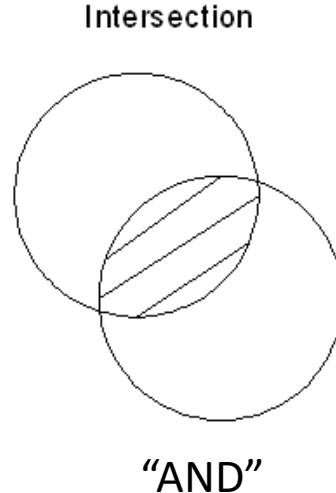
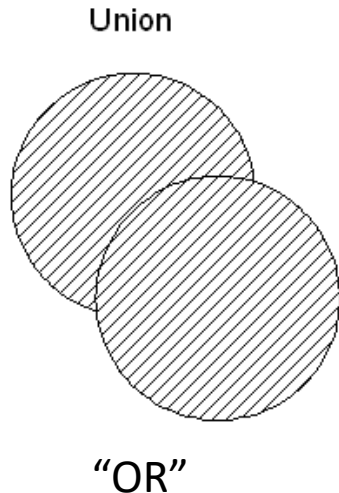


## Division



# Set Operations

- Two relations (R and S) are “union compatible” if they:
  - Have the same degree (number of attributes)
  - Pairs of attributes from R and S have the same domain
- Union compatibility is a requirement for all set operations (union, intersection, and difference)

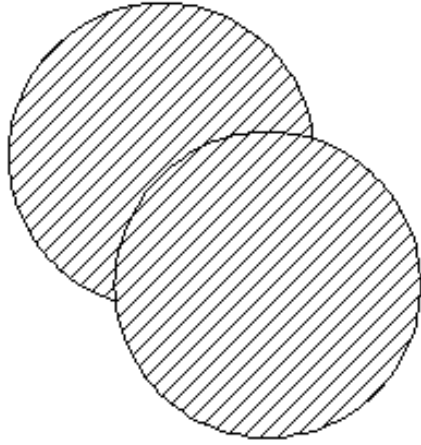


# Set Theory and Relational Algebra

- Database theory is based on set theory
- Manipulations referred to as “Relational Algebra”

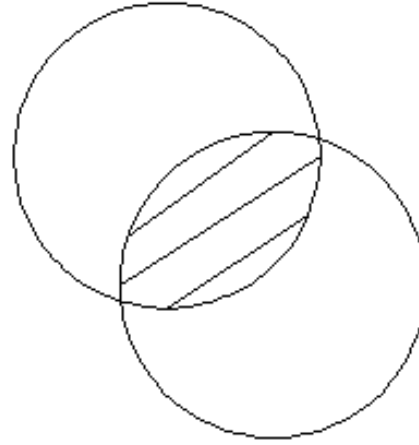
# Set Theory Operators

Union



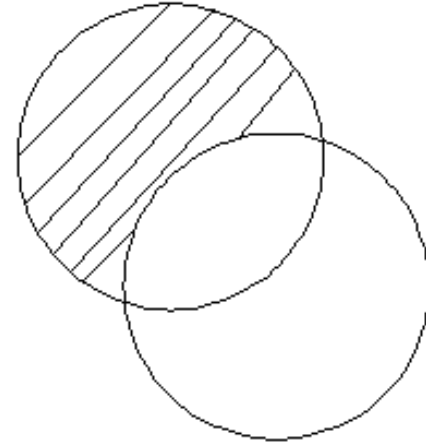
“OR”

Intersection



“AND”

Difference

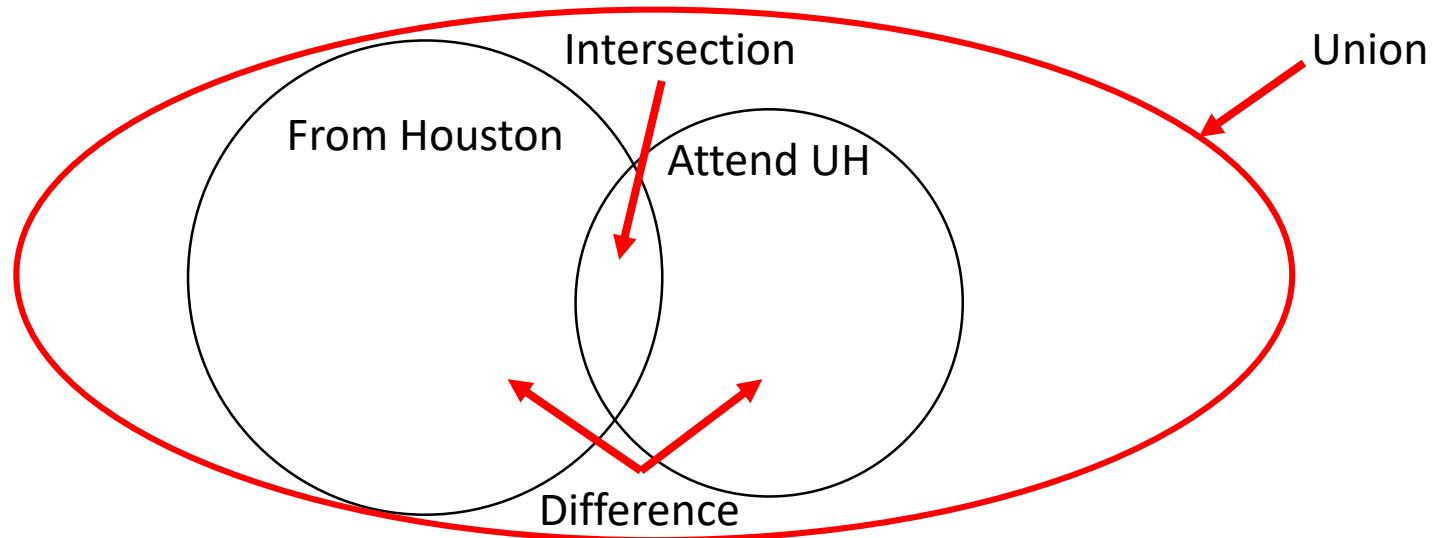


“NOT”

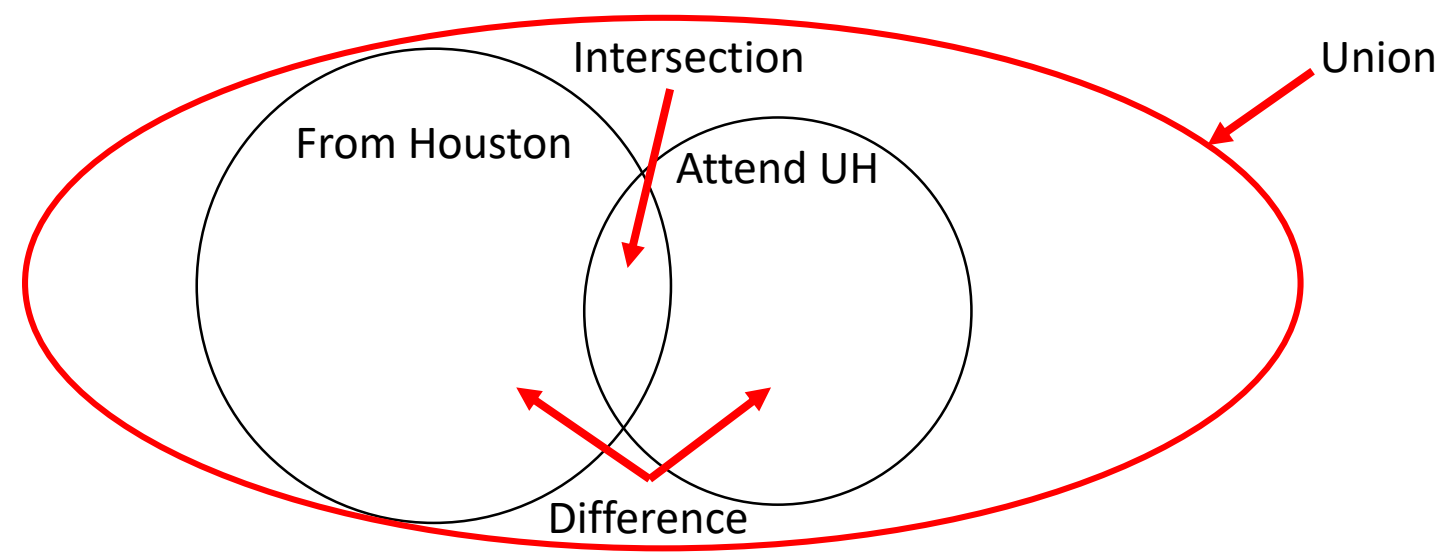
# Set Theory Operators

- Some people are from Houston
- Some people attend UH
- Union: People that are from Houston OR go to UH
- Intersection: People that are from Houston AND go to UH
- Difference: People from Houston but do NOT go to UH

People that go to UH but NOT from Houston



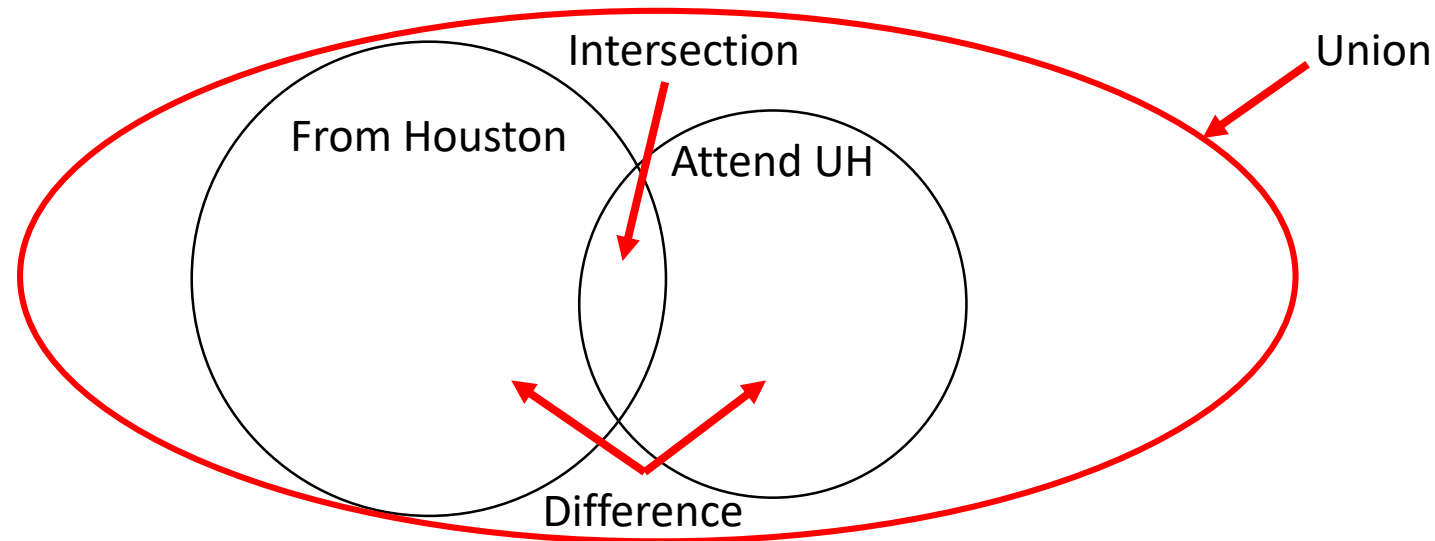
# Union of Houston and UH



| <u>StuID</u> | Name   | Hometown | School |
|--------------|--------|----------|--------|
| 1111         | Adam   | Houston  | UH     |
| 2222         | Beth   | Dallas   | UH     |
| 3333         | Chris  | Houston  | UH     |
| 4444         | Dave   | Austin   | UH     |
| 5555         | Eunice | Houston  | Rice   |
| 6666         | Frank  | Houston  | Rice   |

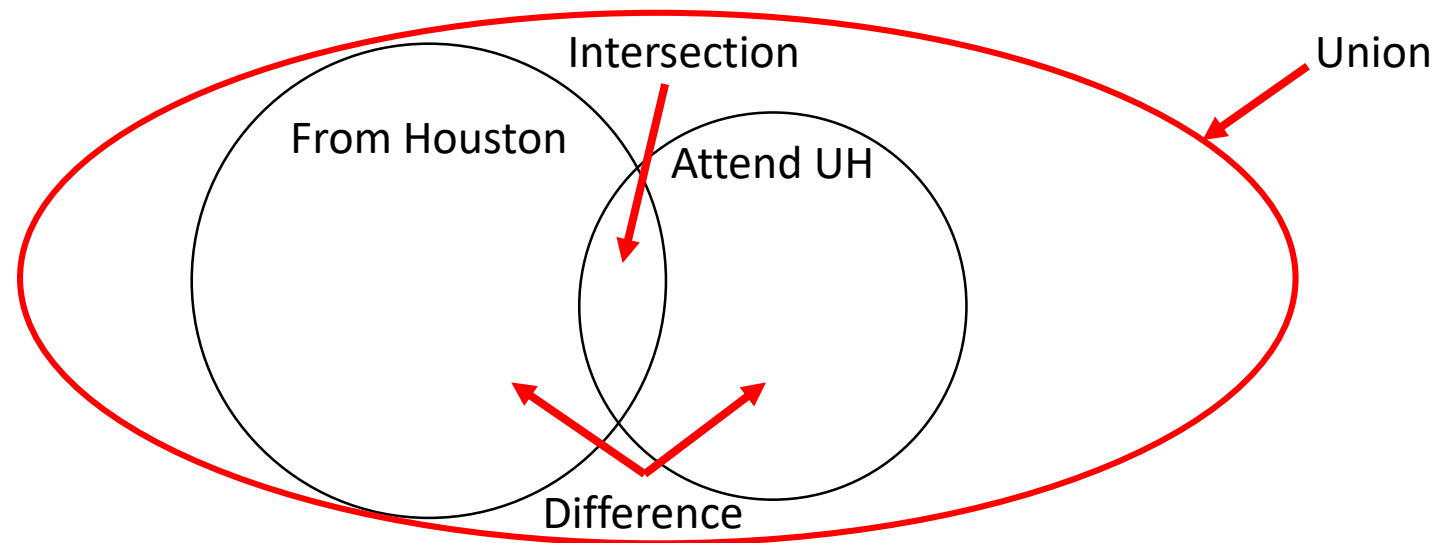


# Intersection of Houston and UH



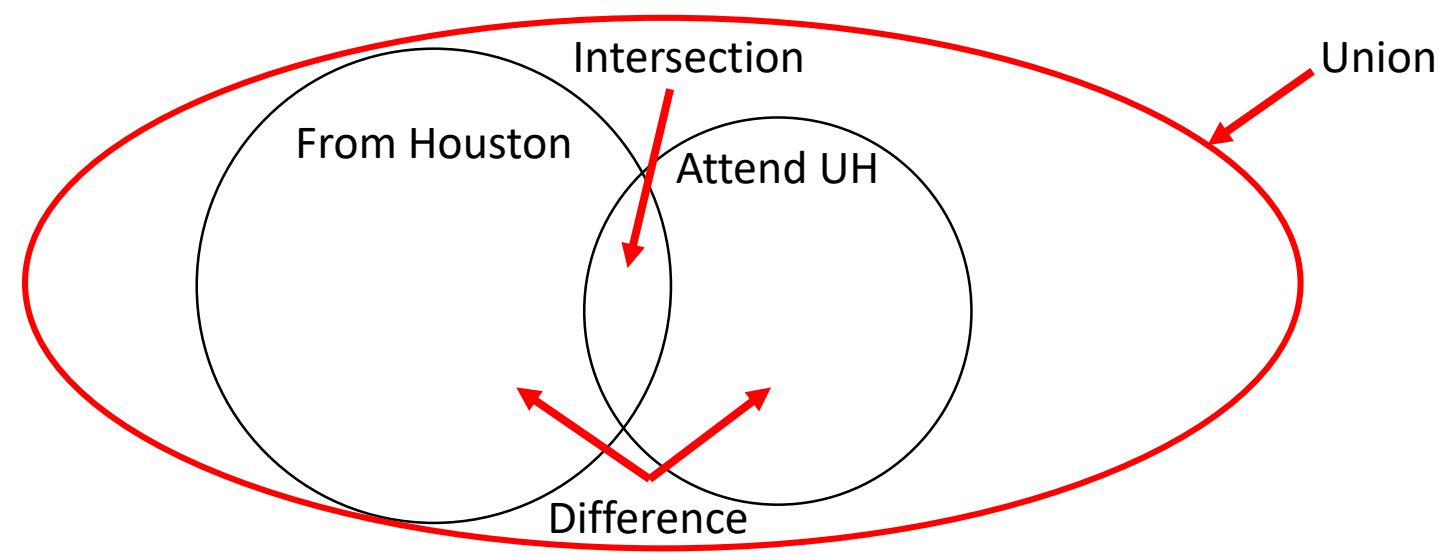
| <u>StuID</u> | Name   | Hometown | School |
|--------------|--------|----------|--------|
| 1111         | Adam   | Houston  | UH     |
| 2222         | Beth   | Dallas   | UH     |
| 3333         | Chris  | Houston  | UH     |
| 4444         | Dave   | Austin   | UH     |
| 5555         | Eunice | Houston  | Rice   |
| 6666         | Frank  | Houston  | Rice   |

# UH minus Houston



| <u>StuID</u> | Name   | Hometown | School |
|--------------|--------|----------|--------|
| 1111         | Adam   | Houston  | UH     |
| 2222         | Beth   | Dallas   | UH     |
| 3333         | Chris  | Houston  | UH     |
| 4444         | Dave   | Austin   | UH     |
| 5555         | Eunice | Houston  | Rice   |
| 6666         | Frank  | Houston  | Rice   |

# Houston minus UH



| <u>StuID</u> | Name   | Hometown | School |
|--------------|--------|----------|--------|
| 1111         | Adam   | Houston  | UH     |
| 2222         | Beth   | Dallas   | UH     |
| 3333         | Chris  | Houston  | UH     |
| 4444         | Dave   | Austin   | UH     |
| 5555         | Eunice | Houston  | Rice   |
| 6666         | Frank  | Houston  | Rice   |

# Union (Binary Operation)

- Creates a third relation containing tuples from either relation
- What plants are either in Texas OR are award winning plants?
- $AW\_Plant \cup TX\_Plant$
- `SELECT * FROM AW_Plant UNION SELECT * FROM TX_Plant`

| AW_PLANT |               |              |
|----------|---------------|--------------|
| Aw_pl_No | Aw_pl_Name    | Aw_pl_Budget |
| 11       | Black Horse   | 2500000      |
| 13       | Mayde Creek   | 1930000      |
| 12       | Whitefield    | 2910000      |
| 17       | River Oaks    | 1930000      |
| 19       | King's Island | 2500000      |
| 15       | Ashton        | 2500000      |

OR

| TX_PLANT |               |              |
|----------|---------------|--------------|
| Tx_pl_No | Tx_pl_Name    | Tx_pl_Budget |
| 16       | Southern Oaks | 1930000      |
| 17       | River Oaks    | 1930000      |
| 18       | Kingwood      | 1930000      |



| R_AW_PLANT |               |              |
|------------|---------------|--------------|
| Aw_pl_No   | Aw_pl_Name    | Aw_pl_Budget |
| 11         | Black Horse   | 2500000      |
| 13         | Mayde Creek   | 1930000      |
| 12         | Whitefield    | 2910000      |
| 17         | River Oaks    | 1930000      |
| 19         | King's Island | 2500000      |
| 15         | Ashton        | 2500000      |
| 16         | Southern Oaks | 1930000      |
| 18         | Kingwood      | 1930000      |

# Intersection (Binary Operation)

- Creates a third relation containing tuples present in both relations
- What plants are in Texas AND are award winning plants?
- $AW\_Plant \cap TX\_Plant$
- `SELECT * FROM AW_Plant INTERSECT SELECT * FROM TX_Plant`

| AW_PLANT |               |              |
|----------|---------------|--------------|
| Aw_pl No | Aw_pl_Name    | Aw_pl_Budget |
| 11       | Black Horse   | 2500000      |
| 13       | Mayde Creek   | 1930000      |
| 12       | Whitefield    | 2910000      |
| 17       | River Oaks    | 1930000      |
| 19       | King's Island | 2500000      |
| 15       | Ashton        | 2500000      |

AND

| TX_PLANT |               |              |
|----------|---------------|--------------|
| Tx_pl No | Tx_pl_Name    | Tx_pl_Budget |
| 16       | Southern Oaks | 1930000      |
| 17       | River Oaks    | 1930000      |
| 18       | Kingwood      | 1930000      |



| R_AW_PLANT |            |              |
|------------|------------|--------------|
| Aw_pl No   | Aw_pl_Name | Aw_pl_Budget |
| 17         | River Oaks | 1930000      |

# Difference (Binary Operation)

- Creates a third relation containing tuples present in one relation but not the other relation
- What plants are in Texas, but are not award winning?
- TX\_Plant – AW\_Plant
- `SELECT * FROM TX_Plant MINUS SELECT * FROM AW_Plant`

| TX_PLANT        |               |              |
|-----------------|---------------|--------------|
| <u>Tx_pl No</u> | Tx_pl_Name    | Tx_pl_Budget |
| 16              | Southern Oaks | 1930000      |
| 17              | River Oaks    | 1930000      |
| 18              | Kingwood      | 1930000      |

NOT

| AW_PLANT        |               |              |
|-----------------|---------------|--------------|
| <u>Aw_pl No</u> | Aw_pl_Name    | Aw_pl_Budget |
| 11              | Black Horse   | 2500000      |
| 13              | Mayde Creek   | 1930000      |
| 12              | Whitefield    | 2910000      |
| 17              | River Oaks    | 1930000      |
| 19              | King's Island | 2500000      |
| 15              | Ashton        | 2500000      |



| TX_PLANT        |               |              |
|-----------------|---------------|--------------|
| <u>Tx_pl No</u> | Tx_pl_Name    | Tx_pl_Budget |
| 16              | Southern Oaks | 1930000      |
| 18              | Kingwood      | 1930000      |

# Difference (Binary Operation)

- Creates a third relation containing tuples present in one relation but not the other relation
- What plants are award winning but are not in Texas?
- $AW\_Plant - TX\_Plant$
- `SELECT * FROM AW_Plant MINUS SELECT * FROM TX_Plant`

| AW_PLANT |               |              |
|----------|---------------|--------------|
| Aw_pl_No | Aw_pl_Name    | Aw_pl_Budget |
| 11       | Black Horse   | 2500000      |
| 13       | Mayde Creek   | 1930000      |
| 12       | Whitefield    | 2910000      |
| 17       | River Oaks    | 1930000      |
| 19       | King's Island | 2500000      |
| 15       | Ashton        | 2500000      |

Not



| R_AW_PLANT |               |              |
|------------|---------------|--------------|
| Aw_pl_No   | Aw_pl_Name    | Aw_pl_Budget |
| 11         | Black Horse   | 2500000      |
| 13         | Mayde Creek   | 1930000      |
| 12         | Whitefield    | 2910000      |
| 19         | King's Island | 2500000      |
| 15         | Ashton        | 2500000      |

| TX_PLANT |               |              |
|----------|---------------|--------------|
| Tx_pl_No | Tx_pl_Name    | Tx_pl_Budget |
| 16       | Southern Oaks | 1930000      |
| 17       | River Oaks    | 1930000      |
| 18       | Kingwood      | 1930000      |

# Natural Join

- Combines two relations into a third by matching values that come from the same domain for attributes in the two relations.
- What award winning plant is each project assigned to?
- `SELECT * FROM AW_Plant NATURAL JOIN Projects`
- `SELECT * FROM AW_Plant INNER JOIN Projects ON Aw_pl_No = Prj_aw_pl_no`

| AW_PLANT |               |              |
|----------|---------------|--------------|
| Aw_pl_No | Aw_pl_Name    | Aw_pl_Budget |
| 11       | Black Horse   | 2500000      |
| 13       | Mayde Creek   | 1930000      |
| 12       | Whitefield    | 2910000      |
| 17       | River Oaks    | 1930000      |
| 19       | King's Island | 2500000      |
| 15       | Ashton        | 2500000      |

| PROJECTS       |        |              |              |
|----------------|--------|--------------|--------------|
| Prj_name       | Prj_no | Prj_location | Prj_aw_pl_no |
| Solar Heating  | 41     | Sealy        | 11           |
| Lunar Cooling  | 17     | Yoakum       | 17           |
| Synthetic Fuel | 29     | Salem        | 17           |
| Nitro-Cooling  | 23     | Parthi       | 12           |
| Robot Sweeping | 31     | Ponca City   | 11           |
| Robot Painting | 37     | Poakum       | 19           |
| Ozone Control  | 13     | Parthi       | 19           |



| PROJECTS       |        |              |              |               |              |
|----------------|--------|--------------|--------------|---------------|--------------|
| Prj_name       | Prj_no | Prj_location | Prj_aw_pl_no | Prj_pl_Name   | Aw_pl_Budget |
| Solar Heating  | 41     | Sealy        | 11           | Black Horse   | 2500000      |
| Lunar Cooling  | 17     | Yoakum       | 17           | River Oaks    | 1930000      |
| Synthetic Fuel | 29     | Salem        | 17           | River Oaks    | 1930000      |
| Nitro-Cooling  | 23     | Parthi       | 12           | Whitefield    | 2910000      |
| Robot Sweeping | 31     | Ponca City   | 11           | Black Horse   | 2500000      |
| Robot Painting | 37     | Poakum       | 19           | King's Island | 2500000      |
| Ozone Control  | 13     | Parthi       | 19           | King's Island | 2500000      |



# Module 6.4

## Relational Algebra

- What are the eight operations in relational algebra?
- Describe the difference in select and project
- Describe the difference in the Union, Intersection, and Difference operations
- What is a natural join?

# Module 6.5

## Views and Materialized Views

- What do views do and why are they useful?
- How do growth and restructuring impact views?
- What is the difference in a normal view and a materialized view?

# Views

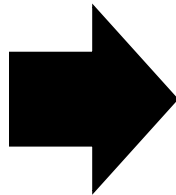
- Last time we talked about different types of joins and projections – these are “views” of the data
- Unlike a relation schema, a view does not contain data
  - Is a logical window to the attributes and tuples from one or a set of relations
- Views provide many benefits:
  - Allows the same data to be seen in different ways by different users (or applications)
  - Provides security by restricting access to data
  - Hides complexity by making data from several relations appear as a single object

# Horse view - Selection

- For a doctor that only needs to see female horses:
  - `SELECT * FROM Horses WHERE Gender='F'`

Relation

| Horses |        |       |        |        |
|--------|--------|-------|--------|--------|
| Name   | Color  | Spots | Gender | Weight |
| Amy    | Yellow | Yes   | F      | 920    |
| Dave   | Grey   | No    | M      | 1800   |
| Ed     | Yellow | No    | M      | 1100   |
| Sally  | Black  | No    | F      | 1200   |
| Sarah  | Grey   | No    | F      | 1700   |
| Tom    | Red    | Yes   | M      | 1500   |
| Jane   | Red    | No    | F      | 1300   |



View

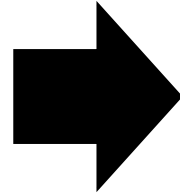
| Horses |        |       |        |        |
|--------|--------|-------|--------|--------|
| Name   | Color  | Spots | Gender | Weight |
| Amy    | Yellow | Yes   | F      | 920    |
| Sally  | Black  | No    | F      | 1200   |
| Sarah  | Grey   | No    | F      | 1700   |
| Jane   | Red    | No    | F      | 1300   |

# Horse View - Projection

- For someone allocating stables who only needs to know the name and weight of the horses:
  - `SELECT Name, Weight FROM Horses`

Relation

| Horses |        |       |        |        |
|--------|--------|-------|--------|--------|
| Name   | Color  | Spots | Gender | Weight |
| Amy    | Yellow | Yes   | F      | 920    |
| Dave   | Grey   | No    | M      | 1800   |
| Ed     | Yellow | No    | M      | 1100   |
| Sally  | Black  | No    | F      | 1200   |
| Sarah  | Grey   | No    | F      | 1700   |
| Tom    | Red    | Yes   | M      | 1500   |
| Jane   | Red    | No    | F      | 1300   |



View

| Horses |        |
|--------|--------|
| Name   | Weight |
| Amy    | 920    |
| Dave   | 1800   |
| Ed     | 1100   |
| Sally  | 1200   |
| Sarah  | 1700   |
| Tom    | 1500   |
| Jane   | 1300   |

# Views

- We're showing users (or applications) only the data they need to see
  - Provides security by restricting access to data
  - Hides complexity by making data from several relations appear as a single object
  - Allows the same data to be seen in different ways by different users

- What does this sound like?
  - External Schema!

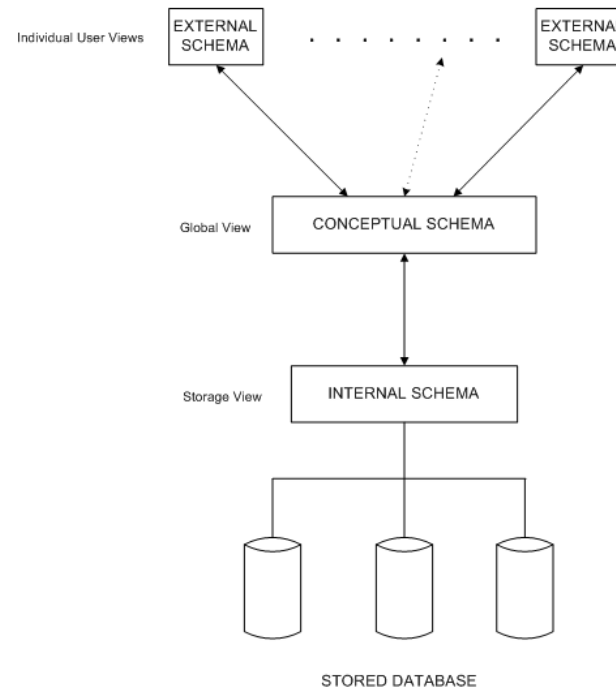


Figure 1.2 The ANSI/SPARC three-schema Architecture

# Growth and Restructuring

- By using views, changes to the database structure may be made without affecting users

- What does this sound like?

- Two primary types of changes:

- **Growth:** Adding new attributes to a relation, or new relations to a data model

- **Restructuring:** Changing the conceptual schema

- Should be “Information Equivalent” – which facilitates the restructuring being reversible

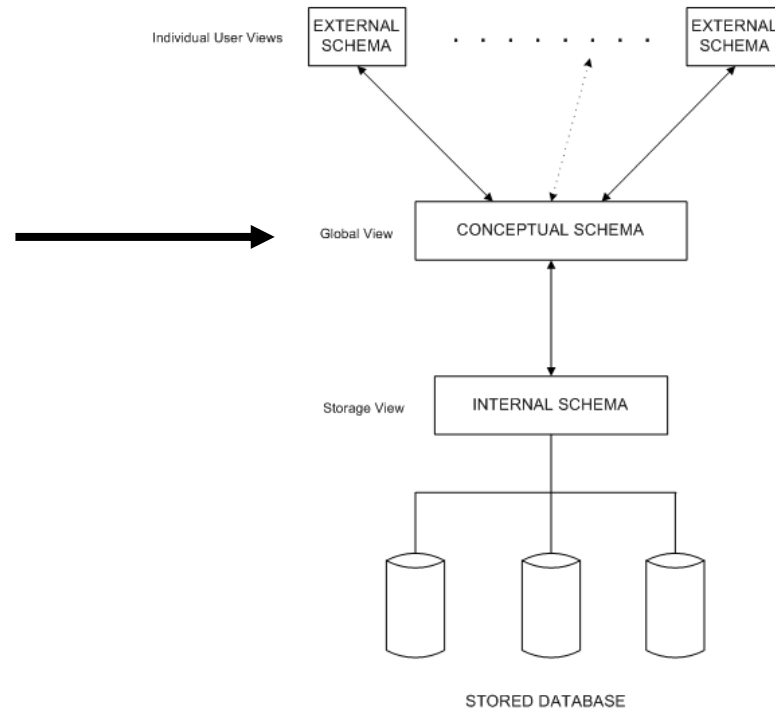


Figure 1.2 The ANSI/SPARC three-schema Architecture

# Growth – Let’s start here

(Note that we’ve added an “owner” attribute)

- We want name, gender, weight, and owner of female horses

```
SELECT Name, Gender, Weight, Owner FROM Horses WHERE Gender='F'
```

| Horses |        |       |        |        |         |
|--------|--------|-------|--------|--------|---------|
| Name   | Color  | Spots | Gender | Weight | Owner   |
| Amy    | Yellow | Yes   | F      | 920    | Mgrimes |
| Dave   | Grey   | No    | M      | 1800   | Mparks  |
| Ed     | Yellow | No    | M      | 1100   | Mgrimes |
| Sally  | Black  | No    | F      | 1200   | Rcooper |
| Sarah  | Grey   | No    | F      | 1700   | Mgrimes |
| Tom    | Red    | Yes   | M      | 1500   | Mgrimes |
| Jane   | Red    | No    | F      | 1300   | Rcooper |



| Horses-View |        |        |         |
|-------------|--------|--------|---------|
| Name        | Gender | Weight | Owner   |
| Amy         | F      | 920    | Mgrimes |
| Sally       | F      | 1200   | Rcooper |
| Sarah       | F      | 1700   | Mgrimes |
| Jane        | F      | 1300   | Rcooper |

What is this part doing? PROJECTING a subset of attributes (Reducing degree)

What is this part doing? SELECTING a subset of tuples (Reducing cardinality)



# Growth - Attributes

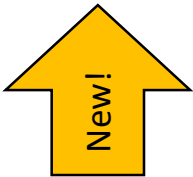
- Adding attributes does not affect the views

SELECT Name, Gender, Weight, Owner FROM Horses WHERE Gender='F'

| Horses |        |       |        |        |         |            |
|--------|--------|-------|--------|--------|---------|------------|
| Name   | Color  | Spots | Gender | Weight | Owner   | Phone#     |
| Amy    | Yellow | Yes   | F      | 920    | Mgrimes | 5551112345 |
| Dave   | Grey   | No    | M      | 1800   | Mparks  | 5552229876 |
| Ed     | Yellow | No    | M      | 1100   | Mgrimes | 5551112345 |
| Sally  | Black  | No    | F      | 1200   | Rcooper | 5553335656 |
| Sarah  | Grey   | No    | F      | 1700   | Mgrimes | 5551112345 |
| Tom    | Red    | Yes   | M      | 1500   | Mgrimes | 5551112345 |
| Jane   | Red    | No    | F      | 1300   | Rcooper | 5553335656 |



| Horses-View |        |        |         |
|-------------|--------|--------|---------|
| Name        | Gender | Weight | Owner   |
| Amy         | F      | 920    | Mgrimes |
| Sally       | F      | 1200   | Rcooper |
| Sarah       | F      | 1700   | Mgrimes |
| Jane        | F      | 1300   | Rcooper |



No change to the view!

# Growth - Relations

- Adding relations to the schema does not affect the views

SELECT Name, Gender, Weight, Owner FROM Horses WHERE Gender='F'

| Horses |        |       |        |        |         |            |
|--------|--------|-------|--------|--------|---------|------------|
| Name   | Color  | Spots | Gender | Weight | Owner   | Phone#     |
| Amy    | Yellow | Yes   | F      | 920    | Mgrimes | 5551112345 |
| Dave   | Grey   | No    | M      | 1800   | Mparks  | 5552229876 |
| Ed     | Yellow | No    | M      | 1100   | Mgrimes | 5551112345 |
| Sally  | Black  | No    | F      | 1200   | Rcooper | 5553335656 |
| Sarah  | Grey   | No    | F      | 1700   | Mgrimes | 5551112345 |
| Tom    | Red    | Yes   | M      | 1500   | Mgrimes | 5551112345 |
| Jane   | Red    | No    | F      | 1300   | Rcooper | 5553335656 |



| Horses-View |        |        |         |
|-------------|--------|--------|---------|
| Name        | Gender | Weight | Owner   |
| Amy         | F      | 920    | Mgrimes |
| Sally       | F      | 1200   | Rcooper |
| Sarah       | F      | 1700   | Mgrimes |
| Jane        | F      | 1300   | Rcooper |

| Owners  |         |         |
|---------|---------|---------|
| OwnerID | OwnName | Balance |
| 101     | Mgrimes | 825.00  |
| 102     | Mparks  | 0       |
| 103     | Rcooper | 210.00  |

No change to the view!



# Restructuring

- Restructuring should have an information equivalent outcome
- The view must be updated, but no impact to users/applications

SELECT Name, Gender, Weight, OwnName as Owner FROM Horses  
JOIN Owners on FKOwnerID = OwnerID WHERE gender = 'F'

| Horses |        |       |        |        |           |
|--------|--------|-------|--------|--------|-----------|
| Name   | Color  | Spots | Gender | Weight | FKOwnerID |
| Amy    | Yellow | Yes   | F      | 920    | 101       |
| Dave   | Grey   | No    | M      | 1800   | 102       |
| Ed     | Yellow | No    | M      | 1100   | 101       |
| Sally  | Black  | No    | F      | 1200   | 103       |
| Sarah  | Grey   | No    | F      | 1700   | 101       |
| Tom    | Red    | Yes   | M      | 1500   | 101       |
| Jane   | Red    | No    | F      | 1300   | 103       |



| Horses-View |        |        |         |
|-------------|--------|--------|---------|
| Name        | Gender | Weight | Owner   |
| Amy         | F      | 920    | Mgrimes |
| Sally       | F      | 1200   | Rcooper |
| Sarah       | F      | 1700   | Mgrimes |
| Jane        | F      | 1300   | Rcooper |

| Owners  |         |            |         |
|---------|---------|------------|---------|
| OwnerID | OwnName | Phone      | Balance |
| 101     | Mgrimes | 5551112345 | 825.00  |
| 102     | Mparks  | 5552229876 | 0       |
| 103     | Rcooper | 5553335656 | 210.00  |

No change to the view!

# Materialized Views

- Normal views are temporary
  - Only exist while the data is being access
  - This is what most views are
- Materialized views (snapshots) are constructed from one or more relations ahead of time
  - Used to “freeze” data at a certain point in time
  - Improves performance
  - Periodically recreated to reflect changes in the data
  - Often deleted if not used for a while, then recreated when needed again

# Module 6.5

## Views and Materialized Views

- What do views do and why are they useful?
- How do growth and restructuring impact views?
- What is the difference in a normal view and a materialized view?

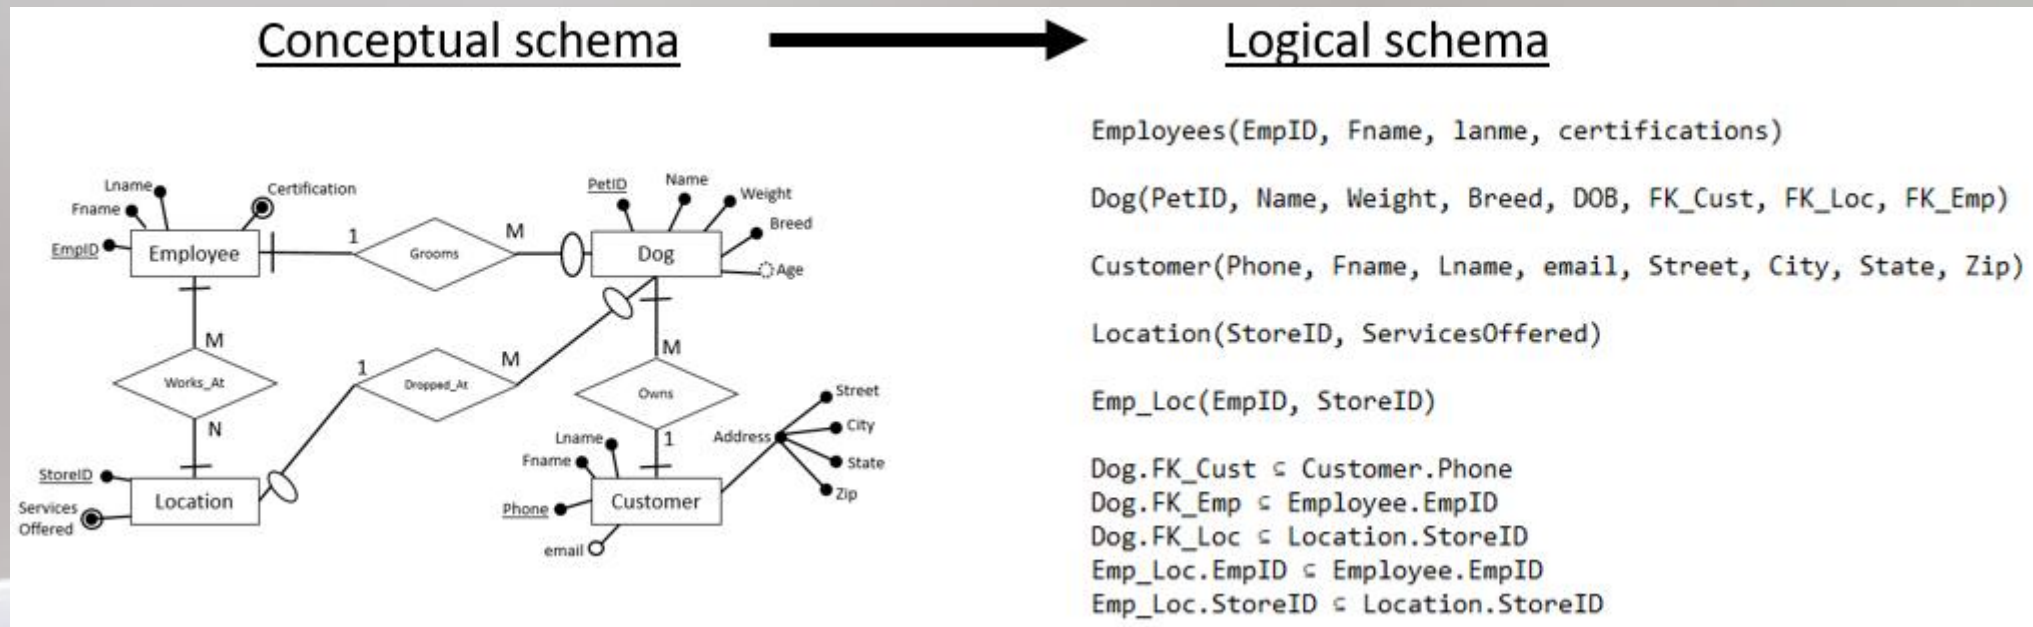
Break

# Module 6.7

## Mapping an ER Model to a Logical Schema

Next step in transforming business rules into an actual DB

ERD → Design Specific ERD → Logical Schema



# Mapping an ER Model to a Logical Schema

- Next step in transforming business rules into an actual DB
  - ERD → Design Specific ERD → Logical Schema
- Unfortunately, it is often difficult to represent all attributes of one schema in another
  - “Information Reducing”
  - Must carry forward all requirements, either by modeling or in semantic integrity constraints

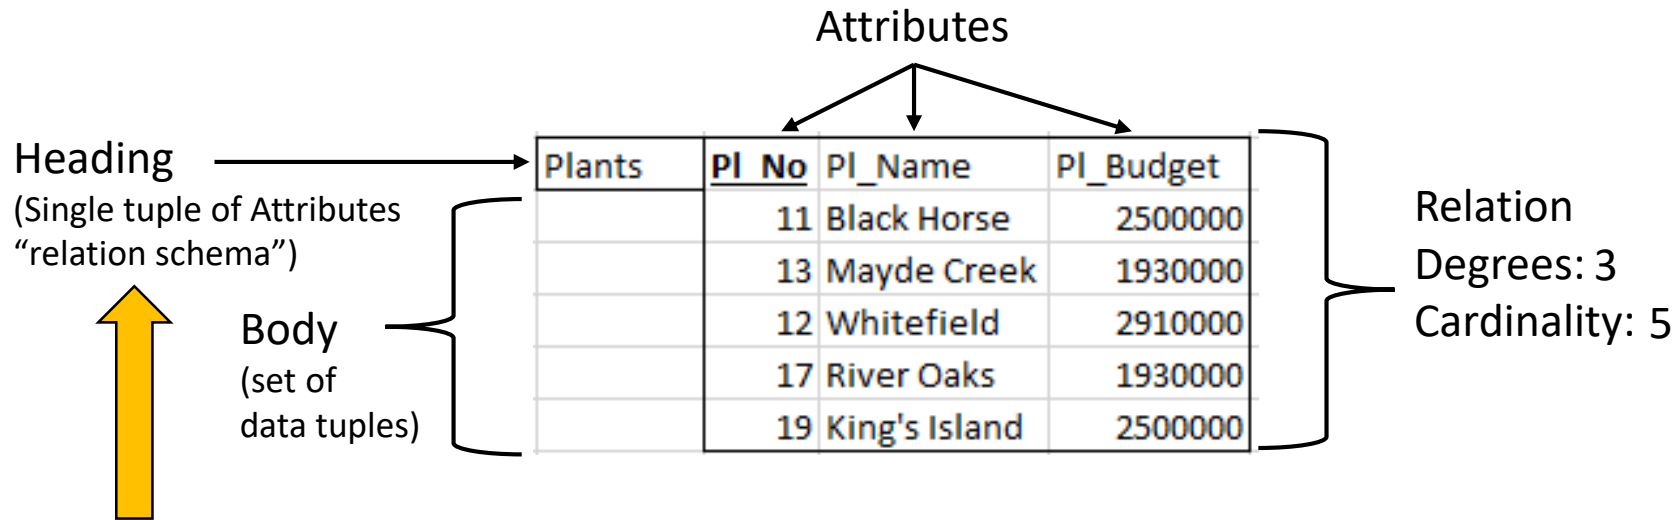


# Steps for Mapping Entity Types

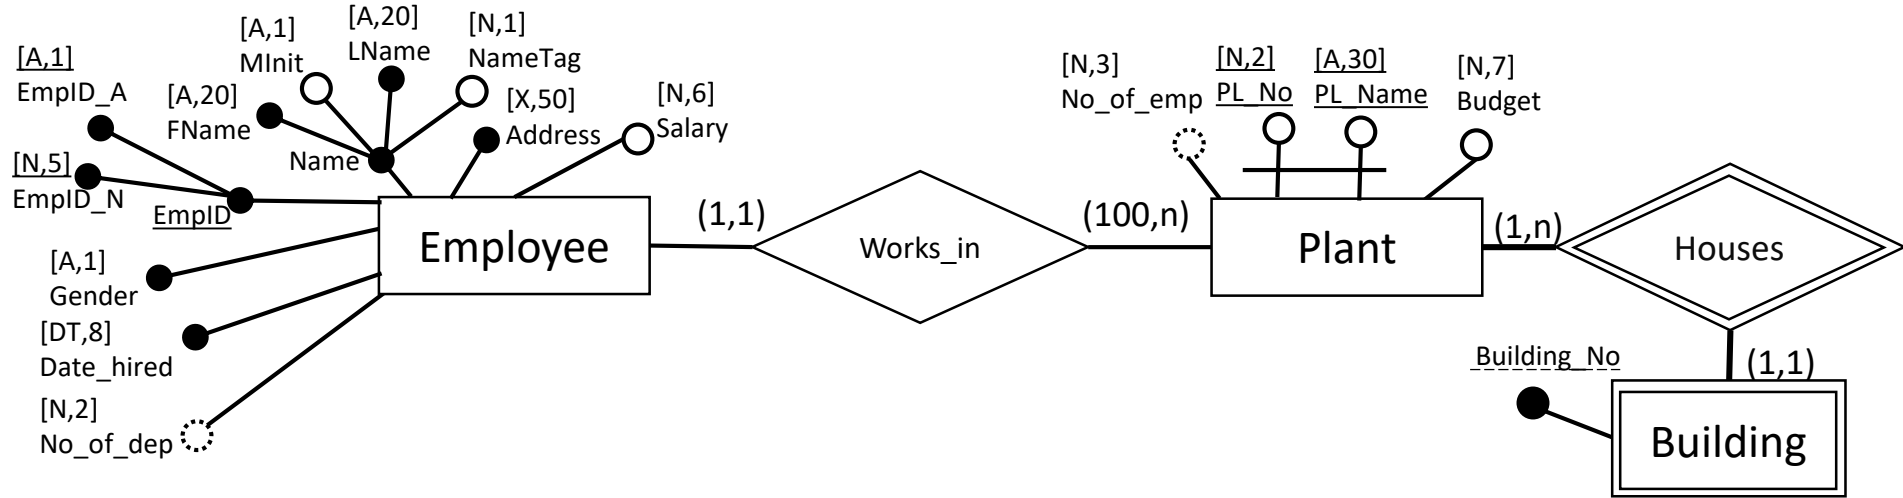
- Create a **relation schema** for each entity type in the ER diagram.
- Create an attribute for every stored attribute. This implies:
  - For composite attributes only their atomic components are recorded
  - Derived attributes are not recorded
  - Multi-valued attributes do not exist in a Design-Specific ER model
- Choose a primary key from among the candidate keys by underlining the attribute(s) making up the primary key.
- For each weak entity type, add the primary key of the identifying parent entity type as attribute(s) in the relation schema.
  - The attribute(s) thus added plus the partial key of the weak entity type form the primary key of the relation schema representing the weak entity type.

# Review: what is a relation?

- A mathematical terms approximated by a two dimensional table:
  - Heading – a single tuple listing the attributes (Relation Schema)
  - Body – collection of data tuples



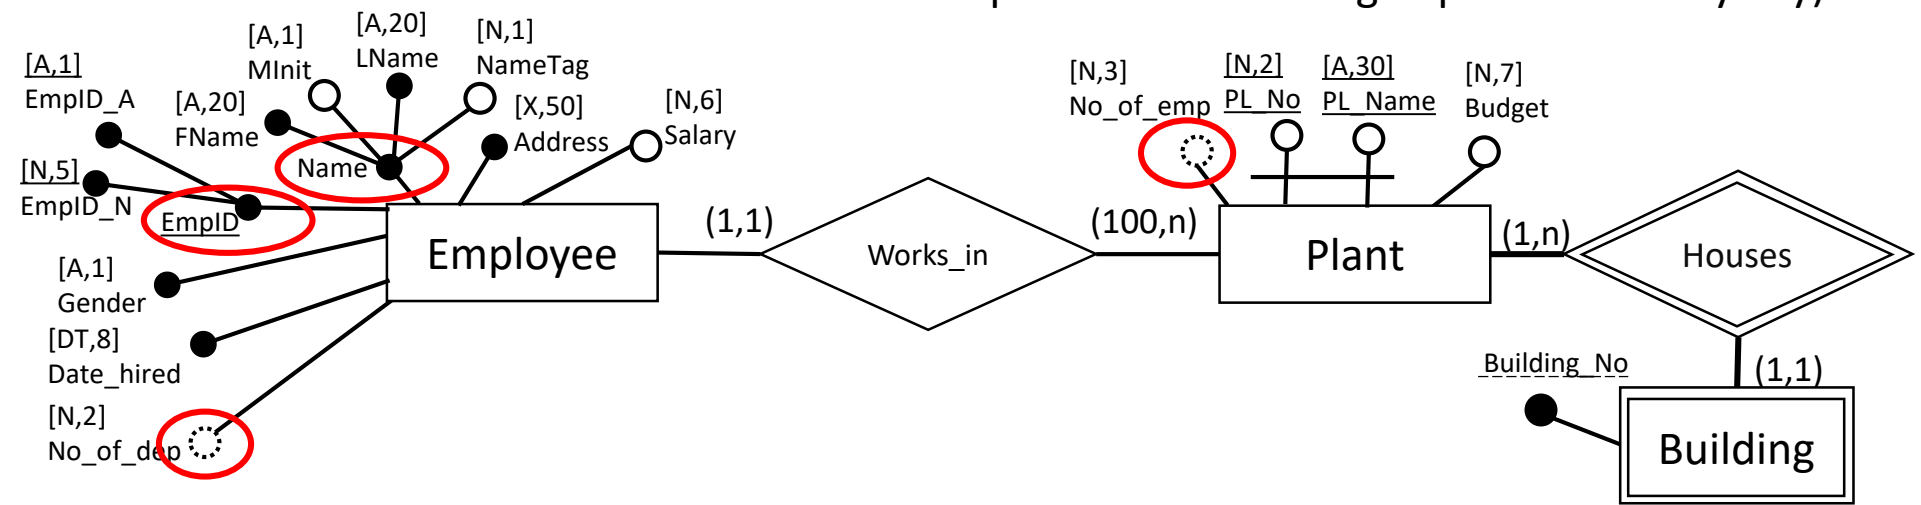
## Step 1: Create relation schemas for entities



|          |  |  |  |  |  |  |  |  |  |  |
|----------|--|--|--|--|--|--|--|--|--|--|
| Employee |  |  |  |  |  |  |  |  |  |  |
| Plant    |  |  |  |  |  |  |  |  |  |  |
| Building |  |  |  |  |  |  |  |  |  |  |

# Step 2: Create attributes

No derived or composite attributes allowed  
No MV attributes or M:N relationship (should not be present in the Design Specific ERD anyway)



|          |         |         |       |       |       |         |        |         |        |            |
|----------|---------|---------|-------|-------|-------|---------|--------|---------|--------|------------|
| Employee | EmpID_A | EmpID_N | FName | MInit | LName | NameTag | Gender | Address | Salary | Date_hired |
|----------|---------|---------|-------|-------|-------|---------|--------|---------|--------|------------|

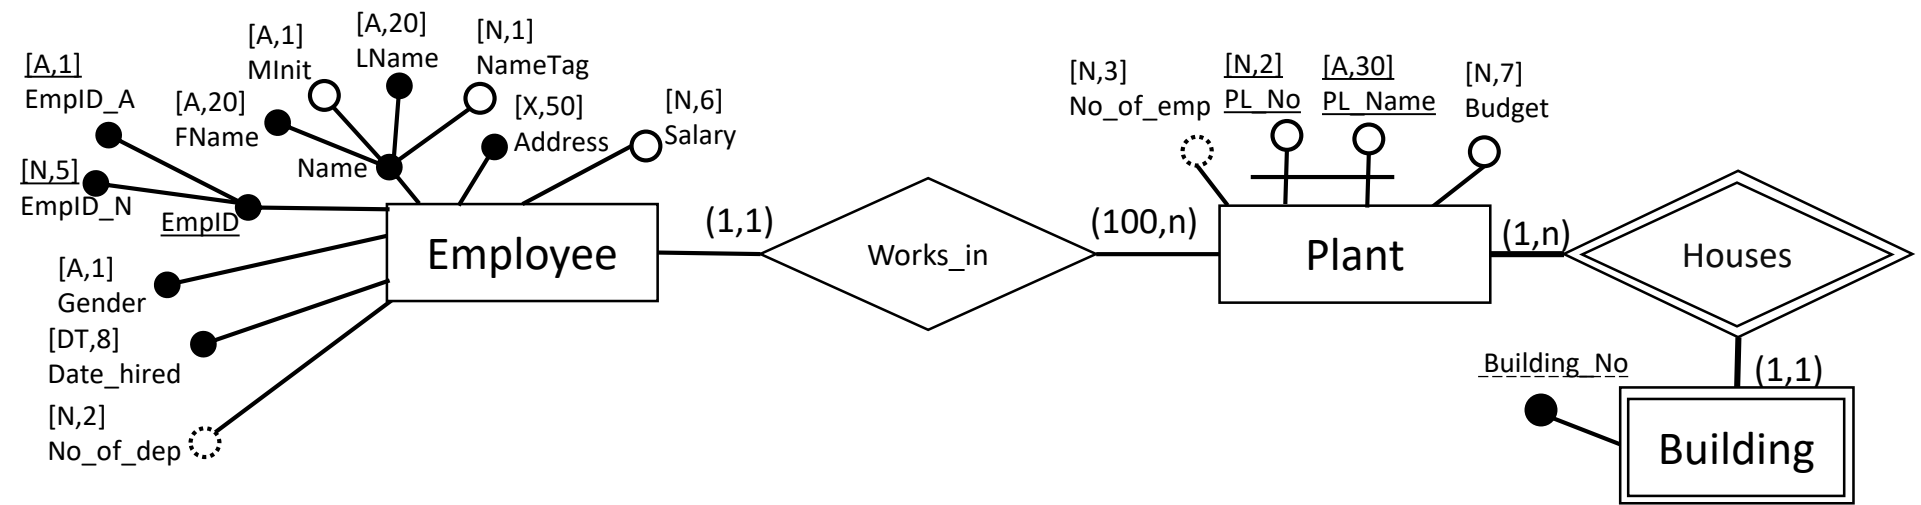
|       |         |       |        |
|-------|---------|-------|--------|
| Plant | PL_Name | PL_No | Budget |
|-------|---------|-------|--------|

|          |             |
|----------|-------------|
| Building | Building_No |
|----------|-------------|

Note 1: Only the atomic attributes constituting EmpID and Name in the ERD are recorded in EMPLOYEE.

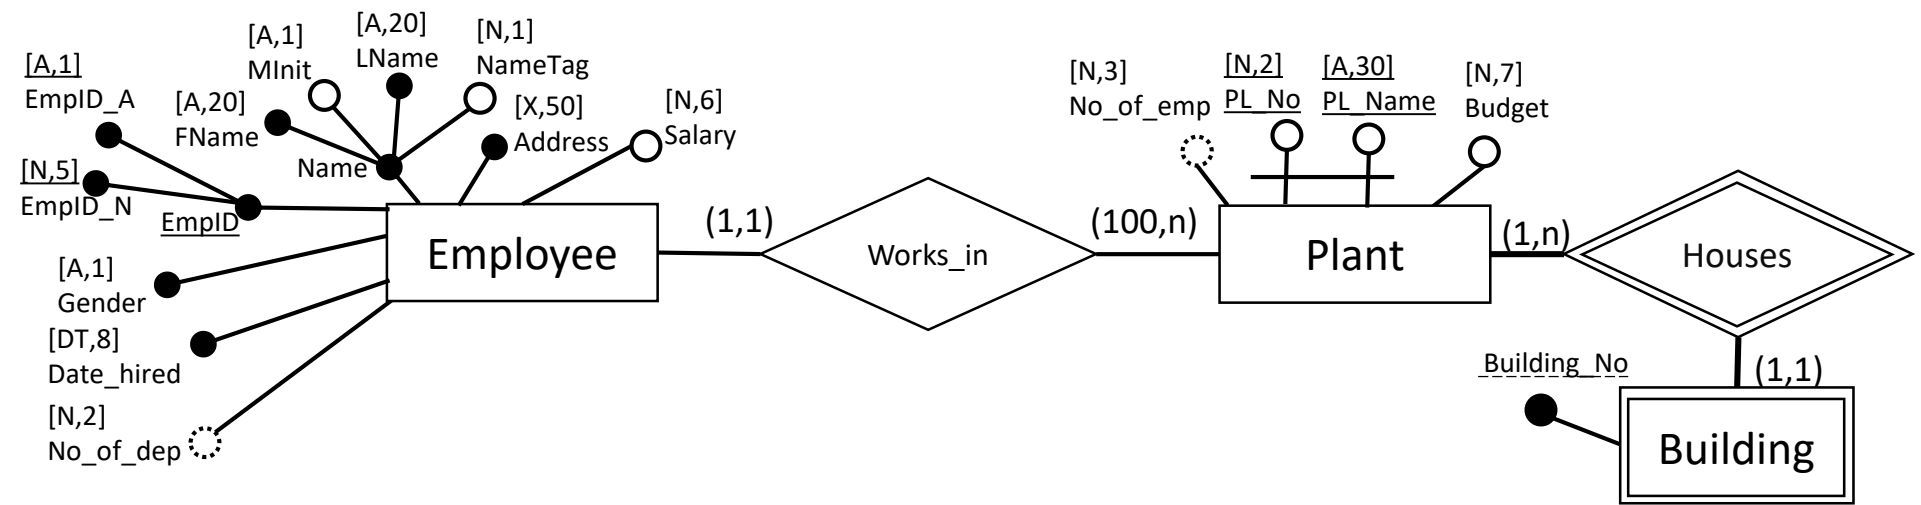
Note 2: The derived attributes, No\_of\_dep in EMPLOYEE and No\_of\_emp in PLANT are not captured here. They will be derived in the external schema (i.e., a view).

# Step 3: Choose PK from the CK by underlining them



|          |                    |                |        |       |       |         |        |         |        |            |
|----------|--------------------|----------------|--------|-------|-------|---------|--------|---------|--------|------------|
| Employee | <u>EmpID_A</u>     | <u>EmpID_N</u> | FName  | MInit | LName | NameTag | Gender | Address | Salary | Date_hired |
| Plant    | PL_Name            | <u>PL_No</u>   | Budget |       |       |         |        |         |        |            |
| Building | <u>Building_No</u> | ???            |        |       |       |         |        |         |        |            |

# Step 4: Add Primary Key of strong entity to weak entities



|          |                |                |       |       |       |         |        |         |        |            |
|----------|----------------|----------------|-------|-------|-------|---------|--------|---------|--------|------------|
| Employee | <u>EmpID_A</u> | <u>EmpID_N</u> | FName | MInit | LName | NameTag | Gender | Address | Salary | Date_hired |
|----------|----------------|----------------|-------|-------|-------|---------|--------|---------|--------|------------|

|       |         |              |        |
|-------|---------|--------------|--------|
| Plant | PL_Name | <u>PL_No</u> | Budget |
|-------|---------|--------------|--------|

|          |                    |                 |
|----------|--------------------|-----------------|
| Building | <u>Building_No</u> | <u>FK_PL_No</u> |
|----------|--------------------|-----------------|

# Logical Schemas Defined!

- Next step is to map the relationships

|          |                |                |       |       |       |         |        |         |        |            |
|----------|----------------|----------------|-------|-------|-------|---------|--------|---------|--------|------------|
| Employee | <u>EmpID_A</u> | <u>EmpID_N</u> | FName | MInit | LName | NameTag | Gender | Address | Salary | Date_hired |
|----------|----------------|----------------|-------|-------|-------|---------|--------|---------|--------|------------|

|       |         |              |        |
|-------|---------|--------------|--------|
| Plant | PL_Name | <u>PL_No</u> | Budget |
|-------|---------|--------------|--------|

|          |                    |                 |
|----------|--------------------|-----------------|
| Building | <u>Building_No</u> | <u>FK_PL_No</u> |
|----------|--------------------|-----------------|

# Mapping Relationship Types

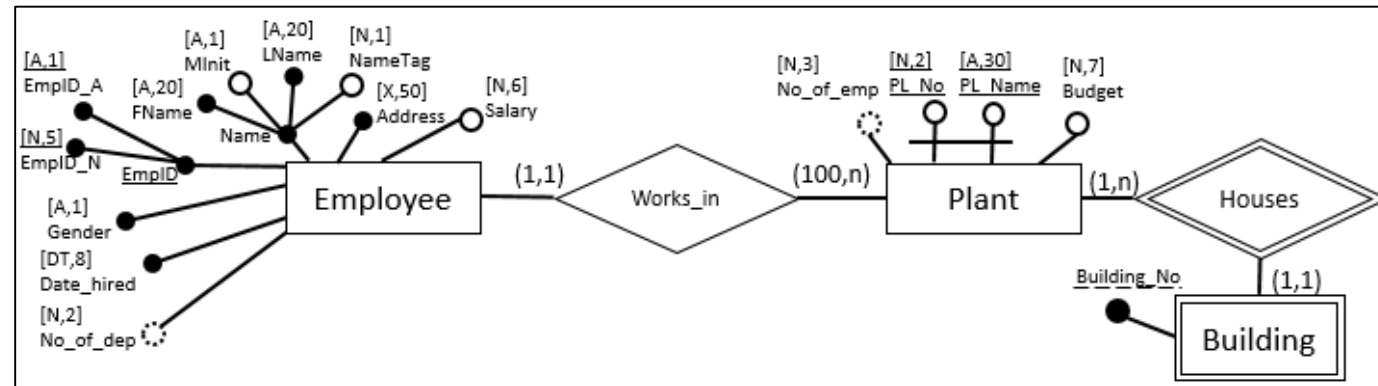
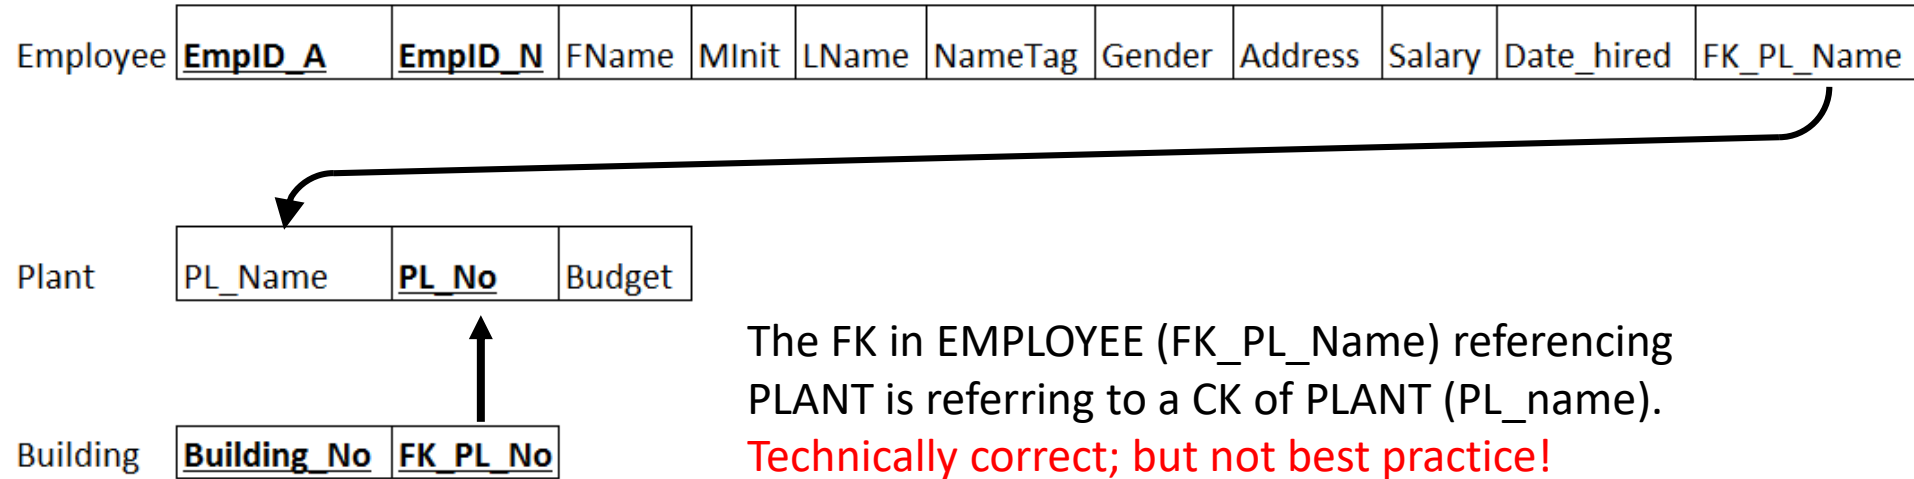
- Characteristic of the Design-Specific ER model:
  - Only binary or recursive relationships are present – why?
  - Only 1:1 or 1:n cardinality constraints are present – why?
- Relationship type mapped by enforcing a foreign key constraint
- The foreign key constraint requires that the referenced attribute(s) be a candidate key of (advisably, primary key) the referenced relation
- Reduction in information 😞
  - Participation constraints are not mapped.
  - Cardinality constraint of (1:n) is implicit by default



# Mapping Relationship Types [1:n]

- Add a foreign key attribute to the referencing schema (Child)
  - Foreign key must share the same domain with a candidate key in the referenced schema (Parent)
  - Best Practice: use primary key as the referenced attribute
- Remember:
  - The entity type on the “many-side” of the relationship type is the referencing relation schema (“child”)
  - The entity type on the “one-side” is the referenced relation schema (“parent”)
- Represented with a directed arc (i.e., a line) with arrowhead pointing at the referenced candidate key

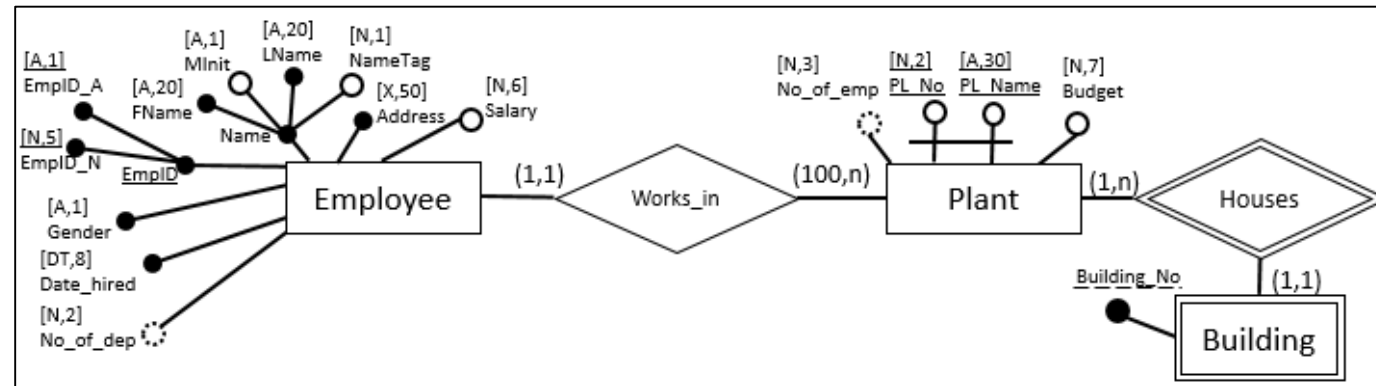
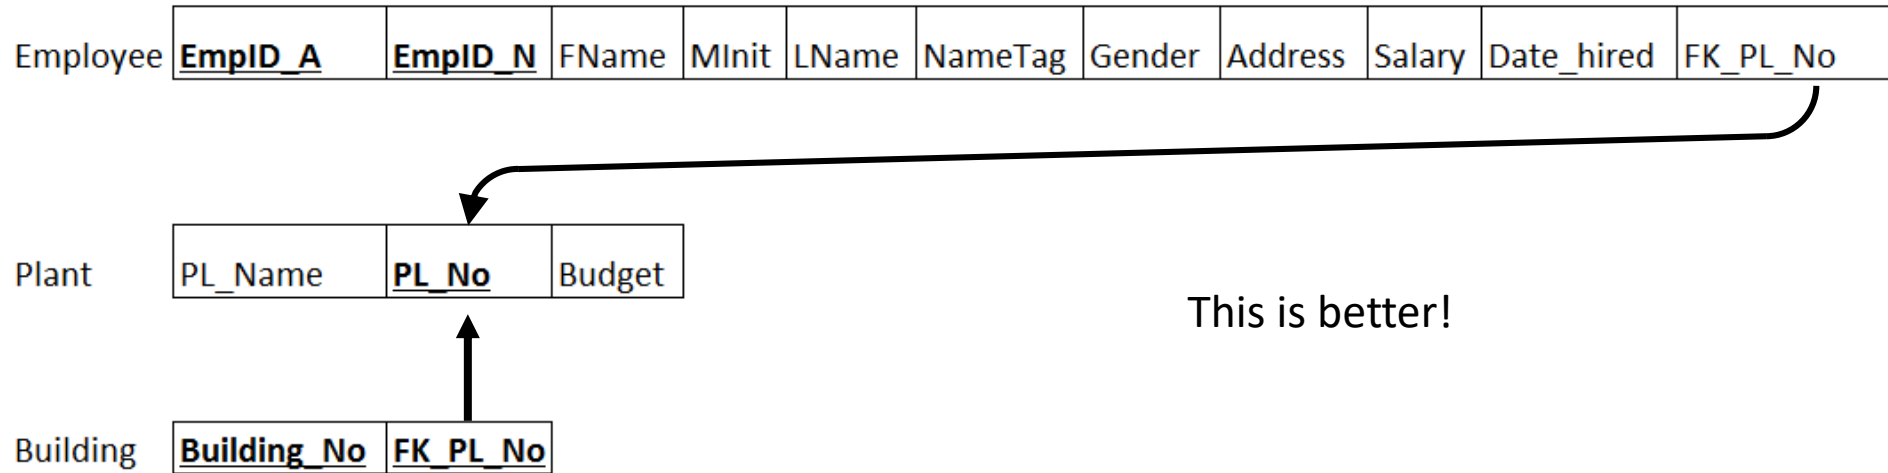
# Mapping Relationship Types [1:n]



Note 1: Participation constraints are not mapped.

Note 2: Cardinality constraint of (1:n) is implicit by default

# Mapping Relationship Types [1:n]

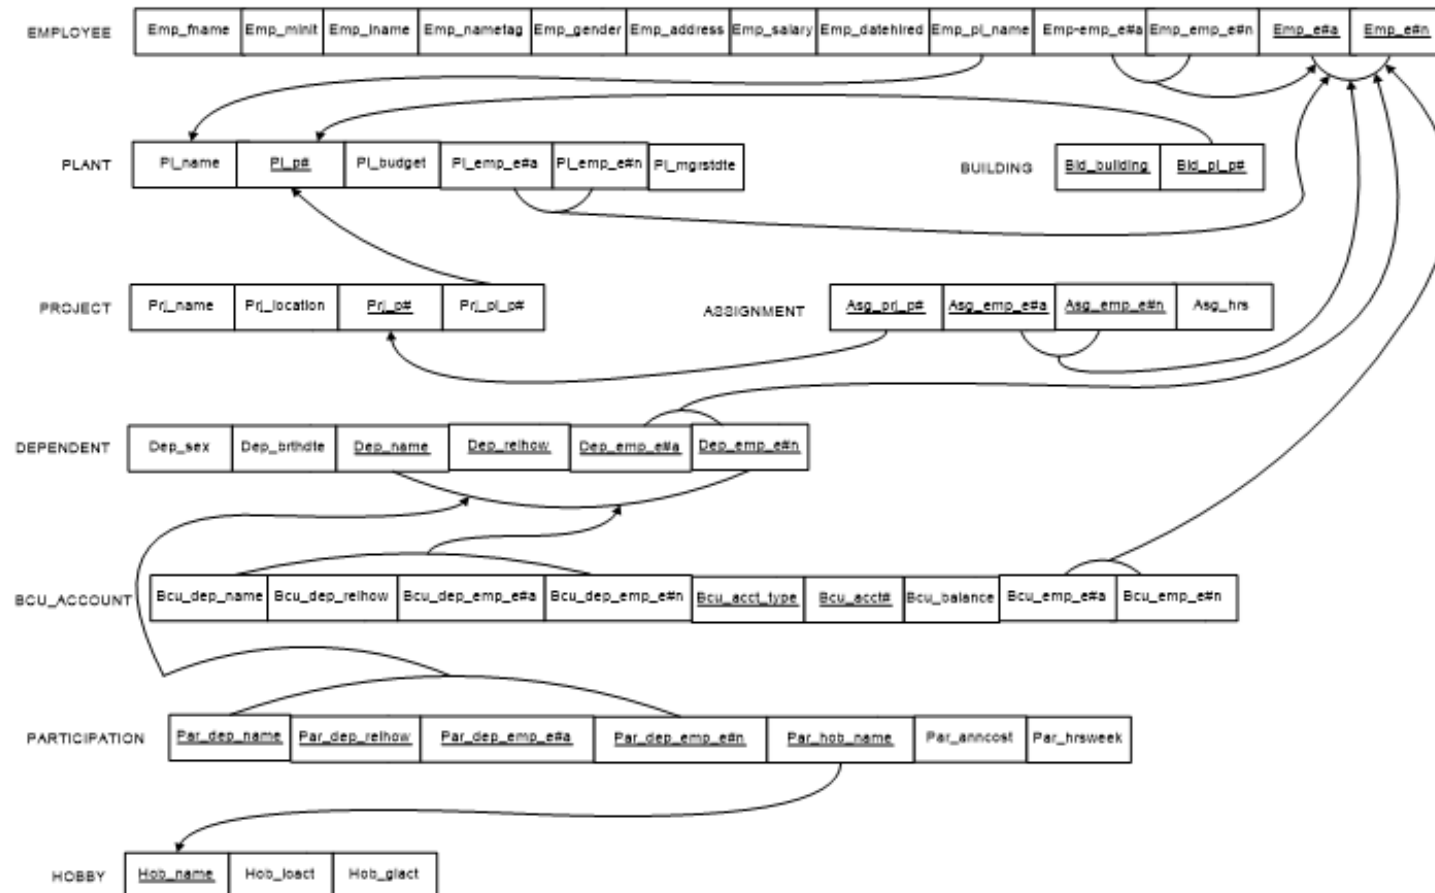


Note 1: Participation constraints are not mapped.

Note 2: Cardinality constraint of (1:n) is implicit by default

# Alternative Notation

- The lines are useful, but may get messy for large models
- We can represent Inclusion Dependency as  $R2.\{A2\} \subseteq R1.\{A1\}$



# Alternative Notation

EMPLOYEE(EmpID\_A, EmpID\_N, FName, MInit, LName, NameTag, Gender, Address, Salary, Date\_hired, FK\_PL\_No)

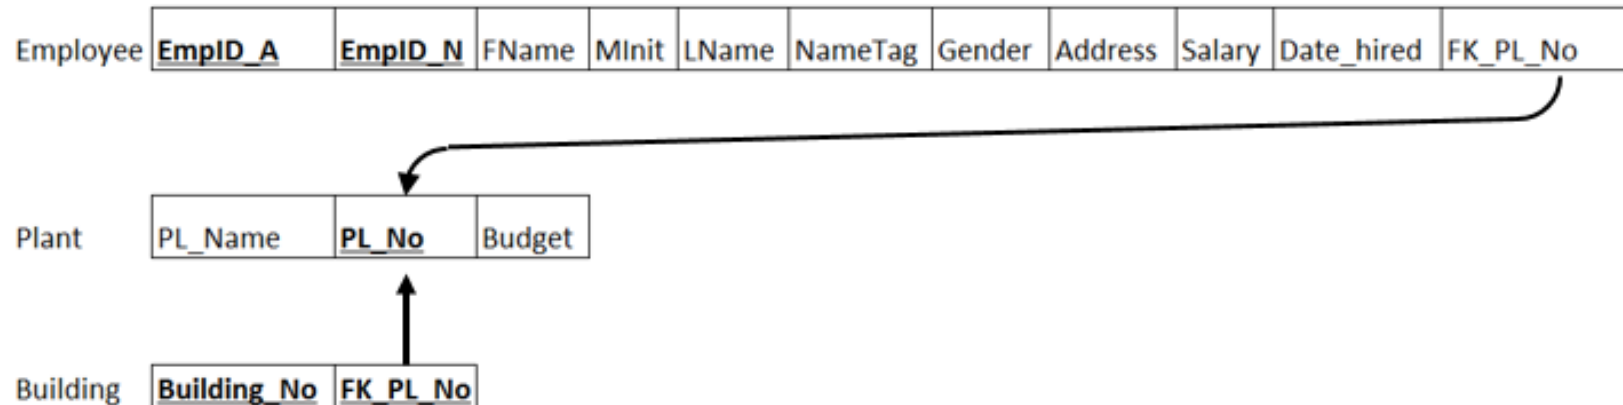
#EMPLOYEE.{FK\_PL\_No}  $\subseteq$  PLANT.{PL\_No}

PLANT (PL\_Name, PL\_No, Budget)

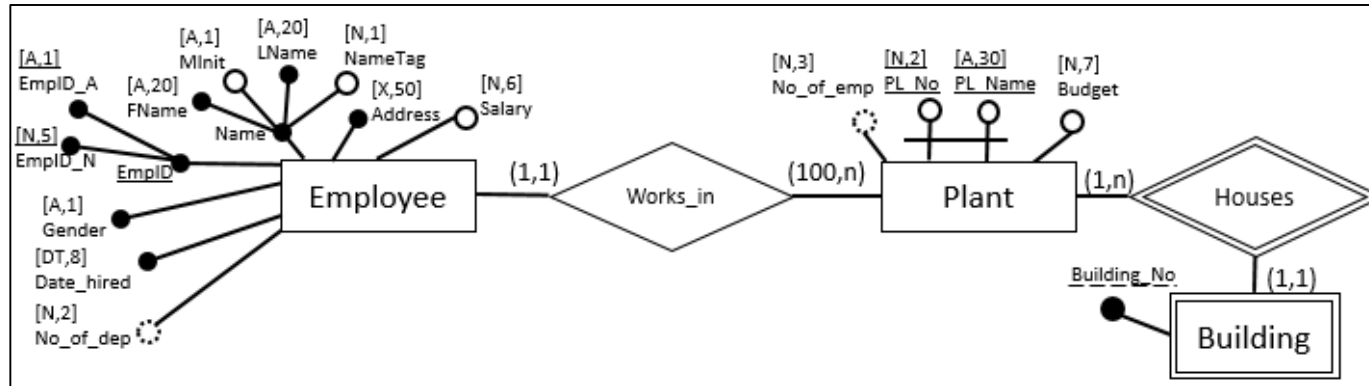
BUILDING (Building\_No, FK\_PL\_No)

#BUILDING.{FK\_PL\_No}  $\subseteq$  PLANT.{PL\_No}

Equivalent

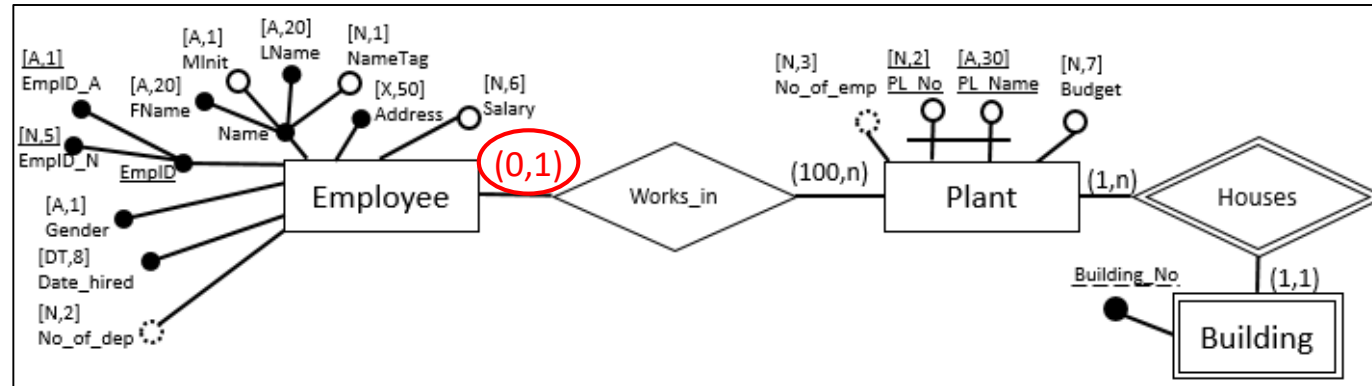


# An alternate mapping for [1:n]

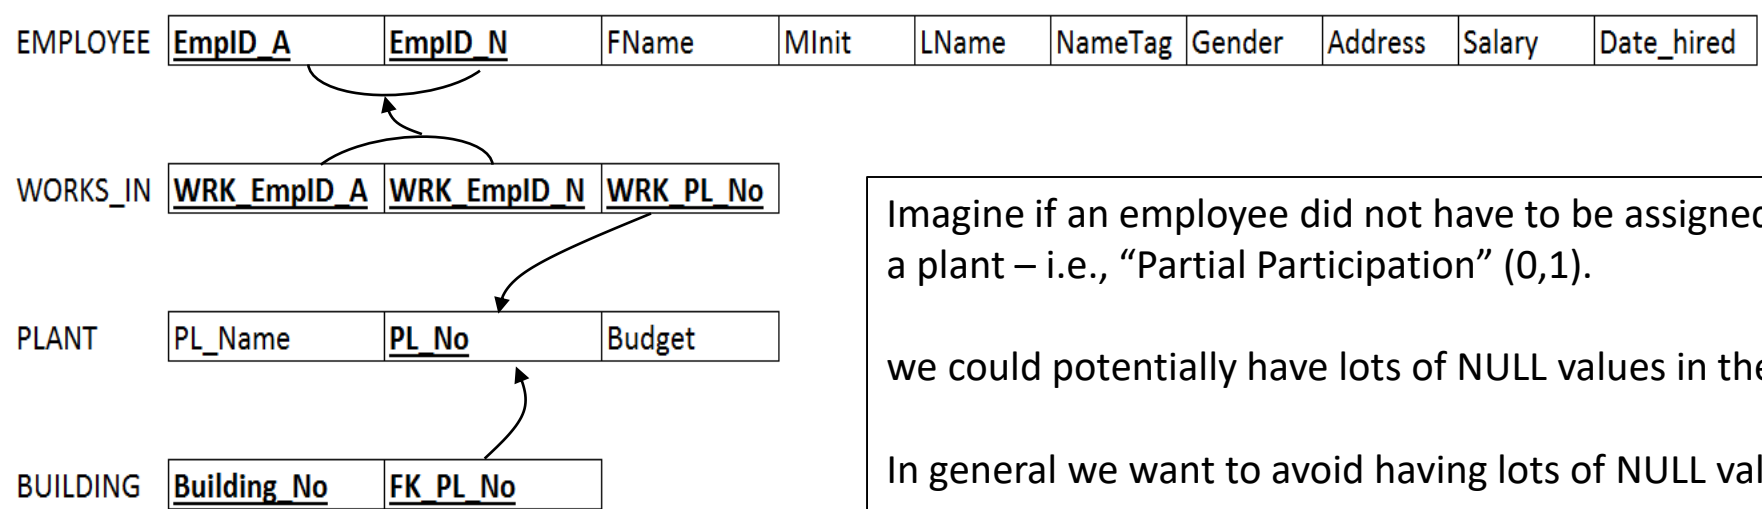


- While not mapped, participation does impact the design
  - Reminder: Participation is “Partial” (optional) or “Total” (mandatory)

# An alternate mapping for [1:n]



- Instead of adding a FK to the child table, create a gerund



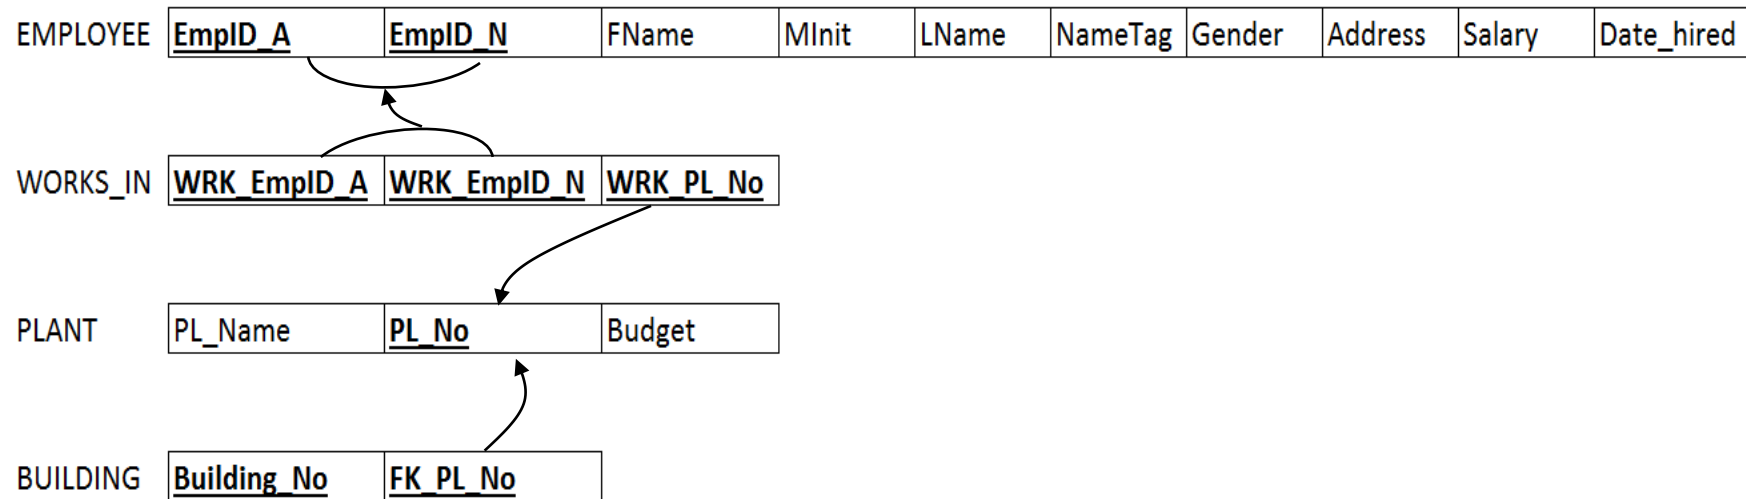
Imagine if an employee did not have to be assigned to a plant – i.e., “Partial Participation” (0,1).

we could potentially have lots of NULL values in the FK.

In general we want to avoid having lots of NULL values

You do not HAVE to model 1:n this way, but it is useful in some scenarios

# An alternate mapping for [1:n]



$\text{WORKS\_IN}.\{\text{WRK\_EmpID\_A}, \text{WRK\_EmpID\_N}\} \subseteq \text{EMPLOYEE}.\{\text{EmpID\_A}, \text{EmpID\_N}\}$

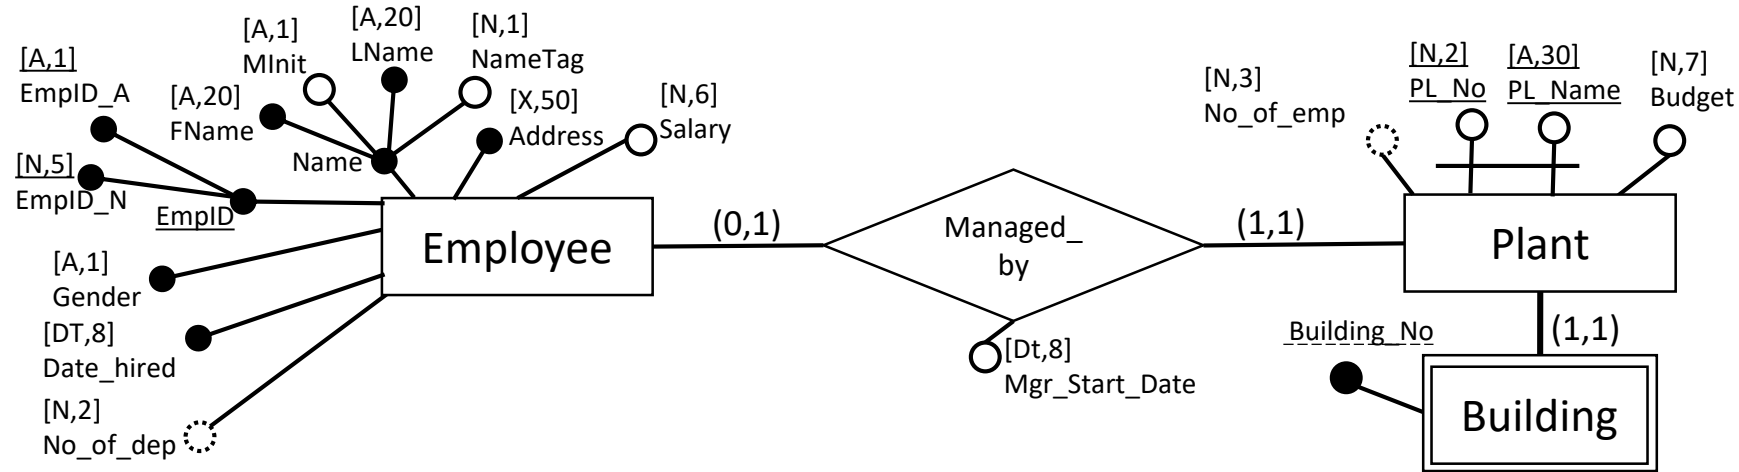
$\text{WORKS\_IN}.\{\text{WRK\_PL\_No}\} \subseteq \text{PLANT}.\{\text{PL\_No}\}$

$\text{BUILDING}.\{\text{FK\_PL\_No}\} \subseteq \text{PLANT}.\{\text{PL\_No}\}$

- Using this design, if an employee is not assigned to a plant, we simply do not create an instance in the WORKS\_IN gerund, rather than having NULL values in the FK in the Employee relation

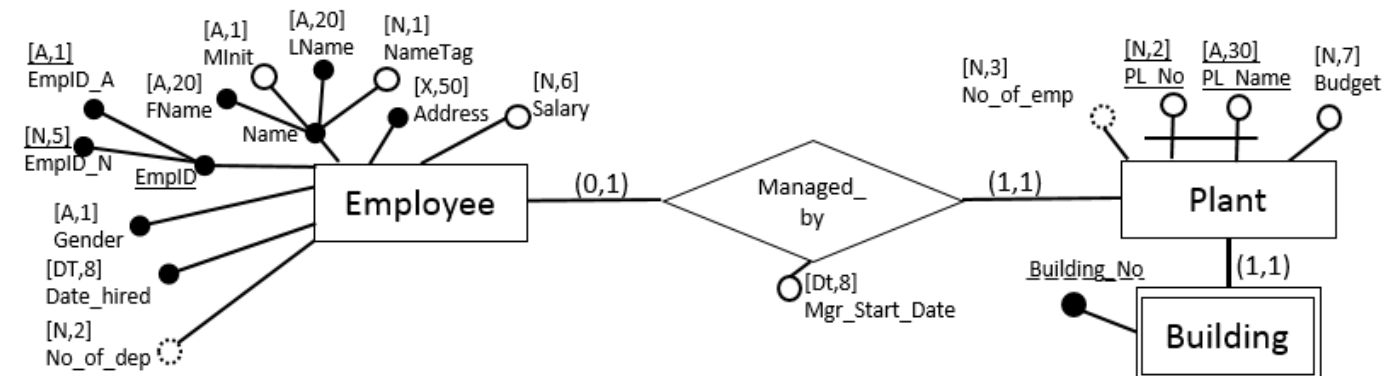


# Mapping a 1:1 relationship



- Which side is the “child” in a 1:1 relationship?
  - They are equivalent, so technically either can have the FK
- So which side to use?
  - The relation with partial participation (employee) will have many NULL values – should be considered the “parent”
  - The relation with total participation (plant) will have no NULL values – should be considered the “Child”

# Mapping a 1:1 relationship



- Both technically work, but option 2 is MUCH better

|          |                |                |       |       |       |         |        |         |        |            |              |             |
|----------|----------------|----------------|-------|-------|-------|---------|--------|---------|--------|------------|--------------|-------------|
| EMPLOYEE | <u>EmpID_A</u> | <u>EmpID_N</u> | FName | MInit | LName | NameTag | Gender | Address | Salary | Date_hired | FK_MGR_PL_No | EMP_MGRDate |
|----------|----------------|----------------|-------|-------|-------|---------|--------|---------|--------|------------|--------------|-------------|

|       |         |              |        |
|-------|---------|--------------|--------|
| PLANT | PL_Name | <u>PL_No</u> | Budget |
|-------|---------|--------------|--------|

|          |                    |          |
|----------|--------------------|----------|
| BUILDING | <u>Building_No</u> | FK_PL_No |
|----------|--------------------|----------|

Option 1

If we put the FK with the relation that has partial participation, we will get many NULL values ☹️

|          |                |                |       |       |       |         |        |         |        |            |
|----------|----------------|----------------|-------|-------|-------|---------|--------|---------|--------|------------|
| EMPLOYEE | <u>EmpID_A</u> | <u>EmpID_N</u> | FName | MInit | LName | NameTag | Gender | Address | Salary | Date_hired |
|----------|----------------|----------------|-------|-------|-------|---------|--------|---------|--------|------------|

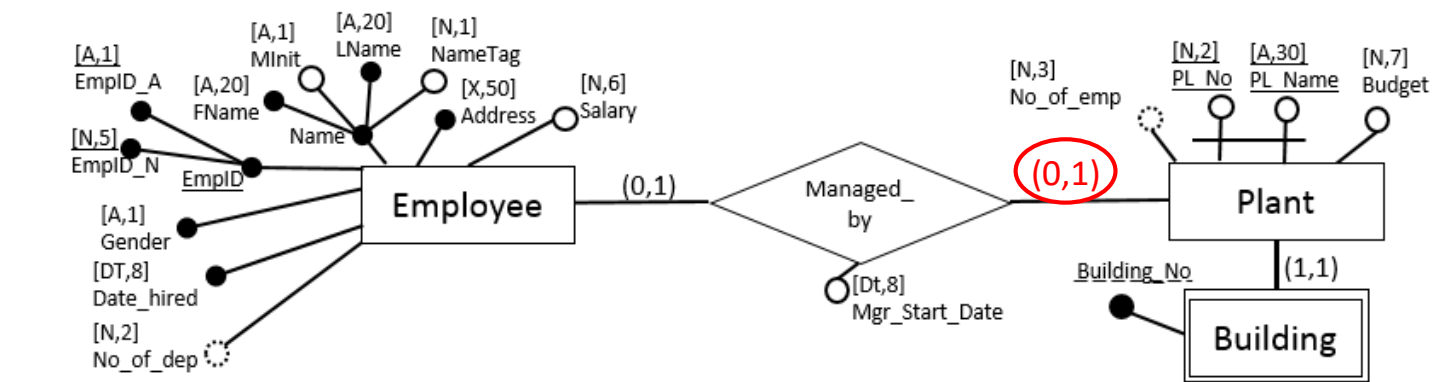
|       |         |              |        |                |                |            |
|-------|---------|--------------|--------|----------------|----------------|------------|
| PLANT | PL_Name | <u>PL_No</u> | Budget | FK_MGR_EmpID_A | FK_MGR_EmpID_N | PL_MGRDate |
|-------|---------|--------------|--------|----------------|----------------|------------|

|          |                    |          |
|----------|--------------------|----------|
| BUILDING | <u>Building_No</u> | FK_PL_No |
|----------|--------------------|----------|

Option 2

If we put the FK with the relation that has mandatory participation, no NULL values 😊

# One more twist on 1:1 relationships



|          |                |                |       |       |       |         |        |         |        |            |
|----------|----------------|----------------|-------|-------|-------|---------|--------|---------|--------|------------|
| EMPLOYEE | <u>EmpID_A</u> | <u>EmpID_N</u> | FName | MInit | LName | NameTag | Gender | Address | Salary | Date_hired |
|----------|----------------|----------------|-------|-------|-------|---------|--------|---------|--------|------------|

|            |                  |             |             |             |
|------------|------------------|-------------|-------------|-------------|
| MANAGED_BY | <u>MAN_PL No</u> | Man_EmpID_A | MAN_EmpID_N | MAN_MGRDate |
|------------|------------------|-------------|-------------|-------------|

|       |         |              |        |
|-------|---------|--------------|--------|
| PLANT | PL_Name | <u>PL No</u> | Budget |
|-------|---------|--------------|--------|

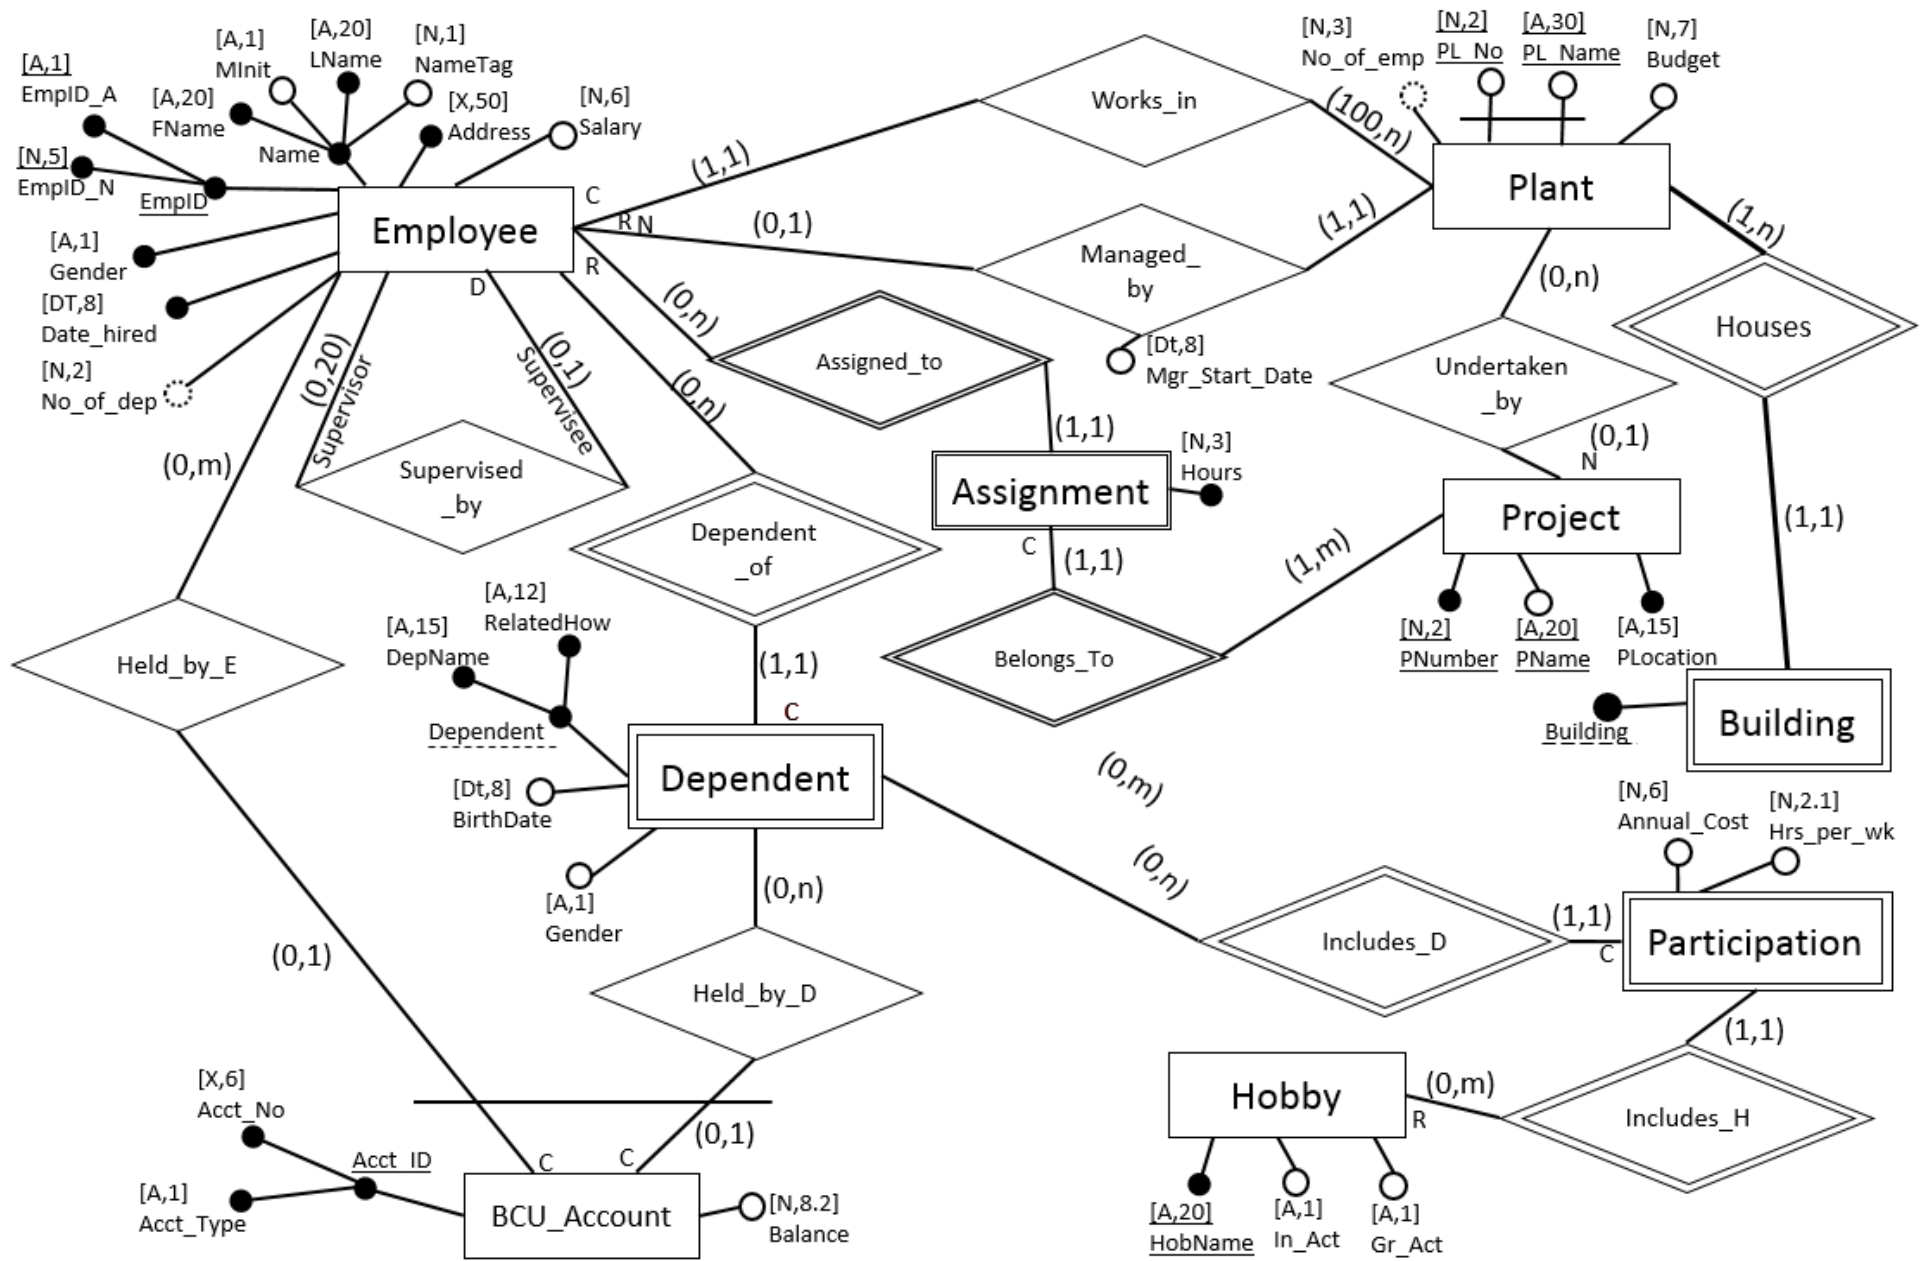
|          |                    |                 |
|----------|--------------------|-----------------|
| BUILDING | <u>Building No</u> | <u>FK PL No</u> |
|----------|--------------------|-----------------|

Imagine participation is partial on both sides (i.e., a plant does not have to have a manager).

We could potentially have lots of NULL values in the FK no matter where we put it.

Again, we can use a gerund to fix this situation.

# Creating a Logical Model – Page 313-315

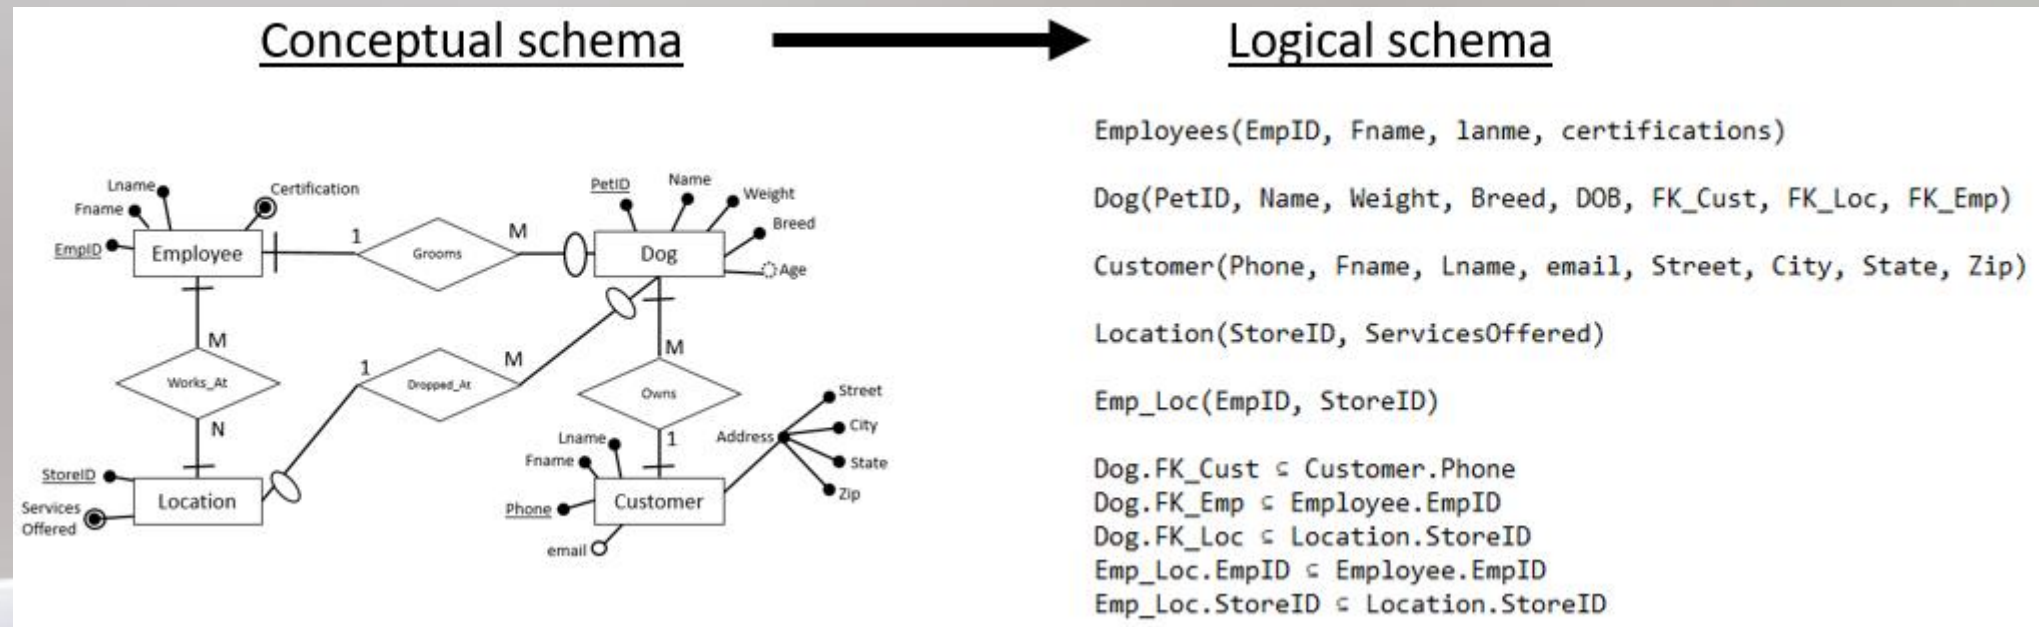


# Module 6.7

## Mapping an ER Model to a Logical Schema

Next step in transforming business rules into an actual DB

ERD → Design Specific ERD → Logical Schema



# Question for you....

```
CREATE TABLE instructors (  
  Inst_id      char(15) CONSTRAINT pk_inst PRIMARY KEY,  
  Inst_name    varchar(50) CONSTRAINT nn_Instname NOT NULL,  
  Inst_office  varchar(8)  
);
```

```
CREATE TABLE courses(  
  Crs_id      char(8) CONSTRAINT pk_crs PRIMARY KEY,  
  Crs_name    varchar(50) CONSTRAINT nn_crs NOT NULL,  
  Crs_hours   numeric(1)  
);
```

```
CREATE TABLE sections (  
  Sec_id      char(5) CONSTRAINT pk_sec PRIMARY KEY,  
  Sec_time    datetime,  
  Sec_location nvarchar(10),  
  Sec_term    char(8) CONSTRAINT nn_secterm NOT NULL,  
  Sec_inst_id char(15) CONSTRAINT fk_inst FOREIGN KEY REFERENCES instructors (inst_id),  
  Sec_crs_id  char(8) CONSTRAINT fk_cou FOREIGN KEY REFERENCES courses (crs_id),  
  CONSTRAINT nn_sec_crs NOT NULL  
);
```

```
CREATE TABLE students (  
  Stu_id      char(15) CONSTRAINT pk_stu PRIMARY KEY,  
  Stu_name    nvarchar(50) CONSTRAINT nn_stuname NOT NULL,  
  Stu_phone   char(12),  
  Stu_GPA     numeric(1,3),  
  CONSTRAINT chk_gpa CHECK (Stu_GPA <= 4.0)  
);
```

```
CREATE TABLE enrollment(  
  Enr_stu_id  char(15),  
  Enr_sec_id  char(8),  
  Enr_Grade   char(2),  
  CONSTRAINT fk_stu FOREIGN KEY Enr_stu_id REFERENCES students (stu_id),  
  CONSTRAINT fk_sec FOREIGN KEY Enr_sec_id REFERENCES sections (sec_id),  
  CONSTRAINT pk_enr PRIMARY KEY (Enr_stu_id, Enr_sec_id),  
);
```

Could you create an ERD or Logical Schema based on this DDL code?

Could you create this code based on an ERD or logical schema?

# Question for you....

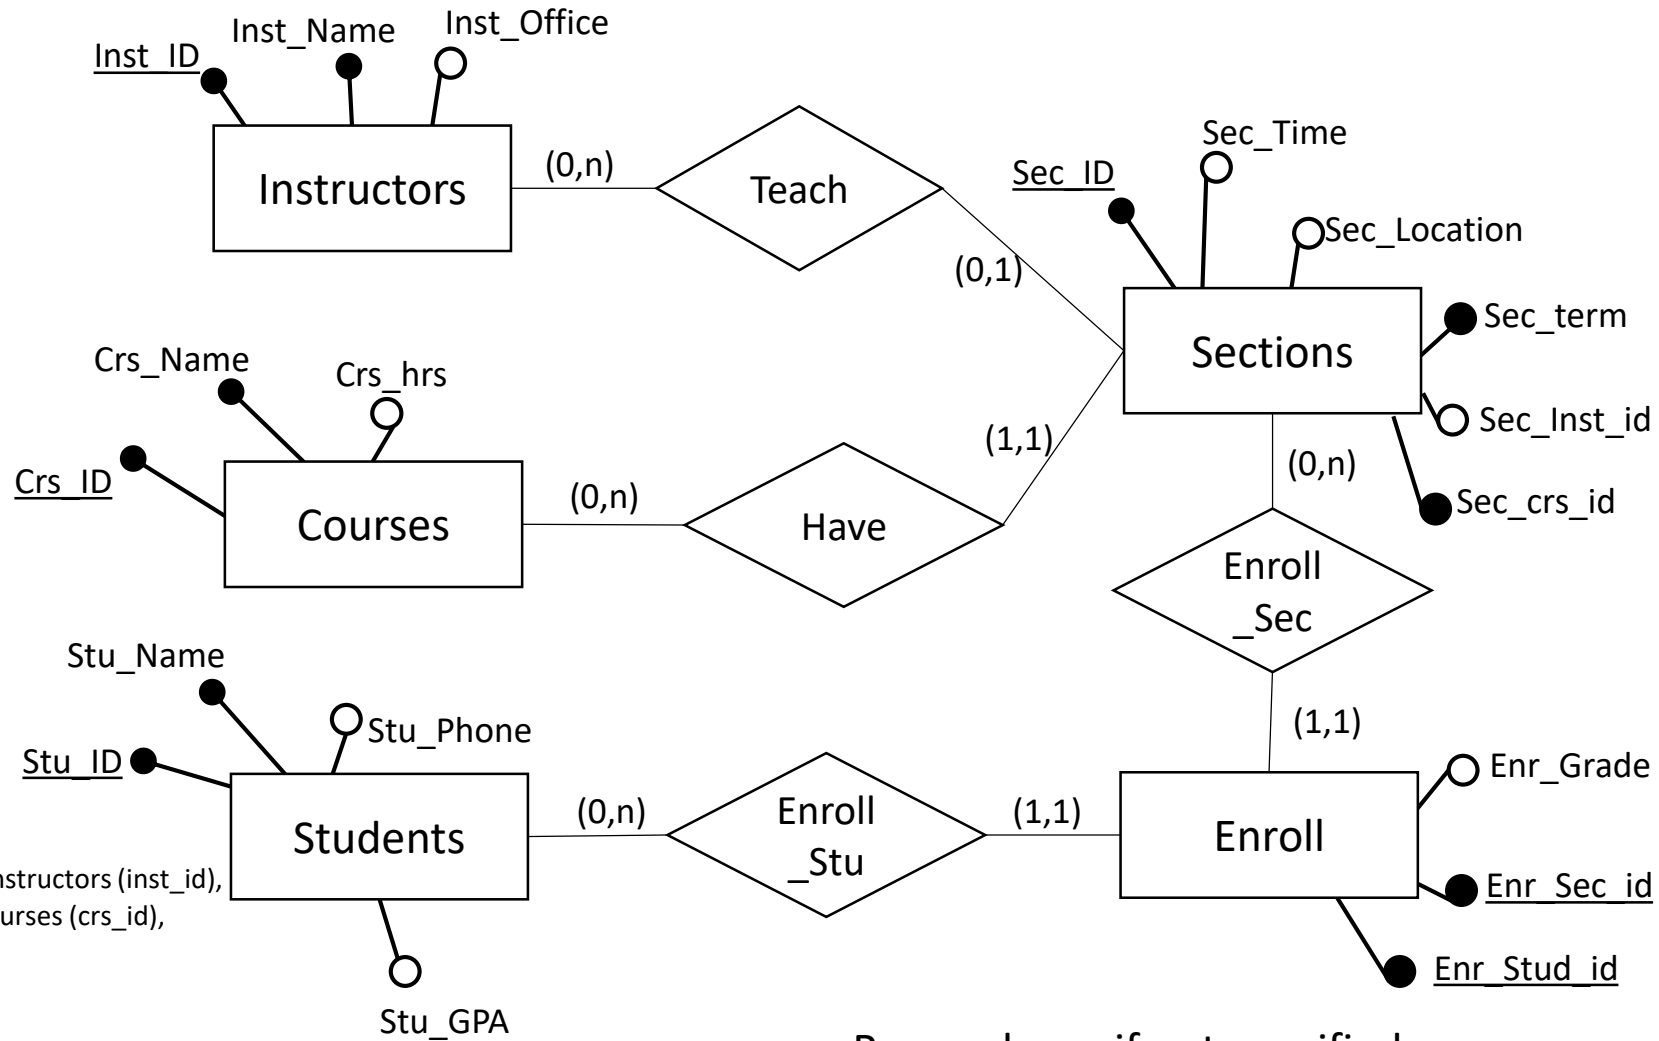
```
CREATE TABLE instructors (
Inst_id      char(15) CONSTRAINT pk_inst PRIMARY KEY,
Inst_name    varchar(50) CONSTRAINT nn_Instname NOT NULL,
Inst_office  varchar(8)
);
```

```
CREATE TABLE courses(
Crs_id      char(8) CONSTRAINT pk_crs PRIMARY KEY,
Crs_name    varchar(50) CONSTRAINT nn_crs NOT NULL,
Crs_hours   numeric(1)
);
```

```
CREATE TABLE sections (
Sec_id      char(5) CONSTRAINT pk_sec PRIMARY KEY,
Sec_time    datetime,
Sec_location nvarchar(10),
Sec_term    char(8) CONSTRAINT nn_secterm NOT NULL,
Sec_inst_id char(15) CONSTRAINT fk_inst FOREIGN KEY REFERENCES instructors (inst_id),
Sec_crs_id  char(8) CONSTRAINT fk_cou FOREIGN KEY REFERENCES courses (crs_id),
CONSTRAINT nn_sec_crs NOT NULL
);
```

```
CREATE TABLE students (
Stu_id      char(15) CONSTRAINT pk_stu PRIMARY KEY,
Stu_name    nvarchar(50) CONSTRAINT nn_stuname NOT NULL,
Stu_phone   char(12),
Stu_GPA     numeric(1,3),
CONSTRAINT chk_gpa CHECK (Stu_GPA <= 4.0)
);
```

```
CREATE TABLE enrollment(
Enr_stu_id  char(15),
Enr_sec_id  char(8),
Enr_Grade   char(2),
CONSTRAINT fk_stu FOREIGN KEY Enr_stu_id REFERENCES students (stu_id),
CONSTRAINT fk_sec FOREIGN KEY Enr_sec_id REFERENCES sections (sec_id),
CONSTRAINT pk_enr PRIMARY KEY (Enr_stu_id, Enr_sec_id),
);
```



Remember – if not specified  
1:n and partial participation is implied

# Progress Quiz Time!

- The Progress Quiz is available in Canvas
  - You MUST complete the quiz on Canvas by 5:00 on Friday – This in-class activity does not count for points!
  - Each week we will discuss the questions, so for those of you that are in class and keeping up with things, you'll have an extra easy time with it!
- Go to <http://kahoot.it> and we'll get started momentarily!



# Exam 1

- 75 minutes, in class, on paper, closed book, individual effort
  - You only need pencil or pen – I will provide the exam, scratch paper, etc.
  - Once you start the exam you cannot leave the room for any reason – plan accordingly (i.e., use the restroom before coming to class)
  - Be on time – once the first person leaves the room, no one new may start the exam
- Will be approximately one third each of:
  - Multiple choice
  - Short answer / Matching / Fill in the Blank / etc.
  - Drawing an ERD
- I won't ask you to WRITE any SQL but you might have to INTERPRET SQL

# Exam 1 (This list of topics is not all inclusive – refer to the slides for more!)

- Topics include anything we have covered up until now:
  - Modules 1, 2, 3, 6, and 10
  - This list of topics is NOT all inclusive – anything we have covered may be on the exam
- All general database concepts
  - Data/information/metadata
  - Schemas
  - File systems vs. databases
- Conceptual data modeling
  - ER Diagrams
  - Characteristics of attributes
  - Cardinality, Participation, Degrees
  - Deletion constraints
  - M:N Relationships / Gerunds

# Exam 1 (This list of topics is not all inclusive – refer to the slides for more!)

- Relational modeling
  - Set operations
  - Relational algebra
  - Keys
  - Foreign Keys / Referential Integrity
  - Logical Models
- Database Creation
  - Data definition language
  - Inserting, updating, and deleting data

Questions, comments, or concerns?

# Exam 1- Commonly missed concepts

- Foreign keys refer to any candidate key – PREFERABLY (and usually) the primary key, but it can be any candidate key
- There is no “child” in a 1:1 relationship – they are equal. In this case we place the foreign key in the entity that has mandatory participation in order to have fewer NULL values
  - If both sides are optional either side is OK, but better to chose the entity with fewer entities OR use a gerund to eliminate NULL vales.
- In a M:N relationship, EITHER side (or neither side, or BOTH) could be considered the parent, since it is made up of two 1:M relationships, in which each entity is the parent and the gerund is the child

# BZAN 6354

## Lecture 6

Go forth and  
do great things!

Dr. Mark Grimes, Ph.D.  
[gmgrimes@bauer.uh.edu](mailto:gmgrimes@bauer.uh.edu)

UNIVERSITY of  
**HOUSTON**

C. T. BAUER COLLEGE of BUSINESS  
Department of Decision & Information Sciences